

# Design and Modeling of a Non-blocking Checkpointing System

Kento Sato<sup>†1,2</sup>, Adam Moody<sup>†3</sup>, Kathryn Mohror<sup>†3</sup>, Todd Gamblin<sup>†3</sup>,  
Bronis R. de Supinski<sup>†3</sup>, Naoya Maruyama<sup>†4,5</sup> and Satoshi Matsuoka<sup>†1,5,6,7</sup>

<sup>†1</sup> Tokyo Institute of Technology

<sup>†2</sup> Research Fellow of the Japan Society for the Promotion of Science

<sup>†3</sup> Lawrence Livermore National Laboratory

<sup>†4</sup> RIKEN Advanced Institute for Computational Science

<sup>†5</sup> Global Scientific Information and Computing Center

<sup>†6</sup> National Institute of Informatics

<sup>†7</sup> JST/CREST



Lawrence Livermore  
National Laboratory


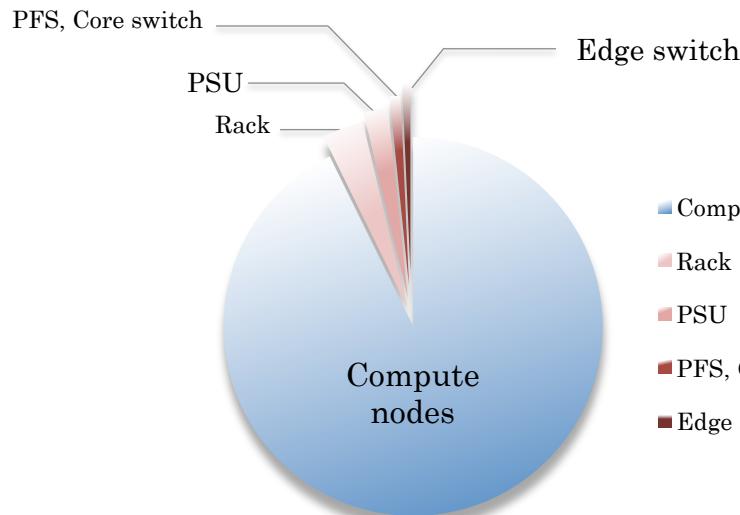


This work performed under the auspices of the U.S. Department of Energy by Lawrence Livermore National Laboratory  
under Contract DE-AC52-07NA27344. LLNL-PRES-599833-DRAFT

November 13<sup>th</sup>, 2012

# Failures on HPC systems

- **Exponential growth in computational power**
  - Enables finer grained scientific simulations
- **Overall failures rate increases accordingly**
  - Due to increasing complexity and system size



TSUBAME2.0, 14<sup>th</sup> in Top500 (June 2012)



2.4 PFlops  
1442 nodes  
2953 CPU sockets  
4264 GPUs  
197 switches  
58 racks

Failure analysis on Tsubame2.0

Period: 1.5 years (Nov 1<sup>st</sup>, 2010 ~ April 6<sup>th</sup> 2012)

Observations: 962 node failures in total

- **System resiliency is becoming more important**
  - Without a viable resilience strategy, applications can not run for even one day on such a large system

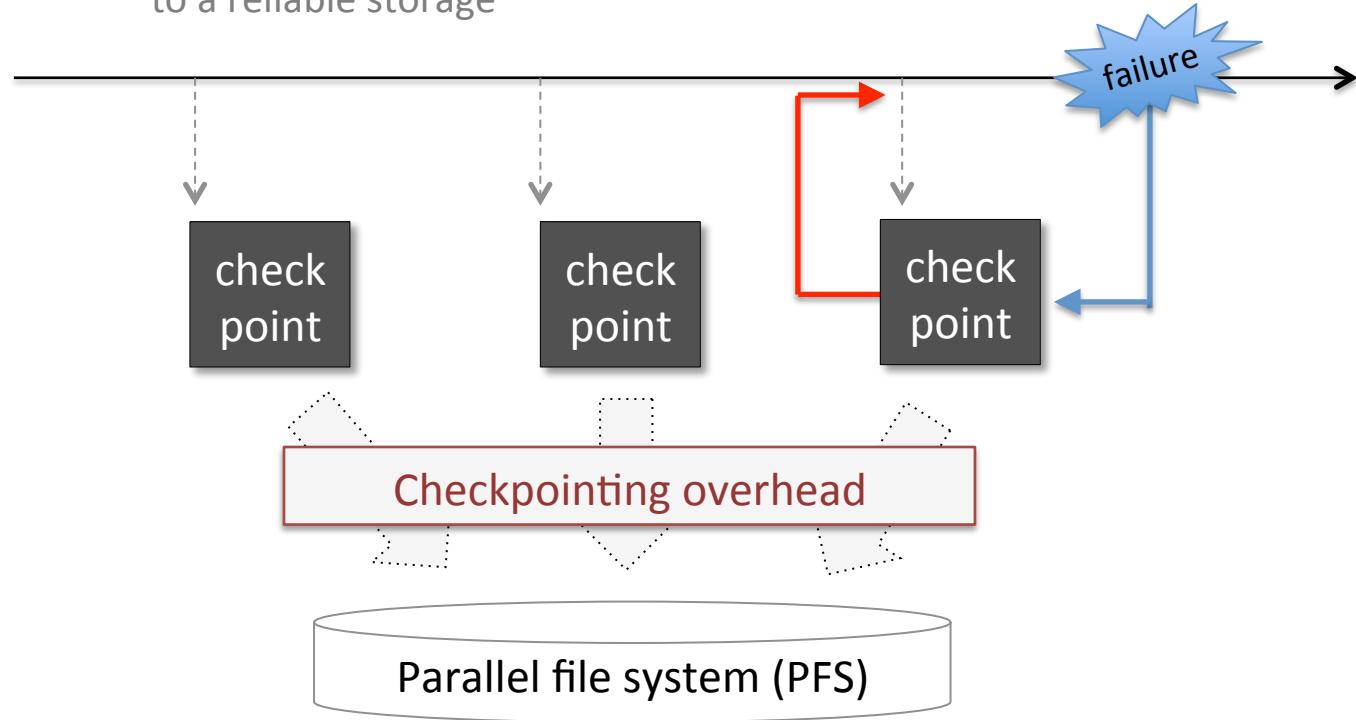
# Traditional Checkpoint/Restart

## Checkpoint

Periodically save a snapshot of an application state to a reliable storage

## Restart

On a failure, restart the execution from the latest checkpoint

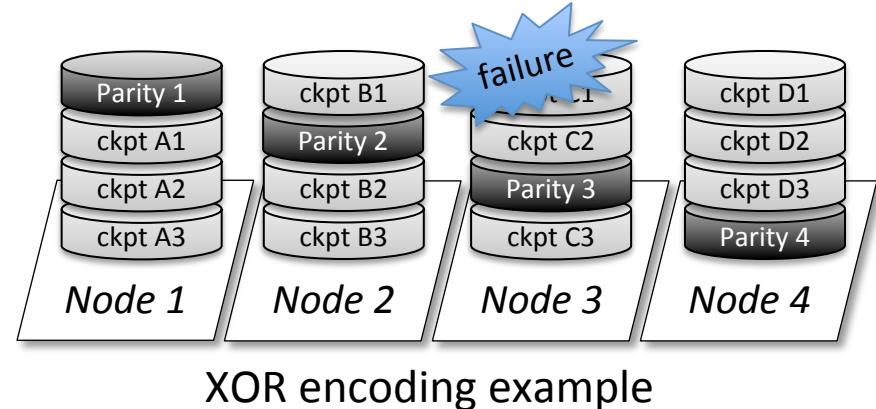


Mostly these checkpoints are stored in the most reliable storage, such as a shared parallel file system(PFS).




# Scalable checkpointing methods

- Diskless checkpoint:
  - Create redundant data across local storages on compute nodes using a encoding technique such as XOR
  - Can restore lost checkpoints on a failure caused by small # of nodes like RAID-5
- Most of failures comes from one node, or can recover from XOR checkpoint
  - e.g. 1) TSUBAME2.0: 92% failures
  - e.g. 2) LLNL clusters: 85% failures



■ LOCAL/XOR/PARTNER checkpoint  
■ PFS checkpoint

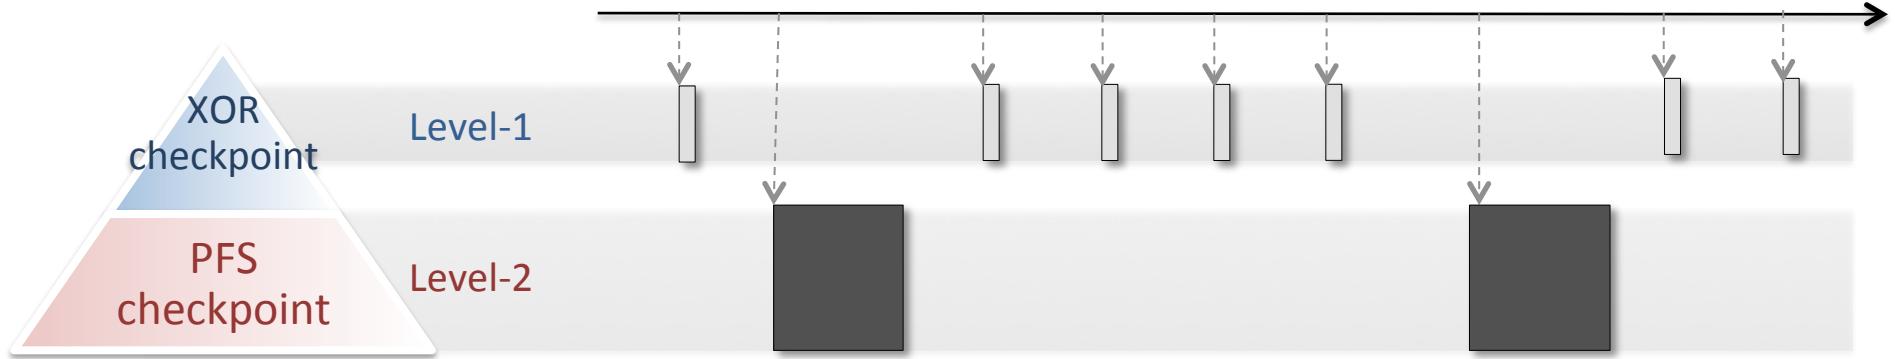
Diskless checkpoint is promising approach



Failure analysis on TSUBAME2.0

Failure analysis on LLNL clusters

# Multi-level checkpointing (MLC)



- Use storage levels hierarchically
  - XOR checkpoint: **Frequently**
    - for **one node** or **a few node** failure
  - PFS checkpoint: **Less frequently**
    - for **multi-node** failure
- **8x efficiency improvement**
  - With MLC implementation called SCR(Scalable Checkpoint/Restart) library developed in LLNL
  - Compared to single-level checkpointing



Source: A. Moody, G. Bronevetsky, K. Mohror, and B. R. de Supinski, "Design, Modeling, and Evaluation of a Scalable Multi-level Checkpointing System," in Proceedings of the 2010 ACM/IEEE International Conference for High Performance Computing, Networking, Storage and Analysis (SC 10).

# MLC Problems on Petascale or larger

## Three potential problems

### 1. PFS checkpoint overhead

- Even with MLC, PFS checkpoint still becomes big overhead


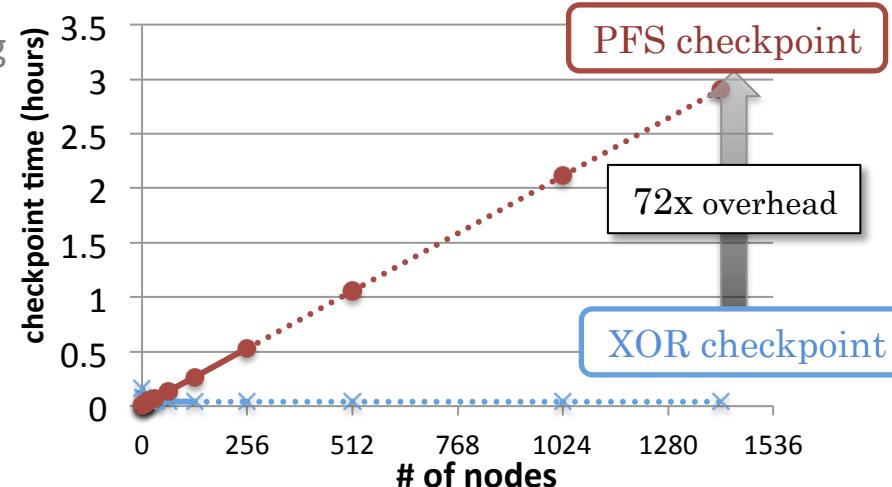
### 2. Inefficient PFS utilization

- Time between PFS checkpoints becomes long, PFS is not utilized during XOR checkpoints

### 3. Failure during PFS checkpoint

- At scale, prolonged PFS checkpointing has a risk of failures during checkpointing

TSUBAME2.0 checkpoint time trend



# Objective, Proposal and Contributions

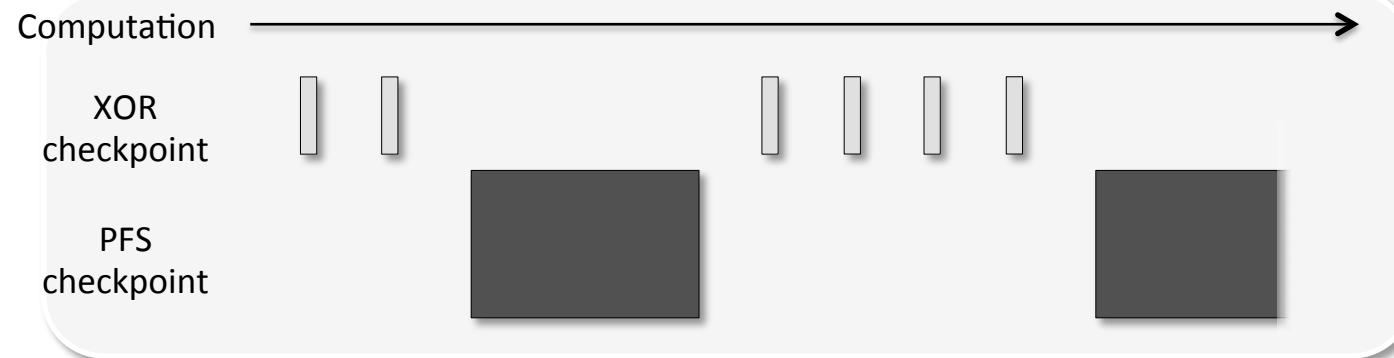
- **Objective:** More efficient MLC
  - Minimize PFS checkpoint overhead
  - Improve PFS utilization
  - Reduce a risk of failure during PFS checkpoint
- **Proposal & Contributions:**
  - Developed a non-blocking checkpointing system as an extension for SCR library
    - PFS checkpoint with  $0.5 \sim 2.5\%$  overhead
  - Modeled the non-blocking checkpointing
    - Determine optimal multi-level checkpoint configuration
    - $1.1 \sim 1.8x$  efficiency on current and future systems

# Outline


- Introduction
- Design of a Non-blocking checkpointing system
- Modeling of the Non-blocking checkpointing
- Evaluation
- Summary

# Non-blocking checkpointing overview

## Blocking multi-level checkpointing



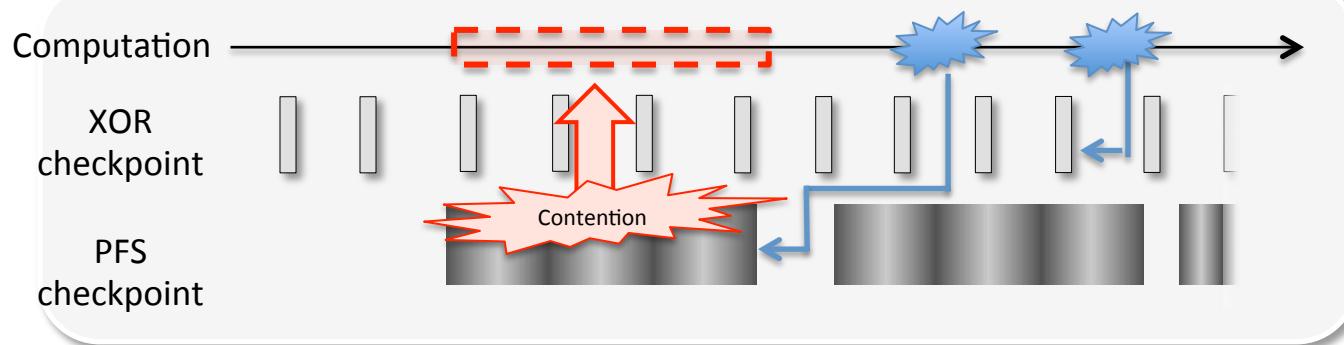
## Non-blocking multi-level checkpointing



- Write PFS checkpoint in the background, minimize overhead
- By initiating next ckpt right after previous one, increase utilization
- Reduce impact of failures requiring XOR checkpoint

# Challenges on Non-blocking checkpointing

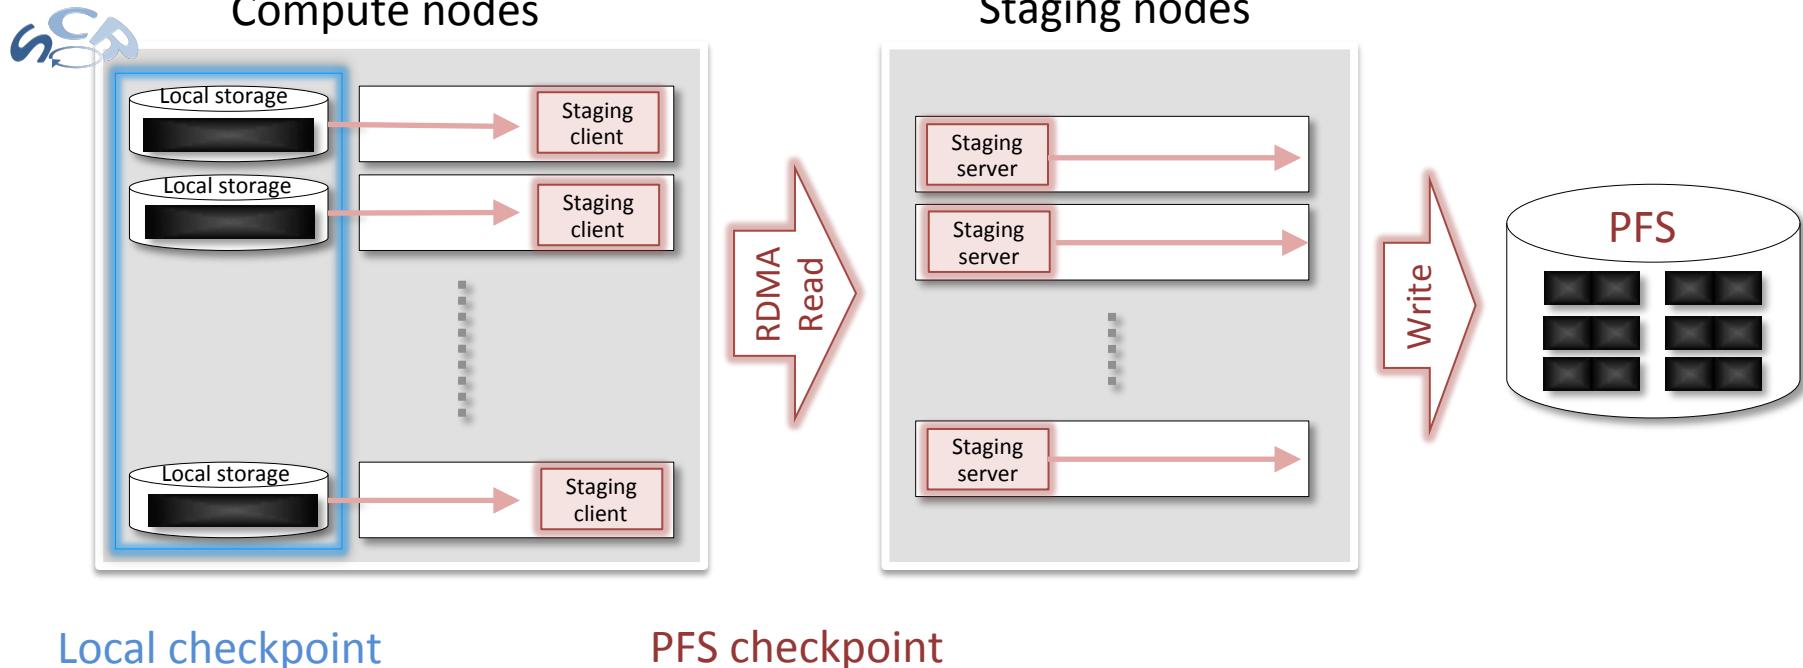
## Non-blocking multi-level checkpointing



- Utilize local SSDs for the additional space
  - Write PFS checkpoint in the background which requires additional storage spaces
- Minimize resource contention
  - PFS checkpointing is running in the background, inflate the runtime due to resource contention  
⇒ Implementation: Use RDMA with checkpoint dedicated nodes
- Optimize configuration (e.g. checkpoint interval)
  - On a failure requiring PFS, need “complete PFS checkpoint”
  - On a failure requiring XOR, need to restore both XOR & PFS ckpt being written  
⇒ Modeling: Model a non-blocking multi-level checkpoint

# Non-blocking checkpointing overview

- Between compute nodes and PFS, use staging nodes
  - Dedicated extra nodes for transferring local checkpoints written by a SCR library
  - Read checkpoints from compute nodes using RDMA, write out to a PFS



# Non-blocking checkpointing using RDMA

## 1. Local storages to Local memory


- After SCR writes checkpoint to a local storage, staging clients running on compute nodes read chunks of the checkpoint from the local storage to a buffer memory

## 2. Local memory to Remote memory

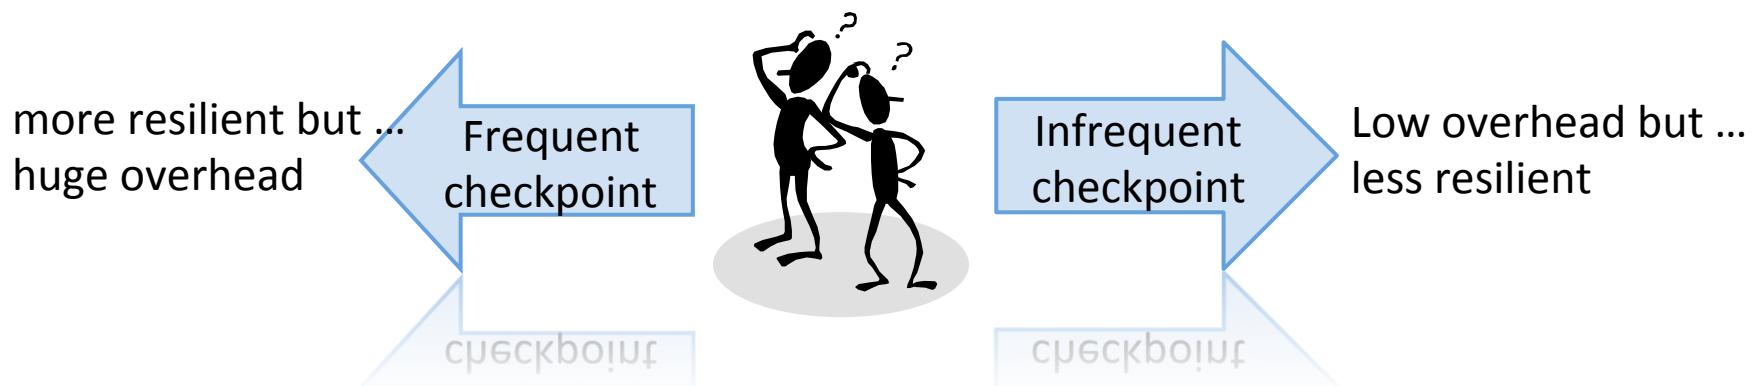
- Send RDMA Read requests to a mapped staging server running on a staging node, staging server read the checkpoints from the buffer using RDMA

## 3. Remote memory to PFS

- Data writer threads running on Staging nodes write checkpoint chunks to PFS in parallel




# Modeling of Non-blocking checkpoint



# Outline


- Introduction
- Design of Non-blocking checkpointing system
- Modeling of Non-blocking checkpointing
- Evaluation
- Summary




# Non-blocking MLC model overview

- Describe an application's state transitions as Markov model
- **Input** (each level of ..)
  - Checkpoint time
  - Restart time
  - Failure rate
  - Interval
- **Output**
  - Expected runtime
- Find checkpoint intervals that minimize runtime

Non-blocking multi-level checkpointing




Non-blocking multi-level checkpoint model



# Assumptions on the model


- Independent and identically distributed failure rate & Poisson distribution
  - One failure does not increase the probability of successive failures
- Stable write & read performance
  - Checkpoint/Restart time significantly does not change during overall the runtime
- Failure on Level- $k$  recovery => Level- $(k+1)$  checkpoint
  - Another one node failure during XOR recovery requires a PFS checkpoint
  - Assume PFS checkpoint can retry infinitely
- Saved checkpoints are never lost on non-failed nodes and a PFS
  - Guarantee failed job can restart from the latest checkpoint



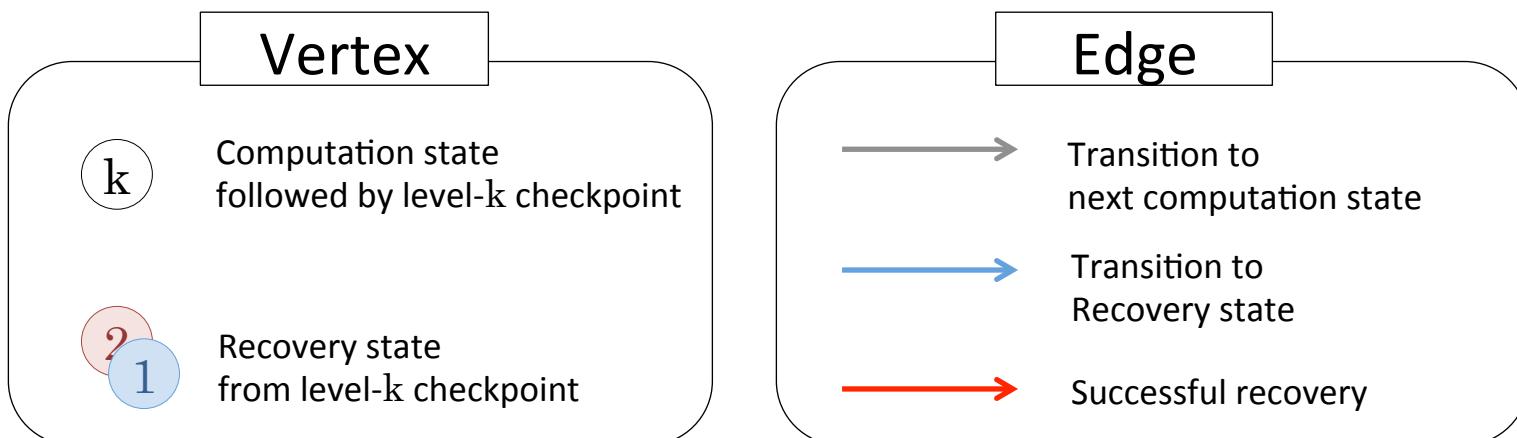
# Two-level checkpoint example

- For simplicity, two-level checkpoint
  - Level-1: XOR checkpoint
  - Level-2: PFS checkpoint


Non-blocking multi-level checkpointing




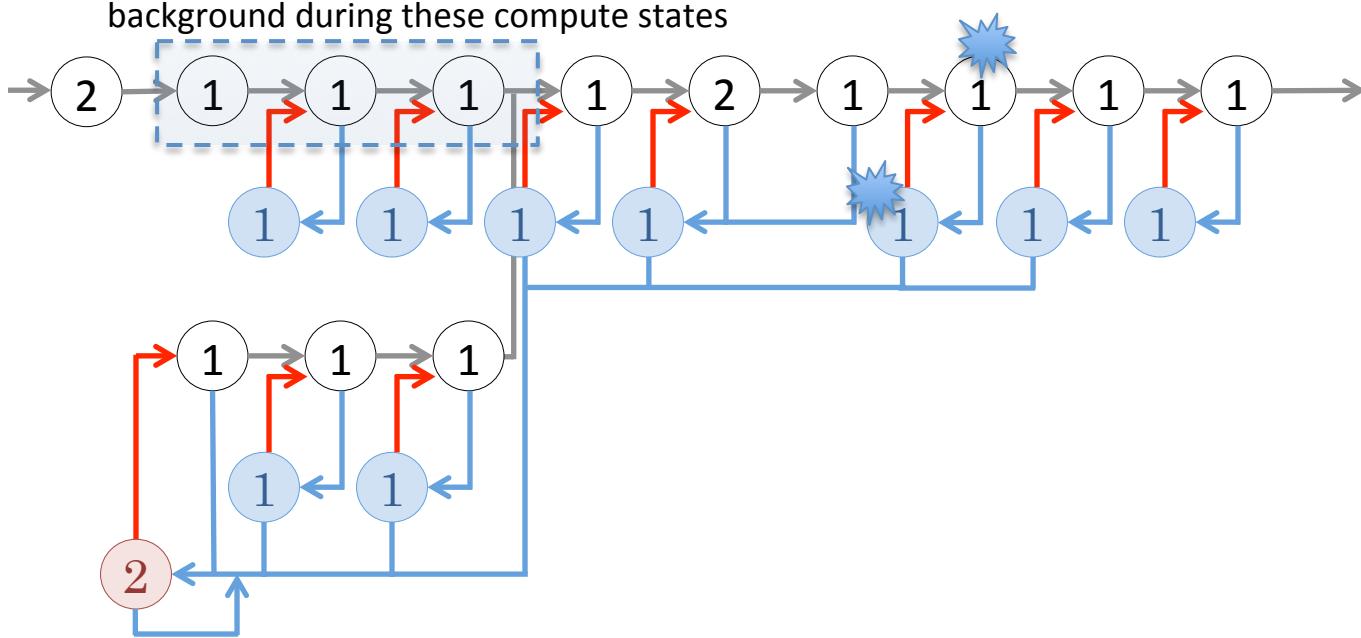
- Describe state transitions as Markov model




# No failure & Level-“1” failure case



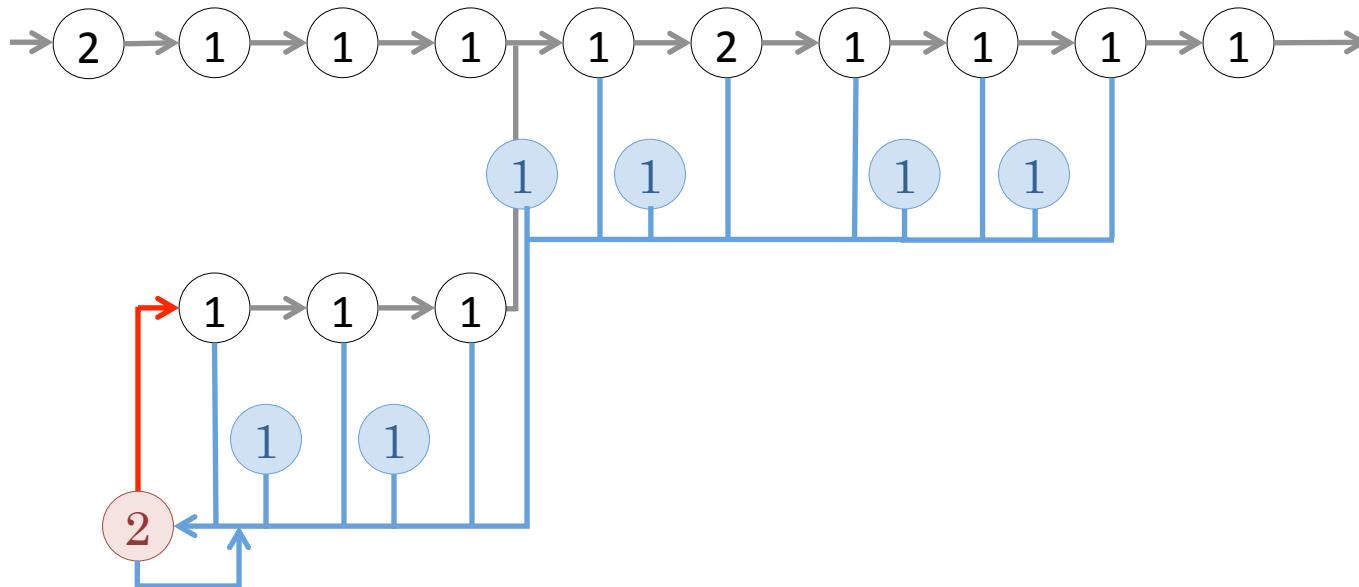
$$\text{L-2 ckpt: L-1 ckpt} = 1: 4$$



# Level-“2” failure case




L-2 ckpt: L-1 ckpt  
= 1: 4



- $k$  Compute state (t sec) + level-k checkpoint
- Successful computation
- $2 \rightarrow 1$  Recovery state from level-k checkpoint
- Recover transition
- Successful recovery

# How to calculate *expected\_runtime* ?



$t$  : Interval  
 $C_c$  :  $c$ -level checkpoint time  
 $r_c$  :  $c$ -level recovery time

|            | Duration  |   |  |
|------------|---|---|--|
|            | $t + c_k$   | $r_k$   |  |
| No failure | <p><math>p_0(t + c_k)</math></p> <p><math>t_0(t + c_k)</math></p> | <p><math>p_0(r_k)</math></p> <p><math>t_0(r_k)</math></p> |  |
| Failure    | <p><math>p_i(t + c_k)</math></p> <p><math>t_i(t + c_k)</math></p> | <p><math>p_i(r_k)</math></p> <p><math>t_i(r_k)</math></p> |  |

$$\begin{aligned}
 p_0(T) &= e^{-\lambda T} \\
 t_0(T) &= T \\
 p_i(T) &= \frac{\lambda_i}{\lambda} (1 - e^{-\lambda T}) \\
 t_i(T) &= \frac{1 - (\lambda T + 1) \cdot e^{-\lambda T}}{\lambda \cdot (1 - e^{-\lambda T})}
 \end{aligned}$$

$\lambda_i$  :  $i$ -level checkpoint time



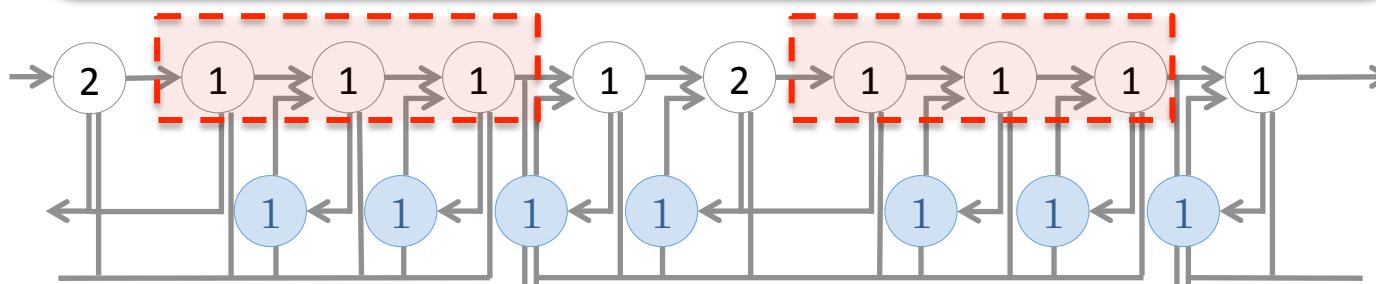

$\left[ \begin{array}{l} p_0(T) \\ t_0(T) \end{array} \right]$  : No failure for  $T$  seconds  
 $\left[ \begin{array}{l} p_i(T) \\ t_i(T) \end{array} \right]$  : Expected time when  $p_0(T) = 0$

$\left[ \begin{array}{l} p_i(T) \\ t_i(T) \end{array} \right]$  :  $i$ -level failure for  $T$  seconds  
 $\left[ \begin{array}{l} p_0(T) \\ t_0(T) \end{array} \right]$  : Expected time when  $p_i(T) = 0$

$$\lambda = \sum \lambda_i$$

# Overhead factor: $\alpha$

- Quantify an overhead by our proposed non-blocking checkpointing system




During these compute states PFS checkpointing is running in the background, inflate the runtime due to resource contention

A diagram of a single node with three stages. A dashed red box encloses the middle stage, which contains a circle labeled "k". An arrow points from the middle stage to the right. Below the node, there are two equations:

$$p_0(t + c_k + \alpha \cdot t)$$
$$t_0(t + c_k + \alpha \cdot t)$$

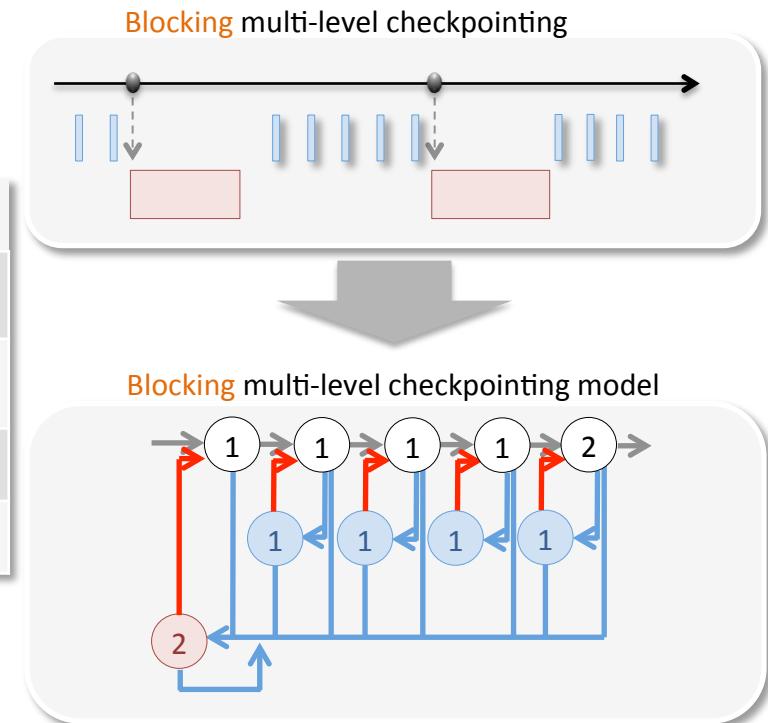
# Arbitrary $N$ -level checkpointing model



# Non-blocking vs. Blocking MLC checkpointing

- **Benchmark:** Himeno benchmark
  - Stencil application solving Poisson's equation using Jacobi iteration method
- **Target System:**  
TSUBAME2.0 Thin nodes (1408 nodes)

|               |  |
|---------------|--|
| CPU           | Intel Xeon X5670 2.93GHz (6cores x 2 sockets)                |
| Memory        | DDR3 1333MHz (58GB)  |
| Network       | Mellanox Technologie<br>Dual rail QDR Infiniband 4x (80Gbps) |
| Local storage | 120GB Intel SSD (RAID0/60GBx2)                               |
| PFS           | Lustre (/work0 )   |



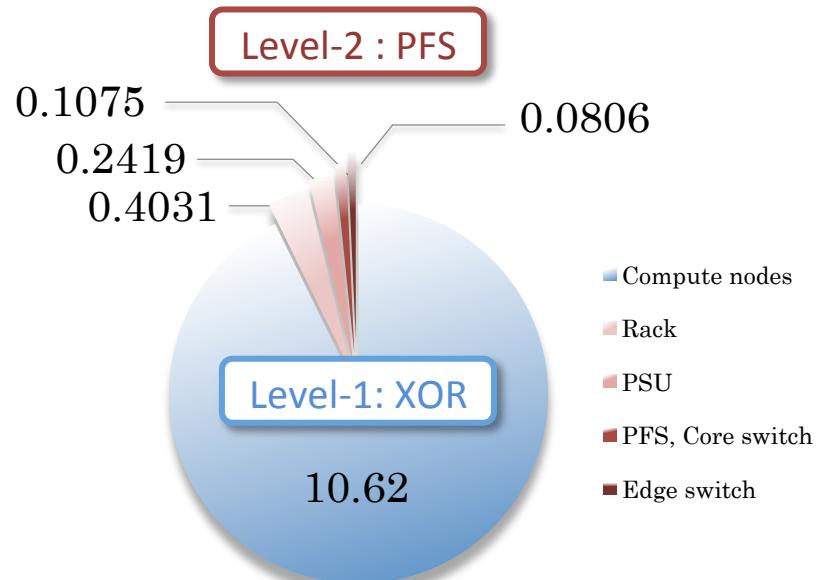
- **Checkpoint Level:** Two-level
  - Level-1: XOR using local SSD
  - Level-2: PFS using Lustre

Source: A. Moody, G. Bronevetsky, K. Mohror, and B. R. de Supinski, "Design, Modeling, and Evaluation of a Scalable Multi-level Checkpointing System," in Proceedings of the 2010 ACM/IEEE International Conference for High Performance Computing, Networking, Storage and Analysis (SC 10).

# Model Parameters

- **Failure rates**

- 1.5 years (Nov 1<sup>st</sup> 2010 ~ Apr 6<sup>th</sup> 2012) failure history




- **Checkpoint size per node: 29GB**

- TSUBAME nodes memory: 58GB

Failure rates (failures/week) on TSUBAME2.0

## Level-1

- **XOR throughput: 400MB/s**



# Model Parameters

- **Failure rates**

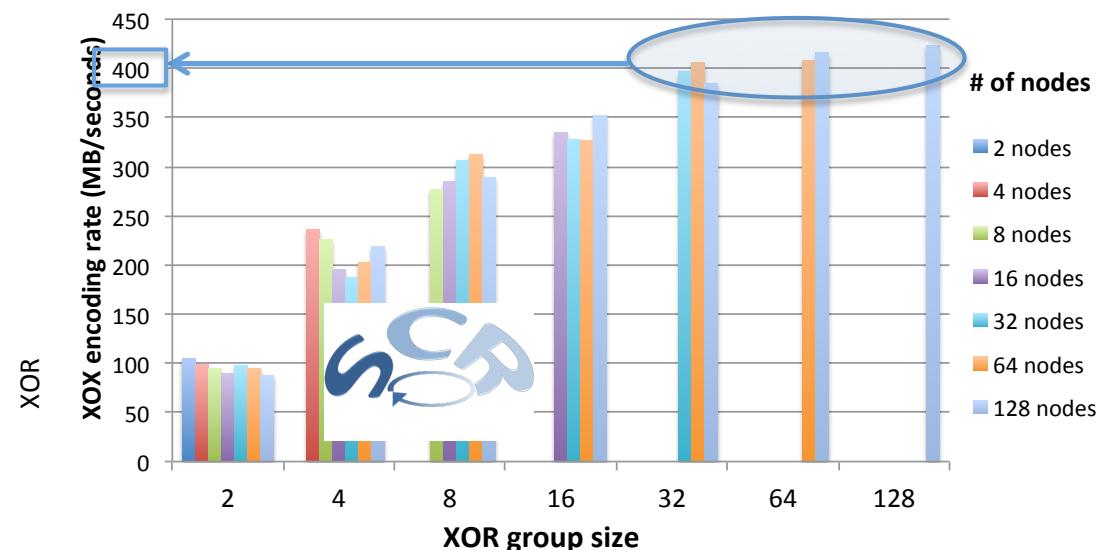

- 1.5 years (Nov 1<sup>st</sup> 2010 ~ Apr 6<sup>th</sup> 2012) failure history

- **Checkpoint size per node: 29GB**

- TSUBAME nodes memory: 58GB


## Level-1

- **XOR throughput: 400MB/s**



# Staging node tuning for TSUBAME2.0

- **# of Staging nodes:** 32 nodes
  - 2.3% of TSUBAME2.0 thin nodes (1408 nodes)
- **PFS throughput:** 6.7GB/seconds
  - 209.5 MB/seconds\* per Staging node
    - \*  $6.7(\text{GB/s}) / 32(\text{nodes}) = 209.5$



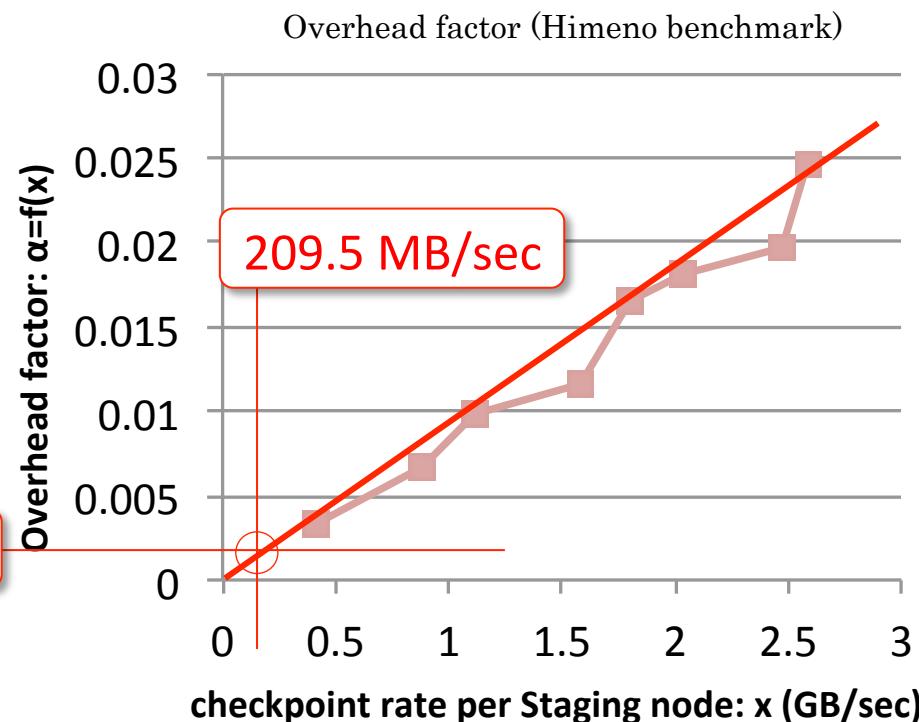
# Overhead factor

- **Overhead factor:** 0.00184 (0.184%)

- For Himeno bechmark

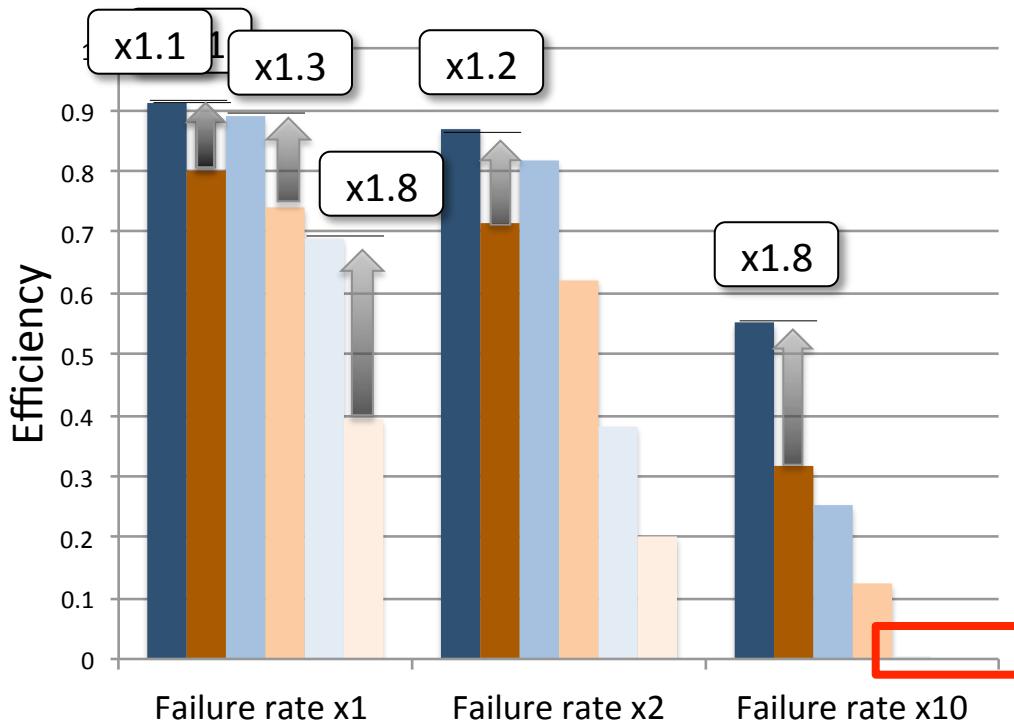
RDMA  $\Rightarrow$  No CPU cycle, No redundant memcpy

RDMA read speed  $\Rightarrow$  209.5MB/s < Network & Memory bandwidth



# Efficiency: Non-blocking vs. blocking

The non-blocking method always achieves higher efficiency than the blocking method



$$\text{Efficiency} = \frac{\text{ideal runtime}}{\text{expected runtime}}$$

*ideal runtime* : No failure and No checkpoint

*expected runtime* : Computed by the models

- PFS cost x1 / Non-blocking
- PFS cost x1 / Blocking
- PFS cost x2 / Non-blocking
- PFS cost x2 / Blocking
- PFS cost x10 / Non-blocking
- PFS cost x10 / Blocking


One TSUBAME2.0 node MTBF: 2.57 years  
# of Nodes: 1408 nodes

x10 scale-out  
→

No computation progresses !!

# Overhead factor: Non-blocking vs. Blocking

Other applications case whose overhead factor becomes bigger



If overhead factor is over 0.2, blocking checkpointing can become more efficient in current system

System scale-out

- Fx1, Cx1, Non-blocking
- Fx1, Cx1, Blocking
- Fx2, Cx2, Non-blocking
- Fx2, Cx2, Blocking
- Fx2, Cx10, Non-blocking
- Fx2, Cx10, Blocking
- Fx10, Cx2, Non-blocking
- Fx10, Cx2, Blocking


F: Failure rate, C: PFS cost

In future systems where the failure rates and cost increase, non-blocking checkpointing can be effective even with a large overhead factor.  
=> Blocking checkpoint overhead dominate the runtime more than overhead factor by non-blocking

# Required PFS performance to meet given application efficiency

When building a reliable data center or supercomputer, two major concerns are monetary cost of the PFS and the PFS throughput required to maintain high efficiency ...

=> predict required PFS performance with the models



# Conclusion

- Developed an non-blocking checkpointing system
  - Write checkpoint data in the background using RDMA
- Markov model of the non-blocking checkpointing
  - Optimal multi-level checkpoint interval
  - Non-blocking v.s. Blocking checkpoint
    - Higher efficiency (1.1 ~ 1.8x) on current and future systems
  - High efficiency (up to 80%) with low PFS throughput

# Q & A

## Speaker:

Kento Sato (佐藤 賢斗)

kent@matsulab.is.titech.ac.jp

Tokyo Institute of Technology (Tokyo Tech)

*Research Fellow of the Japan Society for the Promotion of Science*

[http://matsu-www.is.titech.ac.jp/~kent/index\\_en.html](http://matsu-www.is.titech.ac.jp/~kent/index_en.html)

## Co-authors

Adam Moody, Kathryn Mohror, Todd Gamblin, Bronis R de. Supinski,  
Naoya Maruyama, Satoshi Matsuoka

## Acknowledgement

This work performed under the auspices of the U.S. Department of Energy by Lawrence Livermore National Laboratory under Contract DE-AC52- 07NA27344. LLNL-PRES-599833-DRAFT. This work was also supported by Grant-in-Aid for Research Fellow of the Japan Society for the Promotion of Science (JSPS Fellows) 24008253, and Grant-in-Aid for Scientific Research S 23220003.