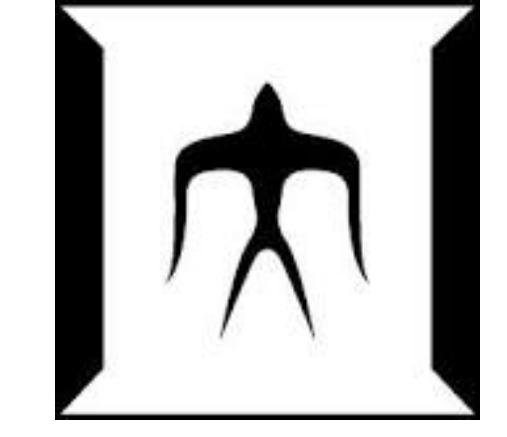


# I/O acceleration with GPU for I/O-bound applications

Kento Sato<sup>†1, †2</sup>, Akira Nukada<sup>†1</sup>, Naoya Maruyama<sup>†3, †1, †5</sup> and Satoshi Matsuoka<sup>†1, †4, †5</sup>



†1 Tokyo Institute of Technology, †2 Research Fellow of the Japan Society for the Promotion of Science,

†3 RIKEN Advanced Institute for Computational Science, †4 National Institute of Informatics, †5 JST/CREST

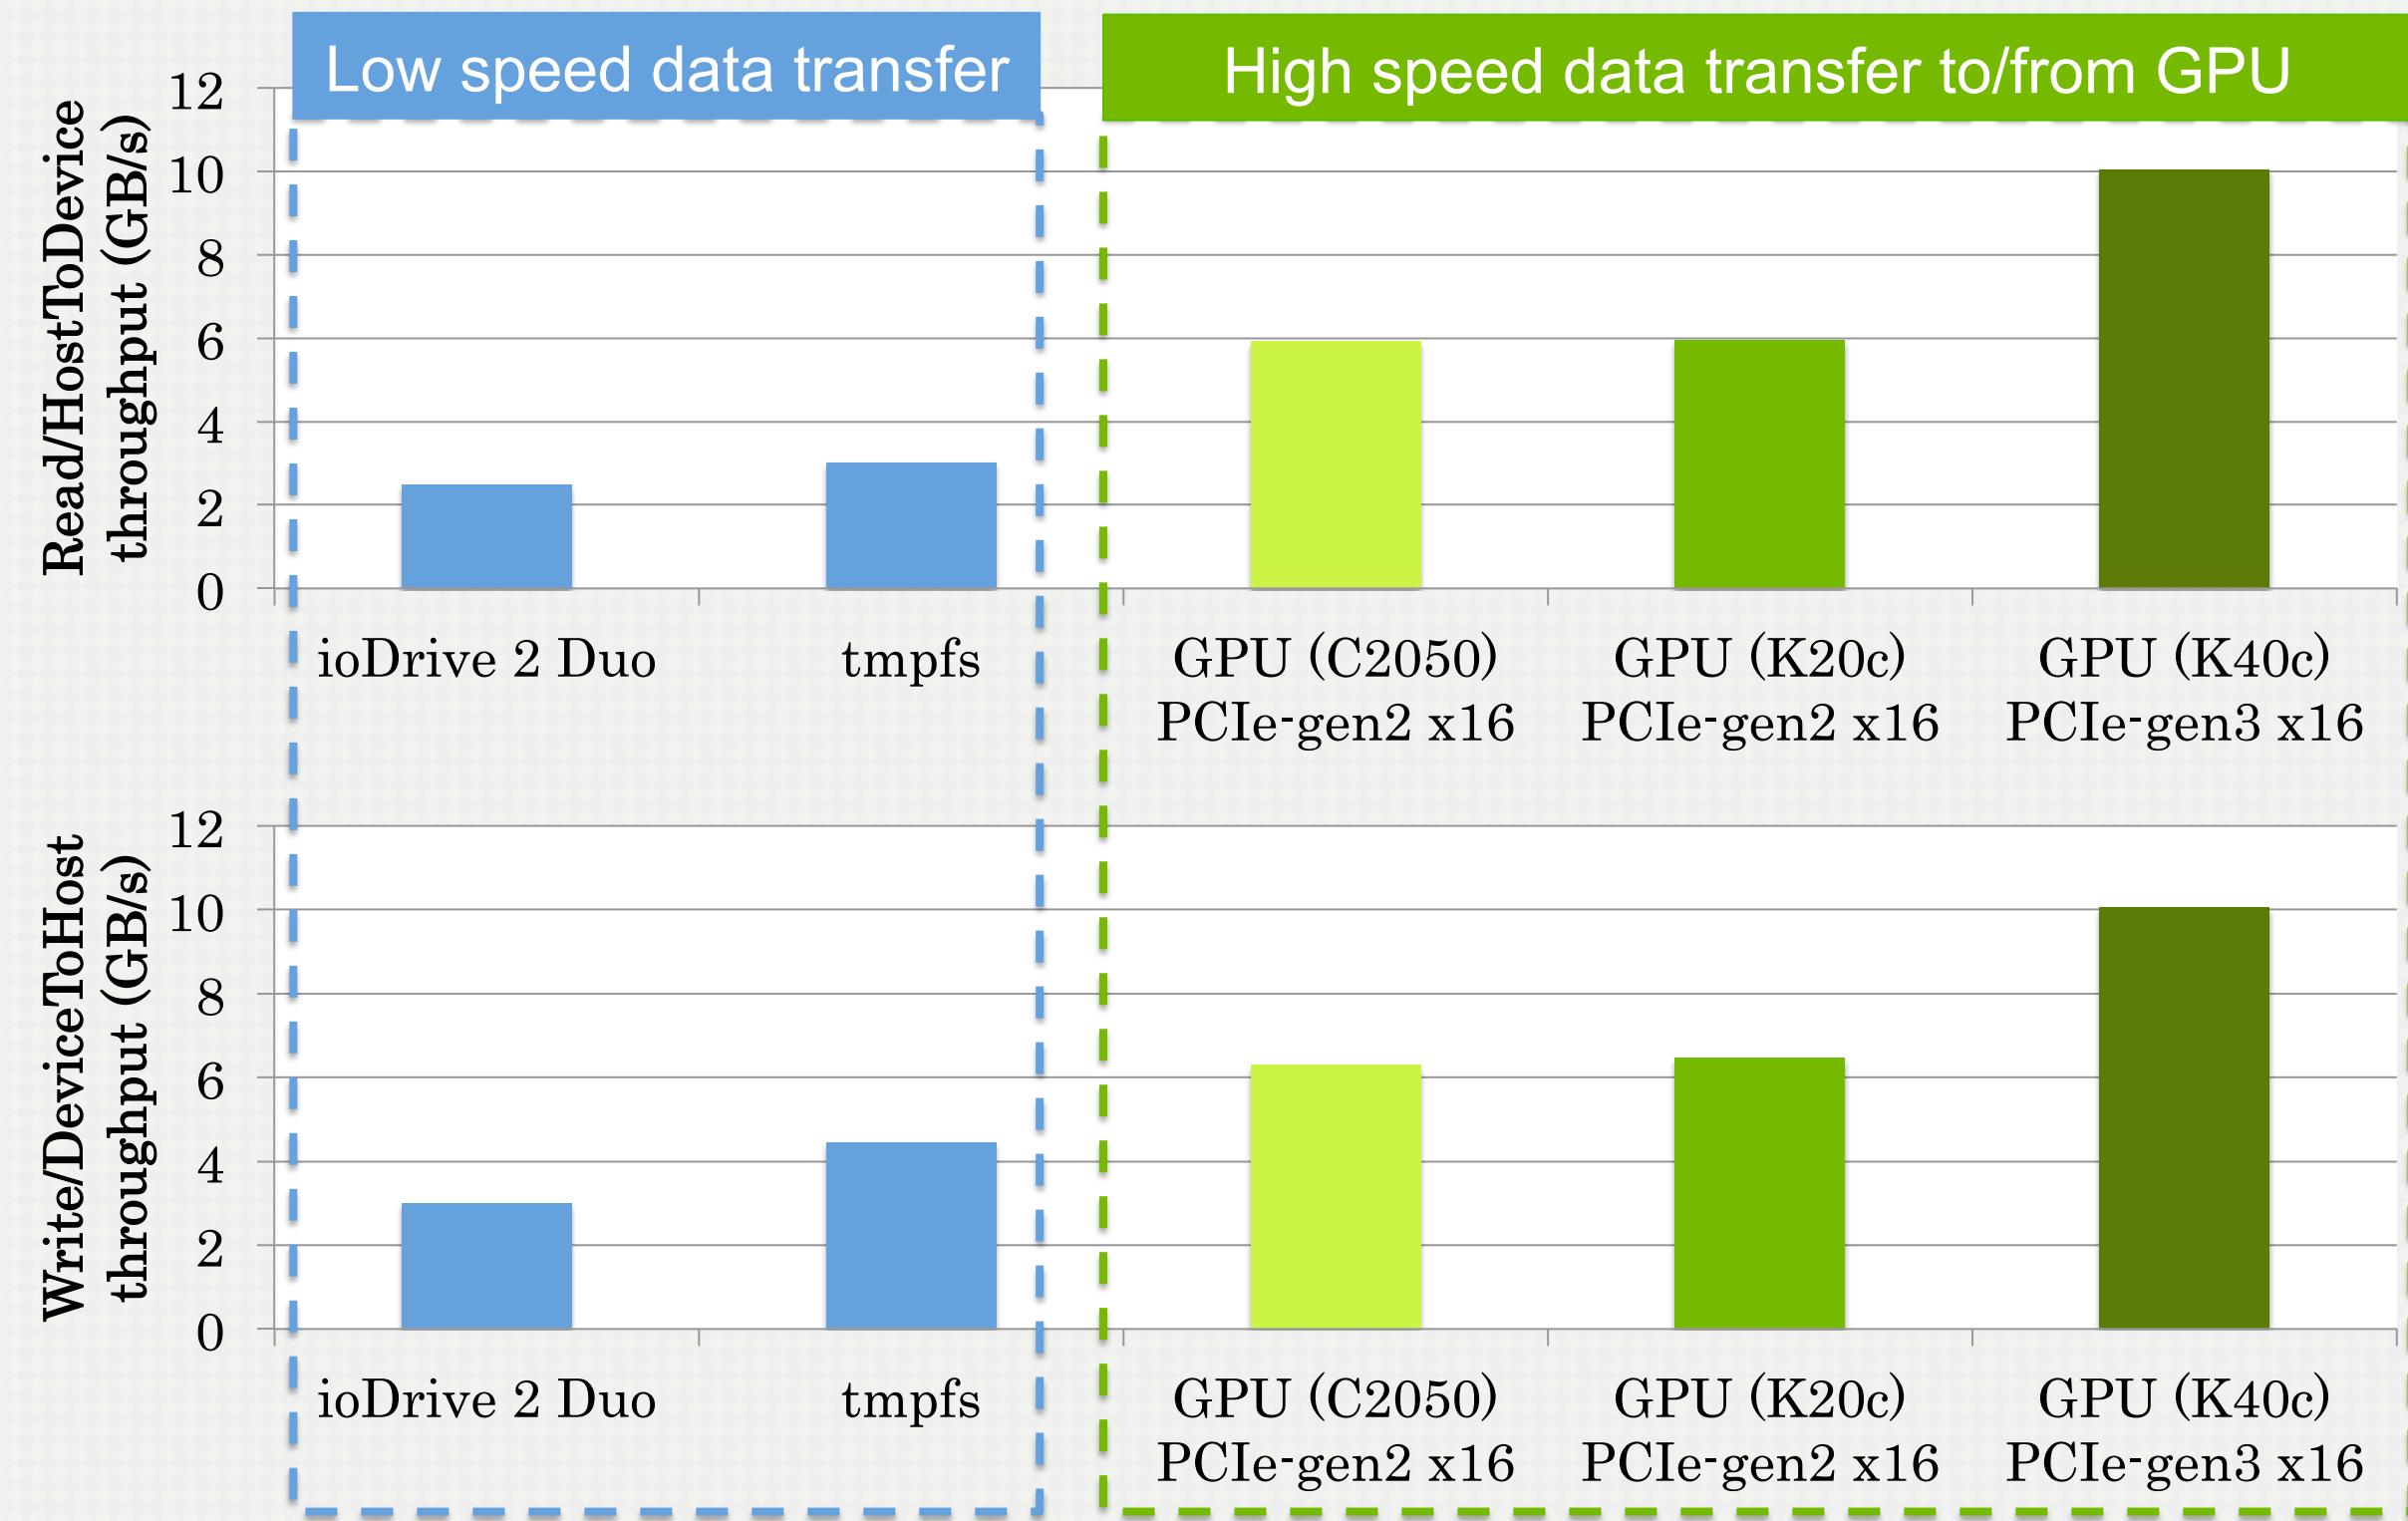
## Background

### GPU-enabled supercomputers

- Several supercomputers usually have GPUs on each compute node to accelerate computation
- Not all applications can be accelerated by GPUs
- Performance of I/O-bound applications is limited by underlining I/O device performance
- Such I/O-bound applications require more I/O bandwidth rather than computational power
- If we execute such non-GPU applications, the GPUs is not utilized, and we waste the resources

### Data transfer throughputs of GPUs

- Data transfer throughputs with POSIX read/write for ioDrive 2 Duo & tmpfs, and with cudaMemcpy for GPUs



- Flash devices on ioDrive can become a bottleneck
- A file system becomes a bottleneck
- GPUs can exploit PCIe's bandwidth
- PCIe-attached GPUs can be used as I/O devices for accelerating I/O-bound applications

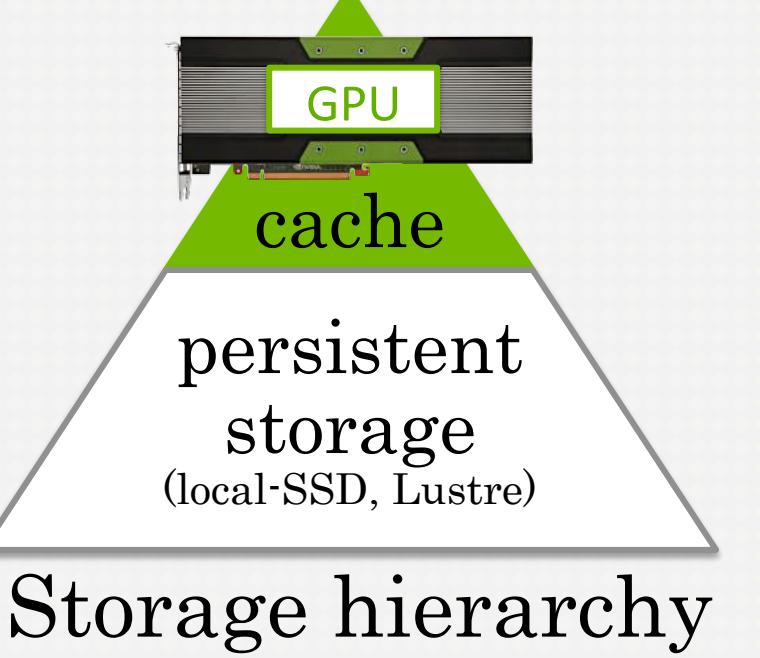
### Objective

- Accelerate I/O operations using GPU device memory
- Exploit PCIe's bandwidth even with I/O operations
- Provide well-known I/O interfaces

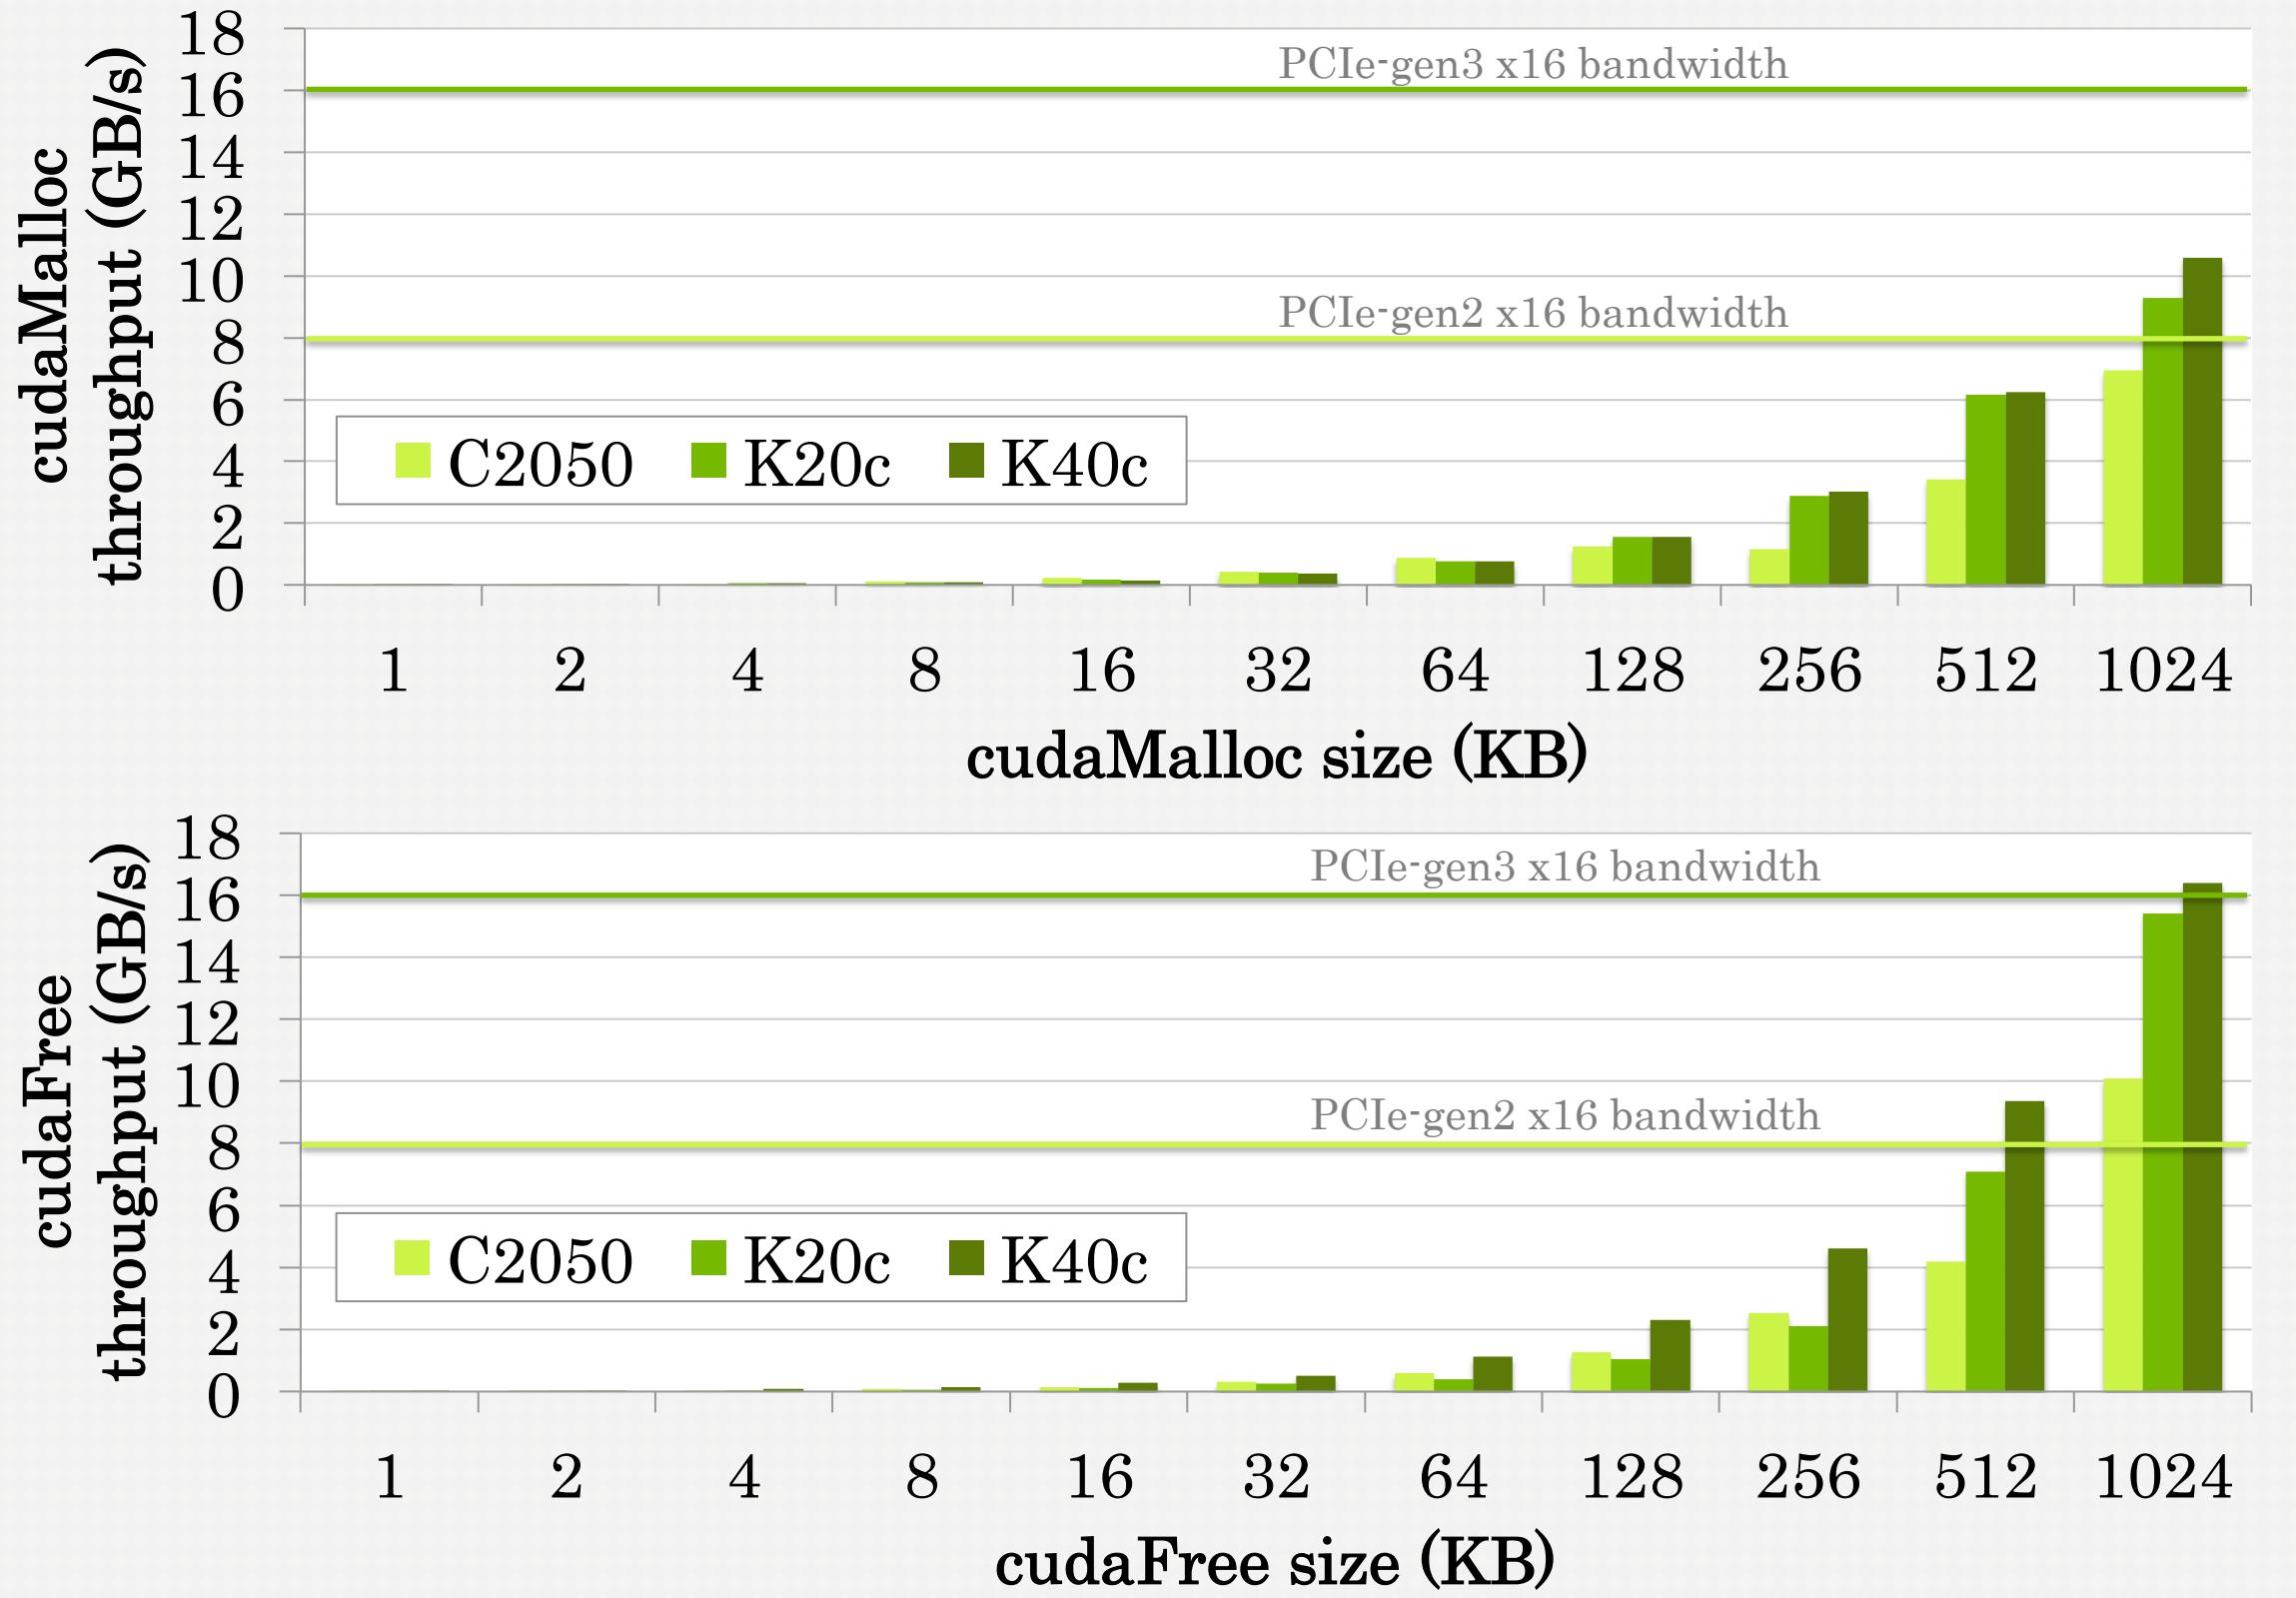
## Methods

### gmfs: GPU-accelerated I/O interface

- gmfs is an I/O interface which uses GPU as buffer cache of persistent storage
- gmfs provides well-known I/O interfaces (open, read, write and close) on top of the hierarchical storage



### Requirement for self memory management

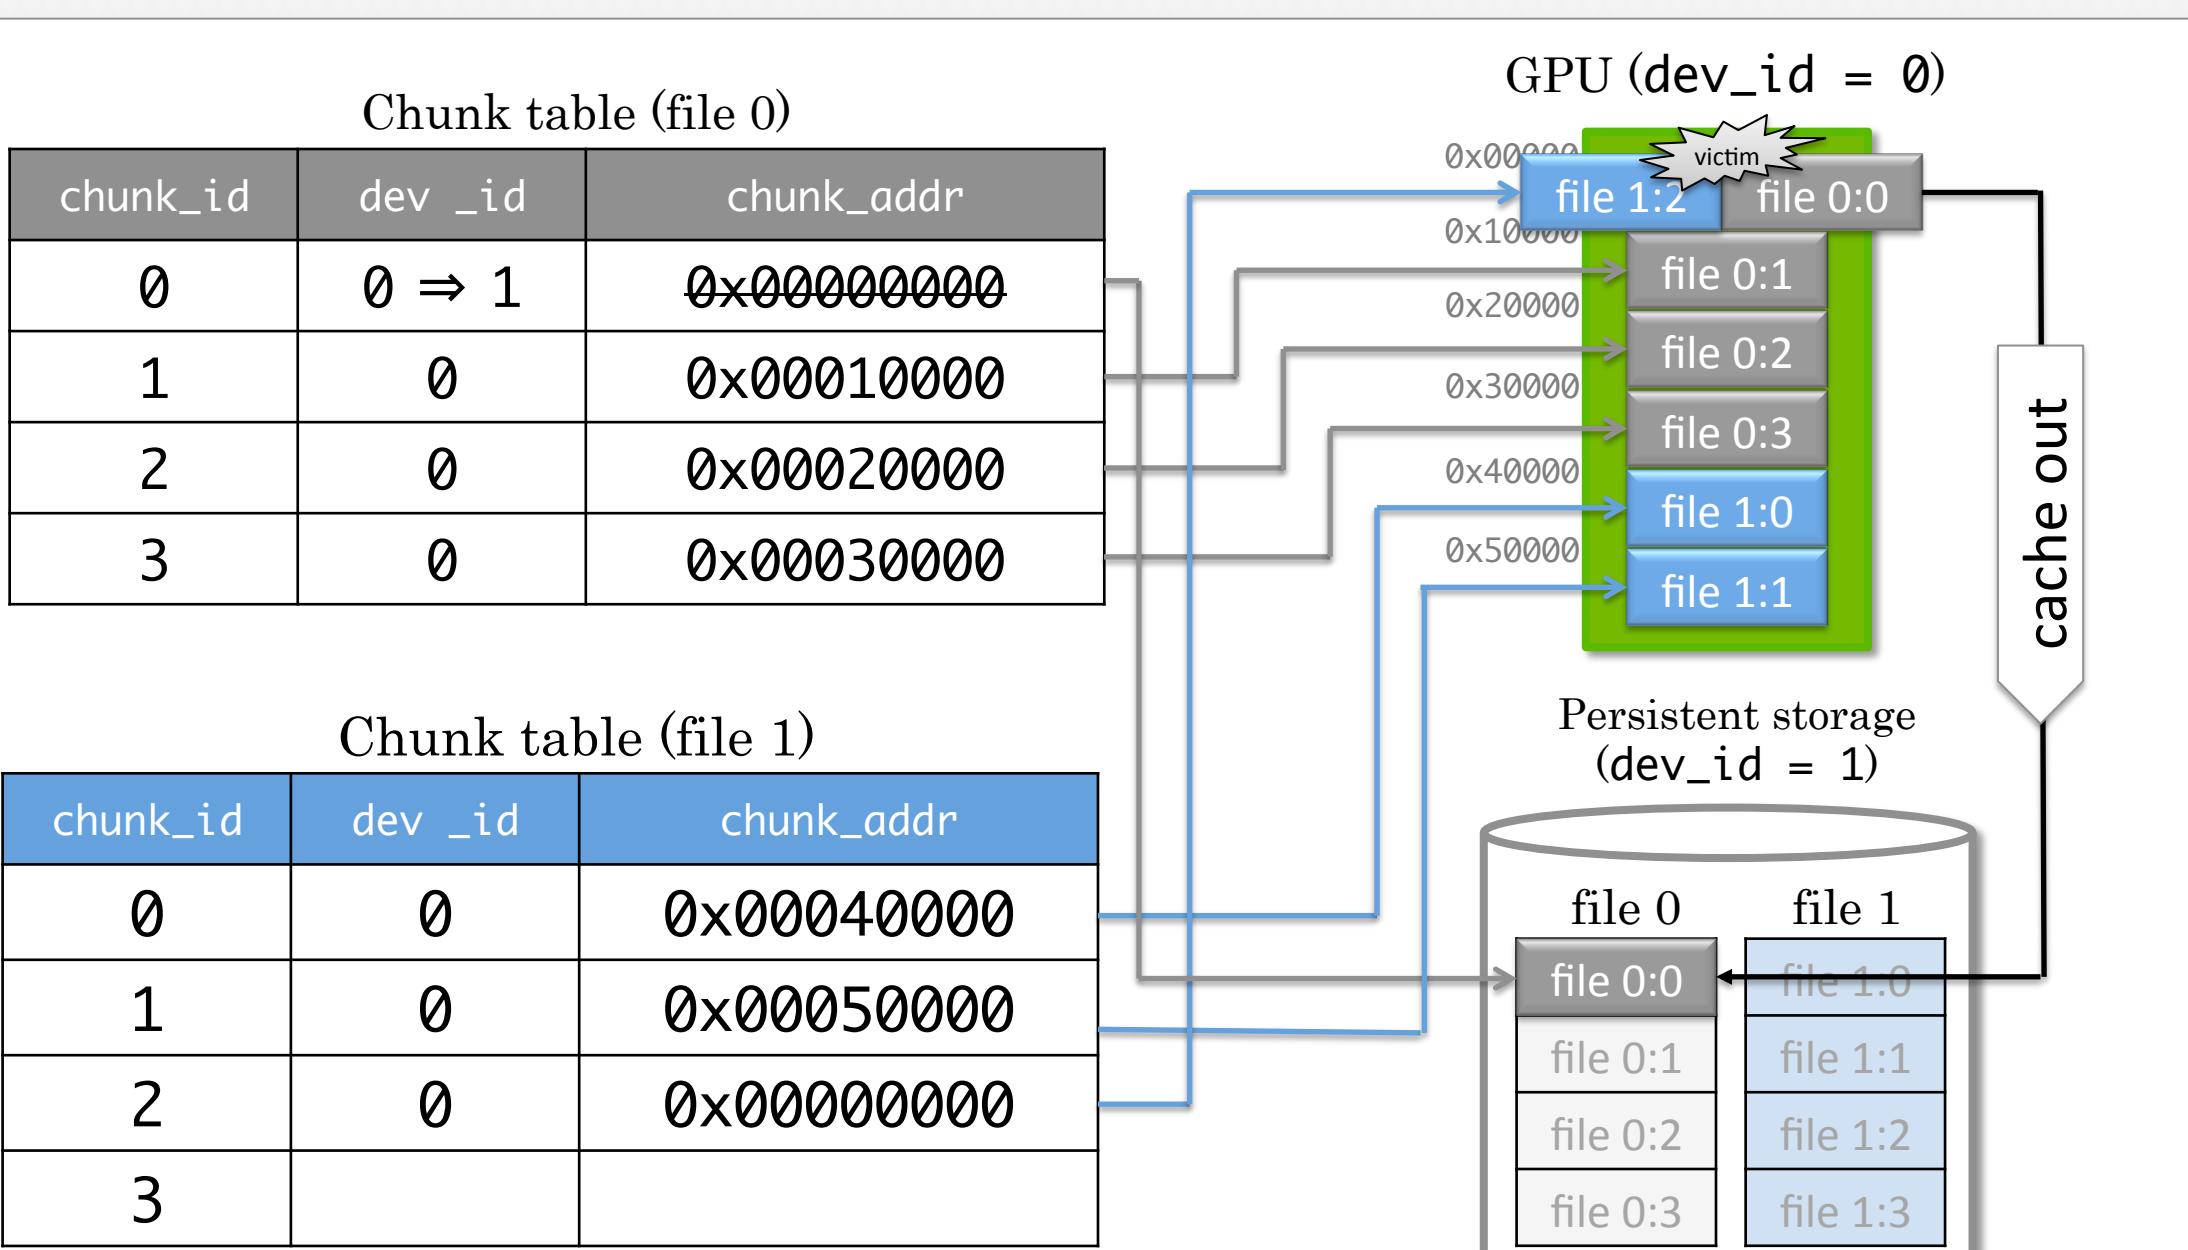


### Challenges

How to exploit PCIe's bandwidth on multiple small writes requiring small cudaMallows and cudaFrees

### gmfs's file management

- First allocate entire GPU memory to gmfs
- gmfs divides the GPU memory into chunks, and chunk tables manage chunk locations (round-robin cache mgmt.)



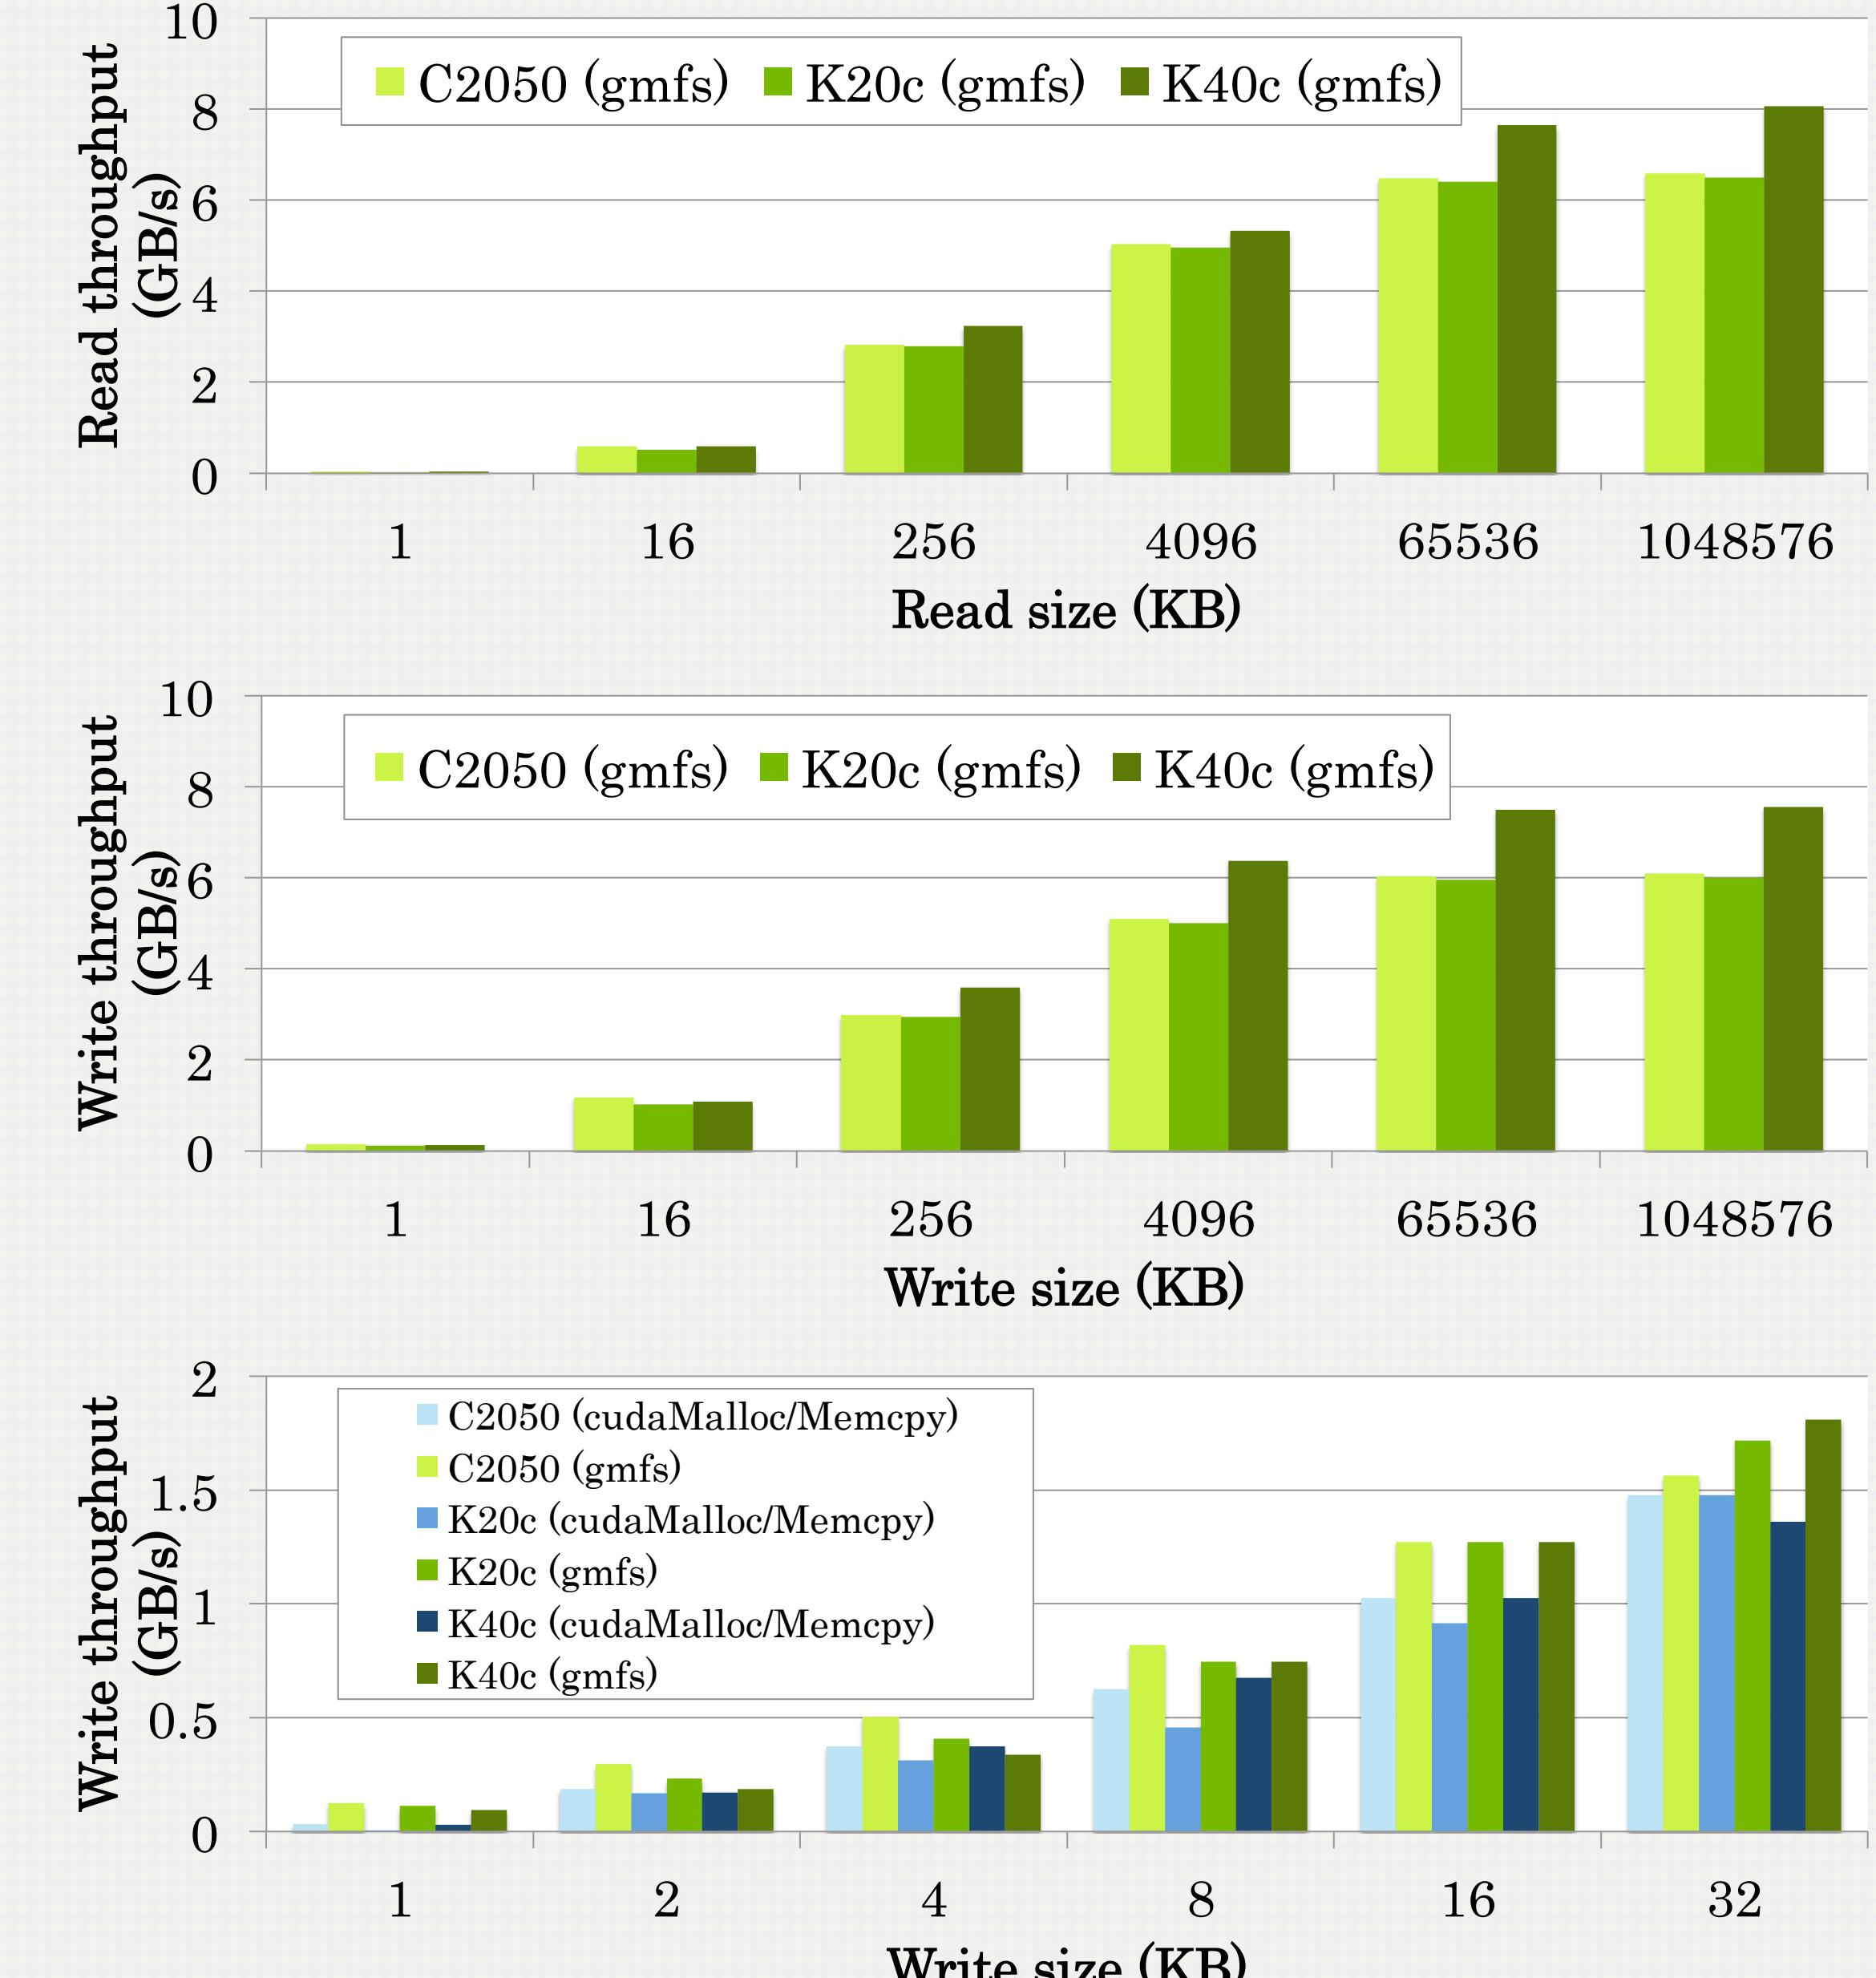
## Experiments

### Experimental environment

M/B	ASUSTeK P9X79 WS
CPU	Intel(R) Core(TM) i7-3930K CPU @ 3.20GHz
Memory	DDR3 DIMM 1333Hz (16GB)
GPU	NVIDIA GPU (Tesla: C2050, K20c, K40c)
CUDA	5.5
Persistent storage	OCZ VTX3-25SAT3-120G (SSD MLC) (Max Read: up to 550MB/s, Max Write: 500MB/s )

## Results

### Sequential Read/Write throughput comparisons



## Conclusion

- gmfs can accelerate sequential read/write, and utilize 82% of PCIe-gen2 peak bw, 50% of PCIe-gen3 peak bw
- gmfs can improve small write operations