

APIs, Architecture and Modeling for Extreme Scale Resilience

Dagstuhl Seminar: Resilience in Exascale Computing

9/30/2014

Kento Sato



LLNL-PRES-661421

This work was performed under the auspices of the U.S. Department of Energy by Lawrence Livermore National Laboratory under Contract DE-AC52-07NA27344. Lawrence Livermore National Security, LLC




Failures on HPC systems

- System resilience is critical for future extreme-scale computing
- 191 failures out of 5-million node-hours
 - A production application using Laser-plasma interaction code (pF3D)
 - Hera, Atlas and Coastal clusters @LLNL => MTBF: 1.2 day
 - C.f.) TSUBAME2.0 => MTBF: a day
- In extreme scale, failure rate will increase
- Now, HPC systems must consider failures as usual events


Motivation to resilience APIs

- Current MPI implementation does not have the capabilities
 - Standard MPI employs a fail-stop model
 - When a failure occurs ...
 - MPI terminates all processes
 - The user locate, replace failed nodes with spare nodes
 - Re-initialize MPI
 - Restore the last checkpoint
 - Applications will use more time for recovery
 - Users manually locate and replace the failed nodes with spare nodes via machinefile
 - The manual recovery operations may introduce extra overhead and human errors
- ⇒ APIs to handle the failures are critical




Resilience APIs, Architecture and the model

- Resilience APIs
 - ⇒ Fault tolerant messaging interface (FMI)



FMI: Fault Tolerant Messaging Interface [IPDPS2014]

FMI overview



- FMI is a survivable messaging interface providing MPI-like interface
 - Scalable failure detection \Rightarrow Overlay network
 - Dynamic node allocation \Rightarrow FMI ranks are virtualized
 - Fast checkpoint/restart \Rightarrow In-memory diskless checkpoint/restart


How FMI applications work ?

FMI example code

```
int main (int *argc, char *argv[]) {  
    FMI_Init(&argc, &argv);  
    FMI_Comm_rank(FMI_COMM_WORLD, &rank);  
    /* Application's initialization */  
    while ((n = FMI_Loop(...)) < numloop) {  
        /* Application's program */  
    }  
    /* Application's finalization */  
    FMI_Finalize();  
}
```

- FMI_Loop enables transparent recovery and roll-back on a failure
 - Periodically write a checkpoint
 - Restore the last checkpoint on a failure
- Processes are launched via fmirun
 - fmirun spawns fmirun.task on each node
 - fmirun.task calls fork/exec a user program
 - fmirun broadcasts connection information (endpoints) for FMI_init(...)


Launch FMI processes



User perspective: No failures

FMI example code

```
int main (int *argc, char *argv[]) {  
    FMI_Init(&argc, &argv);  
    FMI_Comm_rank(FMI_COMM_WORLD, &rank);  
    /* Application's initialization */  
    while ((n = FMI_Loop(...)) < 4) {  
        /* Application's program */  
    }  
    /* Application's finalization */  
    FMI_Finalize();  
}
```




- User perspective when no failures happens
- Iterations: 4
- Checkpoint frequency: Every 2 iterations
- FMI_Loop returns incremented iteration id

User perspective : Failure

FMI example code

```
int main (int *argc, char *argv[]) {  
    FMI_Init(&argc, &argv);  
    FMI_Comm_rank(FMI_COMM_WORLD, &rank);  
    /* Application's initialization */  
    while ((n = FMI_Loop(...)) < 4) {  
        /* Application's program */  
    }  
    /* Application's finalization */  
    FMI_Finalize();  
}
```

- Transparently migrate FMI rank 0 & 1 to a spare node
- Restart form the last checkpoint – 2th checkpoint at iteration 2
- With FMI, applications still use the same series of ranks even after failures



Resilience API: FMI_Loop

```
int FMI_Loop(void **ckpt, size_t *sizes, int len)
```

`ckpt` : Array of pointers to variables containing data that needs to be checkpointed


`sizes` : Array of sizes of each checkpointed variables

`len` : Length of arrays, `ckpt` and `sizes`

returns iteration id


- FMI constructs in-memory RAID-5 across compute nodes
- Checkpoint group size
 - e.g.) group_size = 4

FMI checkpointing



Application runtime with failures

- Benchmark: Poisson's equation solver using Jacobi iteration method
 - Stencil application benchmark
 - MPI_Isend, MPI_Irecv, MPI_Wait and MPI_Allreduce within a single iteration
- For MPI, we use the SCR library for checkpointing
 - Since MPI is not survivable messaging interface, we write checkpoint memory on tmpfs
- Checkpoint interval is optimized by Vaidya's model for FMI and MPI




P2P communication performance

	1-byte Latency	Bandwidth (8MB)
MPI	3.555 usec	3.227 GB/s
FMI	3.573 usec	3.211 GB/s

FMI directly writes checkpoints via memcpy, and can exploit the bandwidth

Even with the high failure rate, FMI incurs only a 28% overhead

Asynchronous multi-level checkpointing (MLC) [SC12]



Source: K. Sato, N. Maruyama, K. Mohror, A. Moody, T. Gamblin, B. R. de Supinski, and S. Matsuoka, "Design and Modeling of a Non-Blocking Checkpointing System," in Proceedings of the International Conference on High Performance Computing, Networking, Storage and Analysis, ser. SC '12. Salt Lake City, Utah: IEEE Computer Society Press, 2012

- Asynchronous MLC is a technique for achieving high reliability while reducing checkpointing overhead
- Asynchronous MLC Use storage levels hierarchically
 - RAID-5 checkpoint: Frequent for one node or a few node failure
 - PFS checkpoint: Less frequent and asynchronous for multi-node failure
- Our previous work model the asynchronous MLC

Failure analysis on Coastal cluster


	MTBF	Failure rate
L1 failure	130 hours	2.13^{-6}
L2 failure	650 hours	4.27^{-7}

Source: A. Moody, G. Bronevetsky, K. Mohror, and B. R. de Supinski, "Design, Modeling, and Evaluation of a Scalable Multi-level Checkpointing System," in Proceedings of the 2010 ACM/IEEE International Conference for High Performance Computing, Networking, Storage and Analysis (SC 10).


Simulation based on Asynchronous MLC

- Checkpoint size: 1 and 10 GB/node
- We increase L1 & L2 failure rates

Async. MLC (Multi-level C/R) model



$$\begin{aligned} t &: \text{Interval} & p_0(T) &= e^{-\lambda T} \\ c_c &: c\text{-level checkpoint time} & t_0(T) &= T \\ r_c &: c\text{-level recovery time} & p_i(T) &= \frac{\lambda_i}{\lambda} (1 - e^{-\lambda T}) \\ \lambda_i &: i\text{-level checkpoint time} & t_i(T) &= \frac{1 - (\lambda T + 1) \cdot e^{-\lambda T}}{\lambda \cdot (1 - e^{-\lambda T})} \end{aligned}$$




High efficiency with current failure rate

If both L1 & L2 failure rate increase, and checkpoint size is large, efficiency decrease faster


Resilience APIs, Architecture and the model

- Resilience APIs
 - In near future, applications must have capabilities of handling failures as usual events
⇒ Fault tolerant messaging interface (FMI)
- Resilience architecture and model
 - Software level approaches are not enough
⇒ Architecture using *Burst buffer*




Burst buffer storage architecture

- Burst buffer
 - A new tier in storage hierarchies
 - Absorb bursty I/O requests from applications
 - Fill performance gap between node-local storage and PFSs in both latency and bandwidth
- If you write checkpoints to burst buffers,
 - Faster checkpoint/restart time than PFS
 - More reliable than storing on compute nodes



Burst buffer storage architecture (cont'd)


Challenges for using burst buffer system



- Exploiting storage bandwidth of burst buffers
 - Burst buffers are connected to networks, networks can be bottleneck
- Analyzing reliability of systems with burst buffers
 - Adding burst buffer nodes increase total system size
 - System efficiency may decrease due to Increased overall failure by added burst buffers

APIs for burst buffers: InfiniBand-based I/O interface (IBIO)

- Provide POSIX-like I/O interfaces
 - Open, read, write and close operations
 - Client can open any files on any servers
 - `open("hostname:/path/to/file", mode)`
- IBIO use ibverbs for communication between clients and servers
 - Exploit network bandwidth of infiniBand



Resilience modeling overview


- To find out the best checkpoint/restart strategy for systems with burst buffers, we model checkpointing strategies

C/R strategy model

$$O_i = \begin{cases} C_i + E_i & (\text{Sync.}) \\ I_i & (\text{Async.}) \end{cases} \quad L_i = C_i + E_i$$


$$C_i \text{ or } R_i = \frac{\langle \text{C/R date size / node} \rangle \times \langle \# \text{ of C/R nodes per } S_i^* \rangle}{\langle \text{write perf. (} w_i \text{) } \rangle \text{ or } \langle \text{read perf. (} r_i \text{) } \rangle}$$

Recursive structured storage model



Storage Model: $H_N \{m_1, m_2, \dots, m_N\}$

Async. MLC model [2]



$t : \text{Interval}$	$p_0(T) = e^{-\lambda T}$
$c_c : c\text{-level checkpoint time}$	$t_0(T) = T$
$r_c : c\text{-level recovery time}$	$p_i(T) = \frac{\lambda_i}{\lambda} (1 - e^{-\lambda T})$
$\lambda_i : i\text{-level checkpoint time}$	$t_i(T) = \frac{1 - (\lambda T + 1) \cdot e^{-\lambda T}}{\lambda \cdot (1 - e^{-\lambda T})}$




Efficiency

Fraction of time an application
spends only
in useful computation

Sequential IBIO read/write performance

- Set chunk size to 64MB for both IBIO and NFS to maximize the throughputs




Node specification	
CPU	Intel Core i7-3770K CPU (3.50GHz x 4 cores)
Memory	Cetus DDR3-1600 (16GB)
M/B	GIGABYTE GA-Z77X-UD5H
SSD	Crucial m4 msata 256GB CT256M4SSD3 (Peak read: 500MB/s, Peak write: 260MB/s)
SATA converter	KOUTECH IO-ASS110 mSATA to 2.5' SATA Device Converter with Metal Fram
RAID Card	Adaptec RAID 7805Q ASR-7805Q Single

Interconnect :Mellanox FDR HCA (Model No.: MCX354A-FCBT)

IBIO achieve the same remote read/write performance as the local read/write performance by using RDMA

Efficiency with Increasing Failure Rates and Checkpoint Costs

- Assuming there is no message logging overhead



In days or a day of MTBF, there is no big efficiency differences

In a few hours of MTBF, with burst buffers, systems can still achieve high efficiency

Even in a hour of MTBF, with uncoordinated, systems can still achieve 70% efficiency

⇒ Partial restart can decrease recovery time from burst buffers and PFS checkpoint

Allowable Message Logging overhead

Message logging overhead allowed in uncoordinated checkpointing to achieve a higher efficiency than coordinated checkpointing

Flat buffer		Burst buffer	
scale factor	Allowable message logging overhead	scale factor	Allowable message logging overhead
1	0.0232%	Coordinated	0.00435%
2	0.0929%		0.0175%
10	2.45%		0.468%
50	84.5%	Uncoordinated	42.0%
100	≈ 100		99.9%


- Logging overhead must be relatively small, less than a few percent in days or a day of MTBF
 - In a few hours or a hour, very high message logging overheads are tolerated

⇒ Uncoordinated checkpointing can be more effective on future systems

Effect of Improving Storage Performance


Flat Buffer-Coordinated
Burst Buffer-Coordinated

Flat Buffer-Uncoordinated
Burst Buffer-Uncoordinated



To see which storage impact to efficiency, we increase performance of level-1 and level-2 storage while keeping MTBF a hour

Improvement of level-1 storage performance does not impact efficiency for both flat buffer and burst buffer systems



Increasing the performance of the PFS does impact system efficiency

L2 C/R overhead is a major cause of degrading efficiency, so reducing level-2 failure rate and improving level-2 C/R is critical on future systems

Summary: Towards extreme scale resiliency

- Resilient APIs
 - Resilient APIs in MPI is critical for fast and transparent recovery in HPC applications
 - In-memory C/R by FMI incurs only a 28% overhead even with the high failure rate
 - Software-level solution may not enough at extreme scale
- Resilient Architecture
 - Burst buffers are beneficial for C/R at extreme scale
 - Uncoordinated C/R
 - When MTBF is days or a day, uncoordinated C/R may not be effective
 - If MTBF is a few hours or less, will be effective
 - Level-2 failure, and Level-2(PFS) performance
 - Reducing Level-2 failure, increasing Level-2 (PFS) performance are critical to improve overall system efficiency

Q & A

Speaker

Kento Sato

Lawrence Livermore National Laboratory
kento@llnl.gov

External collaborators

Satoshi Matsuoka, Tokyo Tech
Naoya Maruyama, RIKEN AICS





**Lawrence Livermore
National Laboratory**