

# *Extreme-Scale Resilience for Billion-Way of Parallelism*

SC14 Workshop: ATIP Workshop on Japanese Research Toward Next-Generation Extreme Computing  
11/17/2014



Kento Sato  
Lawrence Livermore National Laboratory

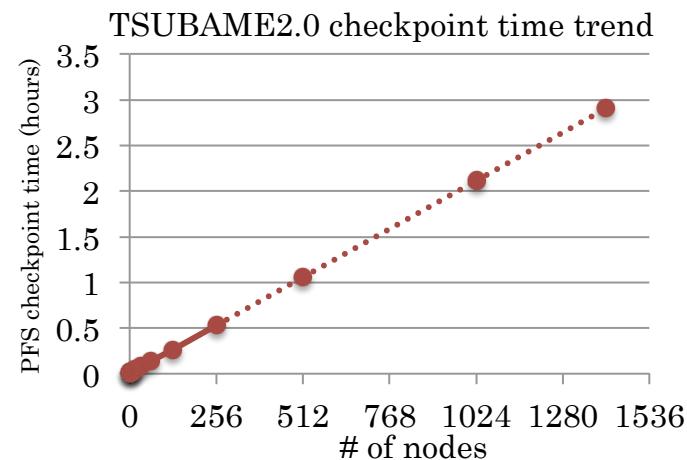


LLNL-PRES-664262

This work was performed under the auspices of the U.S. Department of Energy by Lawrence Livermore National Laboratory under Contract DE-AC52-07NA27344. Lawrence Livermore National Security, LLC

# Failures on HPC systems

- System resiliency is critical for future extreme-scale computing
- MTBF of supercomputers
  - LLNL (Hera, Atlas & Coastal): 1.2 days<sup>[1]</sup>
  - Blue Waters: 8-12 hours<sup>[2]</sup>
  - Titan: 8-12 hours (<= a few failures/day<sup>[2]</sup>)
- MTBF is shrinking
  - MTBF is projected to shrink to a few hours
- Checkpoint/Restart is a popular way for fault tolerance
- Simple checkpoint/restart may not work at extreme scale



[1] A. Moody, G. Bronevetsky, K. Mohror, and B. R. de Supinski, "Design, Modeling, and Evaluation of a Scalable Multi-level Checkpointing System," in Proceedings of the 2010 ACM/IEEE International Conference for High Performance Computing, Networking, Storage and Analysis, ser. SC '10. Washington, DC, USA: IEEE Computer Society, Nov. 2010, pp. 1–11. [Online]. Available: <http://dx.doi.org/10.1109/SC.2010.18>

[2] Yves Robert, "Fault\_Tolerance Techniques for Computing at Scale", Keynote Talk, CCGrid2014

[3] Kento Sato, Adam Moody, Kathryn Mohror, Todd Gamblin, Bronis R. de Supinski, Naoya Maruyama and Satoshi Matsuoka, "Design and Modeling of a Non-blocking Checkpointing System", SC12

# Tokyo Tech. Billion-Way Resilience Project (2011-2015)

- PI: Satoshi Matsuoka
- Current collaborations:
  - ANL (Franck Cappello, FTI), LLNL (Bronis de Spinksi, SCR), ETH Zurich (Torsten Hoefer), RIKEN (Naoya Maruyama), U-Tokyo (Hideyuki Jitsumoto) ...
- Objective: Scalable fault tolerance techniques for extreme scale system
  - API & Software: Encoding/Redundancy technique, Compression, Support for Many-core architecture
  - Architecture: Scalable Storage design, and Resilient network/interconnects
  - Analysis: Failure analysis, and Failure prediction
  - Modeling: Optimal checkpoint interval, Encoding/Redundancy, and I/O model

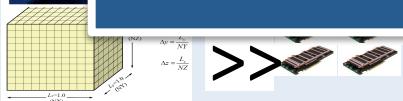


Lawrence Livermore  
National Laboratory

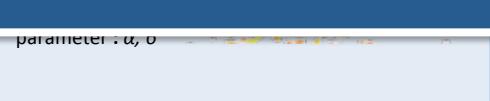


# Tokyo Tech. Billion-Way Resilience Project (2011-2015)

## High Performance Computing Applications



>>>



### API & software

- Encoding/Redundancy technique
- Checkpoint compression
- Asynchronous I/O APIs for checkpointing
- Support for Many-core architecture
- Resource manager & Scheduler for resilience

### Architecture

- Resilient storage design
- Resilient network design

### Analysis

- Failure monitoring & analysis

### Model

- Models of checkpointing / I/O for optimal interval & performance prediction

# Failure monitoring & analysis (2010 ~ present)

Analysis

## Failure history of TSUBAME2.0/2.5

### Findings

- Failures seasonal
  - Largely due to boot failures in peak-shift operations during summer to limit power, despite SW retries
  - Future SCs in Clouds need to cope with this
- GPU vs. CPUs
  - 19 CPU+memory fail-stop failures, 25 replacements, MTBF 118 years,  $2.22^{18}$  FLOP/error
  - 53 GPU+memory ECC fail-stop failures, 57 replacements, MTBF 75 years,  $1.61^{19}$  FLOP/error
  - GPU error rate x7 better / flop vs. CPU, proportional to performance difference per chip
- Failures are Largely Independent
  - Most of failures are a single node
  - Low # of InfiniBand and storage failures

TSUBAME2.5 故障履歴[Failure History of TSUBAME2.5]							
csv形式で出力							
1.障害発生情報[Current Trouble Information of TSUBAME2.5]							
ホスト/機器	キューブ	発生日付	復旧日付	障害状況	原因	対処	影響範囲
t2a000071	S	2014/11/13 06:29		Uncorrectable PCI Express Error (Embedded device, Bus 0, Device 0, Function 0, Error status 0x00000000)			
t2a002004	GV	2014/11/09 12:00-13:00	2014/11/13 11:45	GPU ECC Error SingleBit:0 DoubleBit:0 ECC Error	エラーカウントクリア後、ス トレステストにて問題無し	該当ノード	GPU
t2a000077	S	2014/11/09 11:45	2014/11/13 11:45	Uncorrectable PCI Express Error (Embedded device, Bus 0, Device 0, Function 0, Error status 0x00000000)	Uncorrectable PCI Express Error	GPU1-2swap後、2時間のス トレステストにて問題無し	該当ノード
t2a000177	S			GPU2 ECC Error SingleBit:0 DoubleBit:0 ECC Error	エラーカウントクリア後、ス トレステストにて問題無し	該当ノード	GPU
t2a000056	S			Express Error	Express Error	GPU1-2swap後、2時間のス トレステストにて問題無し	該当ノード
t2a007053						再起動	該当OS書き込みのジ ャンプ
t2a000057							該当ノード
t2a007057		2014/11/09 23:15	05:12				
t2a007059							
t2a000179	S	2014/11/07 16:24	2014/11/13 11:45	Uncorrectable PCI Express Error (Embedded device, Bus 0, Device 0, Function 0, Error status 0x00000000)			
t2a000072							
t2a000073	S	2014/11/07 15:29-15:40	2014/11/13 11:45	Uncorrectable PCI Express Error (Embedded device, Bus 0, Device 0, Function 0, Error status 0x00000000)	Uncorrectable PCI Express Error		
t2a000076							
t2a001084	S	2014/11/07 02:00-03:00	2014/11/12 11:30	/Scratch : Read-only file system	SSD障害	SSD bay交換	
SFA10K sc10		2014/11/06 05:45	2014/11/07 11:30	ディスクFAIL : Enc#3 Slot#17 (Idx: 614) -	ディスク交換	なし	
t2a005131	HX	2014/11/04 22:41	2014/11/13 11:45	Uncorrectable PCI Express Error (Embedded device, Bus 0, Device 0, Function 0, Error status 0x00000000)	I/O Mezz障害	GPU1-2 Swap後、ストレス テストに失敗。GPUPowerボ ード、I/O Mezz交換後、24 時間のス トレステストにて問題無し。	該当ノード
t2a002075	GV	2014/11/04 16:58	2014/11/11 10:00	Uncorrectable PCI Express Error (Embedded device, Bus 0, Device 0, Function 0, Error status 0x00000000)	Uncorrectable PCI Express Error	GPU1-2swap後、2時間のス トレステストにて問題無し。	該当ノード
t2a006176	I	2014/11/04 12:00-13:00	2014/11/07 15:15	GPU LinkSpeed Down - 1 Errors of 3 x16	GPU LinkSpeedDown	GPU1-2swap後、2時間のス トレステストにて問題無し。	該当ノード
t2a006179				GPU			
t2a004070	GV	2014/11/04 12:00-13:00	2014/11/07 11:45	GPU LinkSpeed Down - 1 Errors of 3 x16	GPU LinkSpeedDown	GPU1-2swap後、2時間のス トレステストにて問題無し。	該当ノード
SFA10K sc10		2014/11/04 10:52	2014/11/05 17:30	コントローラ#1から、ディスクエンクロージャー	DEM障害	交換作業中、最大性能 が低下	Disk
		2014/11/04	2014/11/04				

We're are Disclosing the failure history in public

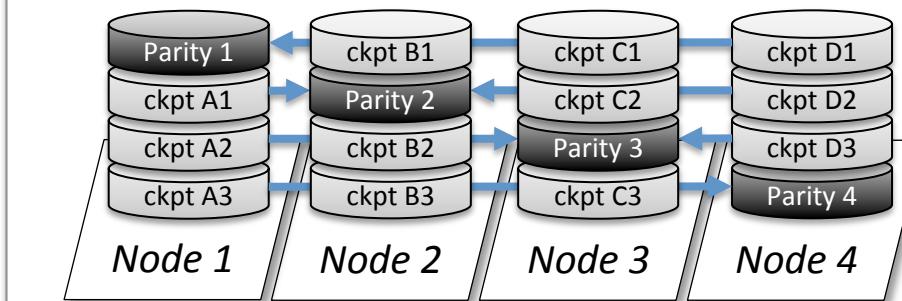
# FTI: High Performance Fault Tolerance Interface

[SC11, EuroPar12 & Cluster12 (Leonardo Bautista-Gomez et al.)]

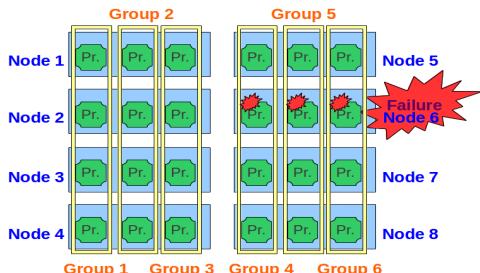
- Diskless checkpoint:

- Create redundant data across local storages on compute nodes using a encoding technique such as Reed-Solomon, XOR
  - Scalable by using distributed disks
- Can restore lost checkpoints on a failure caused by small # of nodes like RAID-5

Diskless checkpointing

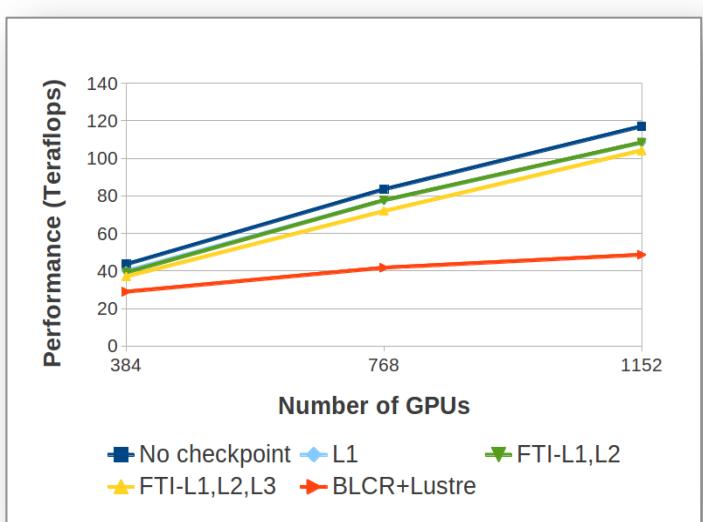


## Diskless checkpoint runtime library using Reed-Solomon encoding



➤ FTI implements a scalable Reed-Solomon encoding algorithm by utilizing local storages such as SSD

➤ FTI analyzes the topology of the system and create encoding clusters that increase the resilience



# Design and Modeling of Async. Checkpointing

[SC12, Kento Sato et al.]

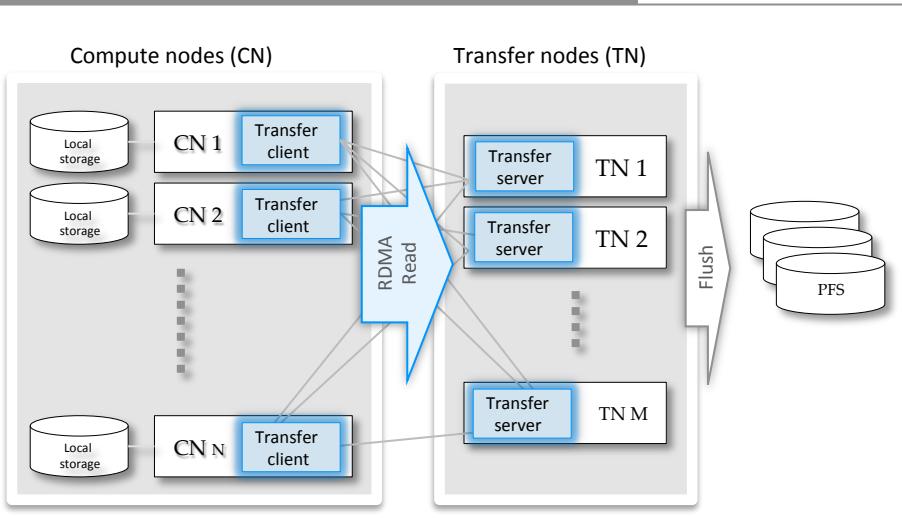
8%

- **Objective:** Minimize checkpoint overhead to PFS
  - Minimize CPU usage, memory and network bandwidth
- **Proposed method:** Implementation and modeling  
Non-blocking checkpointing
  - Asynchronously write checkpoints to PFS through Staging nodes using RDMA
  - Determine the optimal checkpoint interval on the asynchronous checkpoint scheme

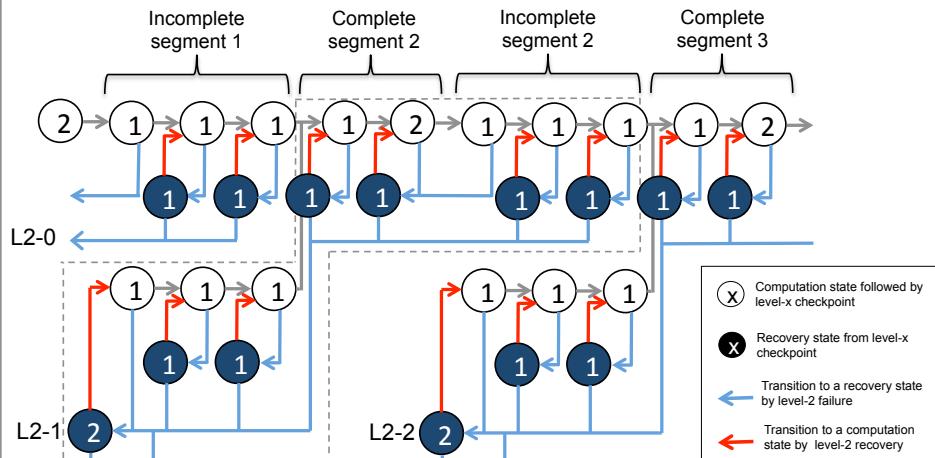
8-12% of failures still requires PFS checkpoint

Failure analysis on TSUBAME2.0

## Async. checkpointing system



## Async. checkpointing model



# Design and Modeling of Async. Checkpointing

[SC12, Kento Sato et al.]

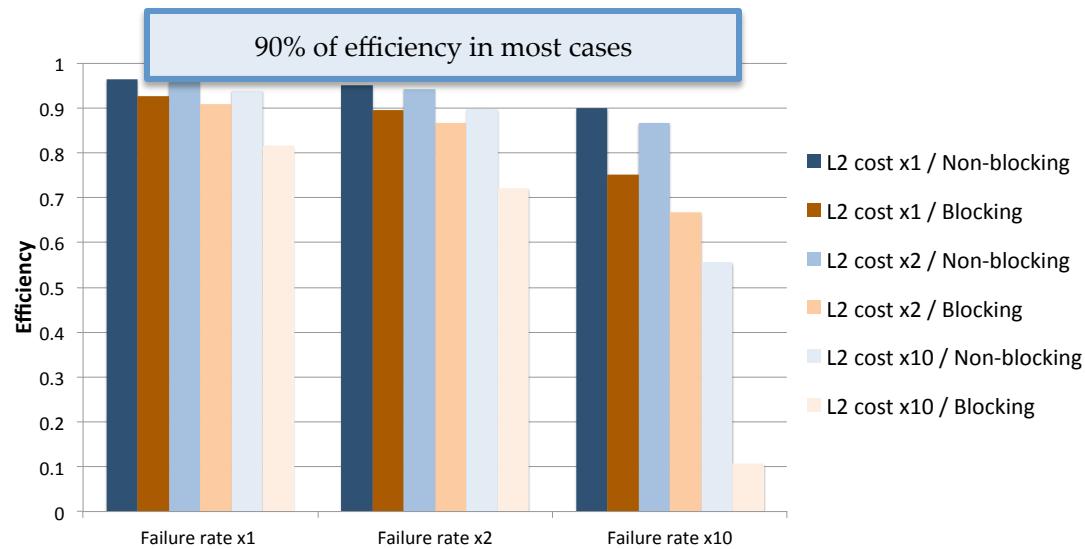
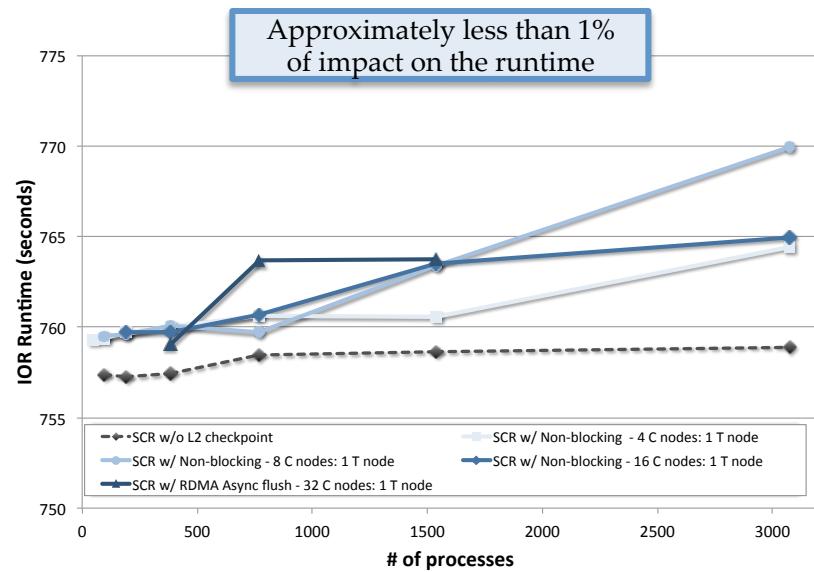
- Experiment:**

- Benchmark: CPU-bound micro benchmark
- Method
  - Non-blocking: proposed method on two level checkpointing
    - L1: XOR checkpoint, L2: Proposed **non-blocking** checkpoint
  - Blocking: Existing two level checkpointing
    - L1: XOR checkpoint, L2: **Blocking** checkpoint

$$\text{Efficiency} = \frac{\text{ideal\_time}}{\text{expected\_time}}$$

- Results:**

- Asynchronous RDMA checkpoint: **About 1 % of overhead** with the proposed checkpointing
- Optimal checkpoint interval: Achieved **high efficiency** even with increasing failure rates



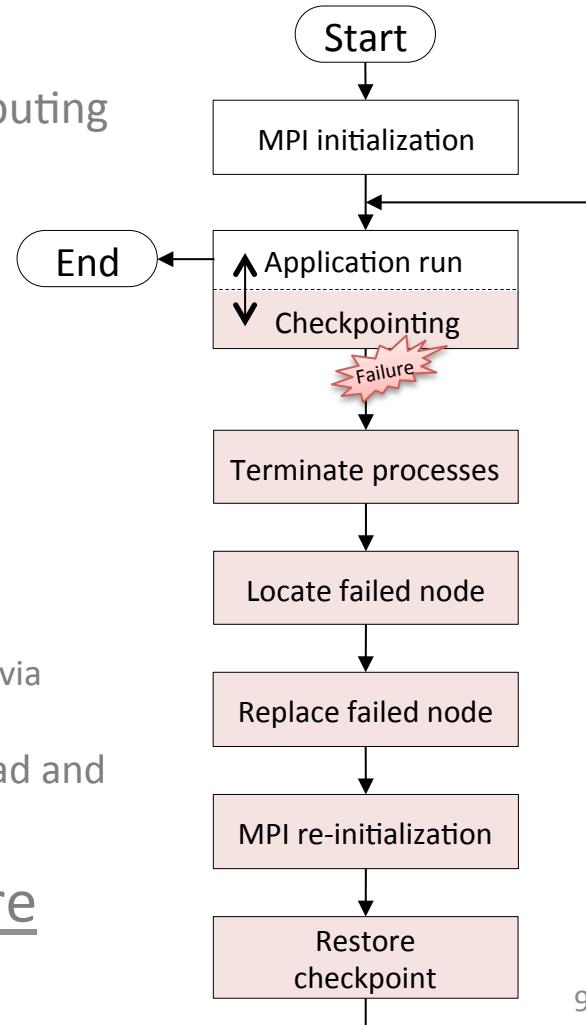
# FMI: Fault Tolerant Messaging Interface

[IPDSP2014, Kento Sato et al.]

## Requirement of fast and transparent recovery

- MPI
  - De-facto communication library enabling parallel computing
  - Standard MPI employs a fail-stop model
- When a failure occurs ...
  - MPI terminates all processes
  - The user locate, replace failed nodes with spare nodes
  - Re-initialize MPI
  - Restore the last checkpoint
- Applications will use more time for recovery
  - Users manually locate and replace the failed nodes with spare nodes via machinefile
  - The manual recovery operations may introduce extra overhead and human errors

⇒ APIs for transparent, but fast recovery are critical

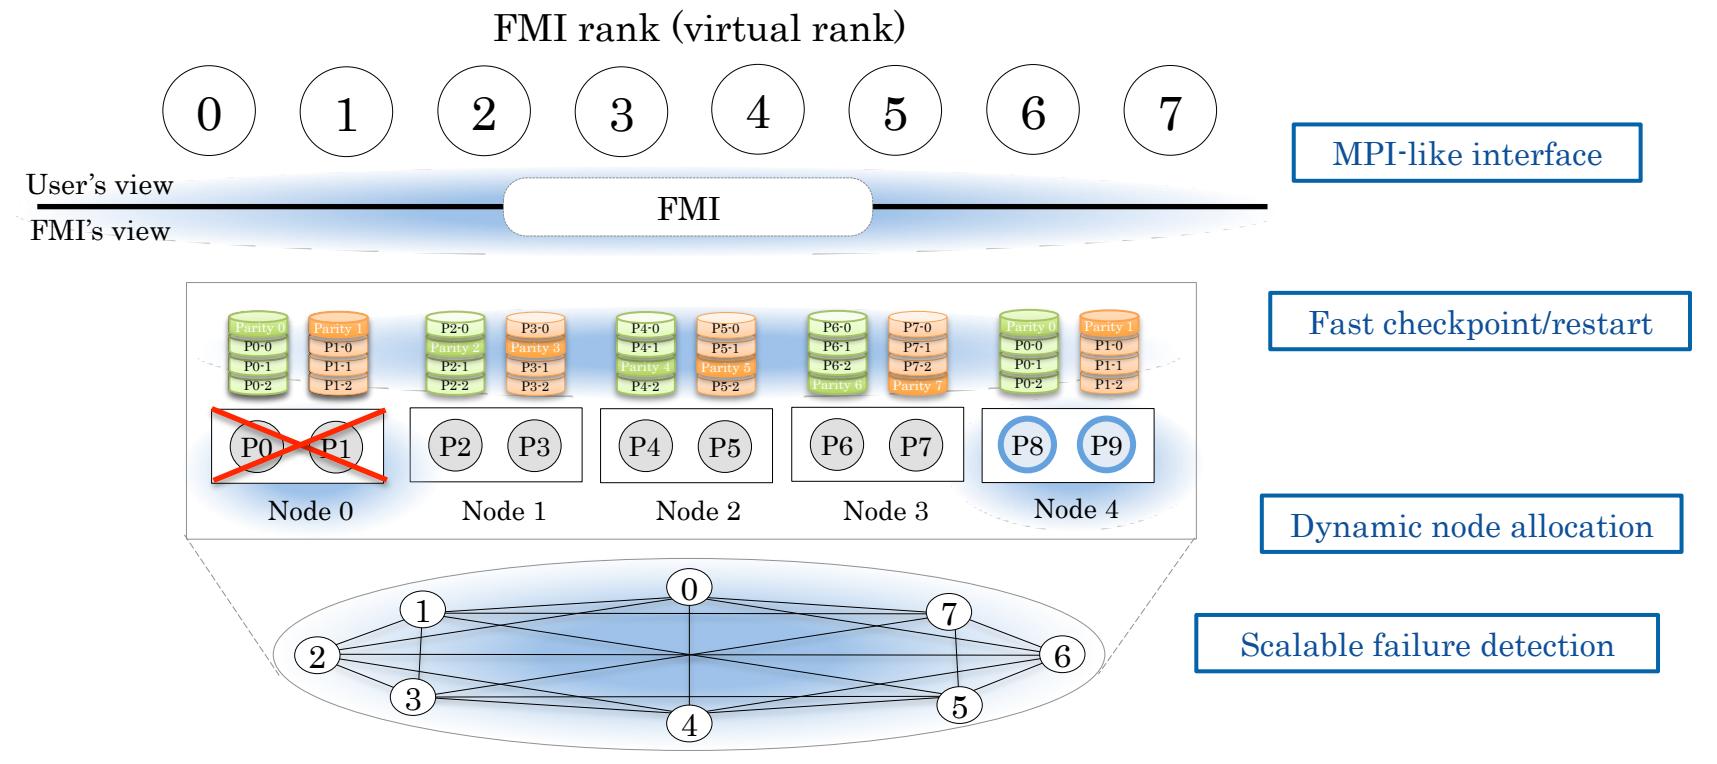


# FMI: Fault Tolerant Messaging Interface

[IPDSP2014, Kento Sato et al.]

## FMI for Fast and transparent recovery

### FMI overview



- FMI is a survivable messaging interface providing MPI-like interface
  - Scalable failure detection  $\Rightarrow$  Overlay network
  - Dynamic node allocation  $\Rightarrow$  FMI ranks are virtualized
  - Fast in-memory checkpoint/restart  $\Rightarrow$  Diskless checkpoint/restart

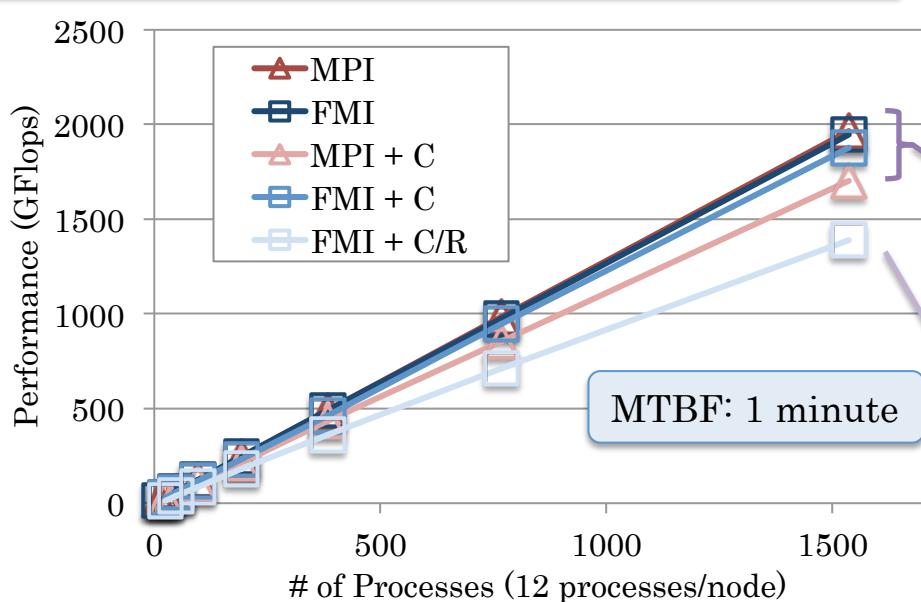
# FMI: Fault Tolerant Messaging Interface

[IPDSP2014, Kento Sato et al.]

## Example code & Evaluation

### FMI example code

```
int main (int *argc, char *argv[]) {
    FMI_Init(&argc, &argv);
    FMI_Comm_rank(FMI_COMM_WORLD, &rank);
    /* Application's initialization */
    while ((n = FMI_Loop(...)) < numloop) {
        /* Application's program */
    }
    /* Application's finalization */
    FMI_Finalize();
}
```



- **FMI\_Loop** enables transparent recovery and roll-back on a failure
  - Periodically write a checkpoint
  - Restore the last checkpoint on a failure

### P2P communication performance

	1-byte Latency	Bandwidth (8MB)
MPI	3.555 usec	3.227 GB/s
FMI	3.573 usec	3.211 GB/s

FMI directly writes checkpoints via memcpy, and can exploit the bandwidth

Even with the high failure rate, FMI incurs only a 28% overhead

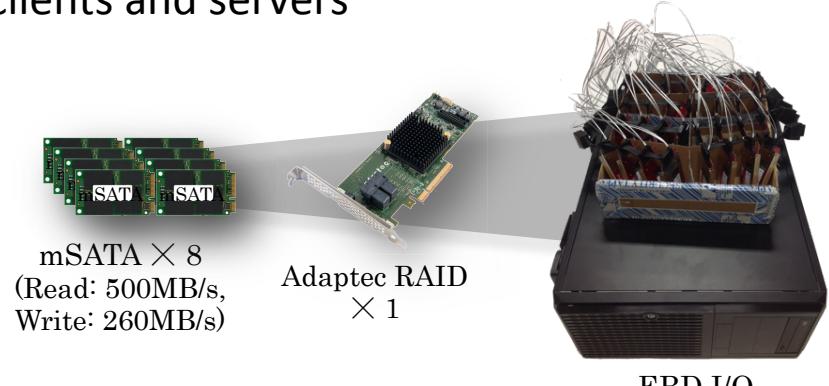
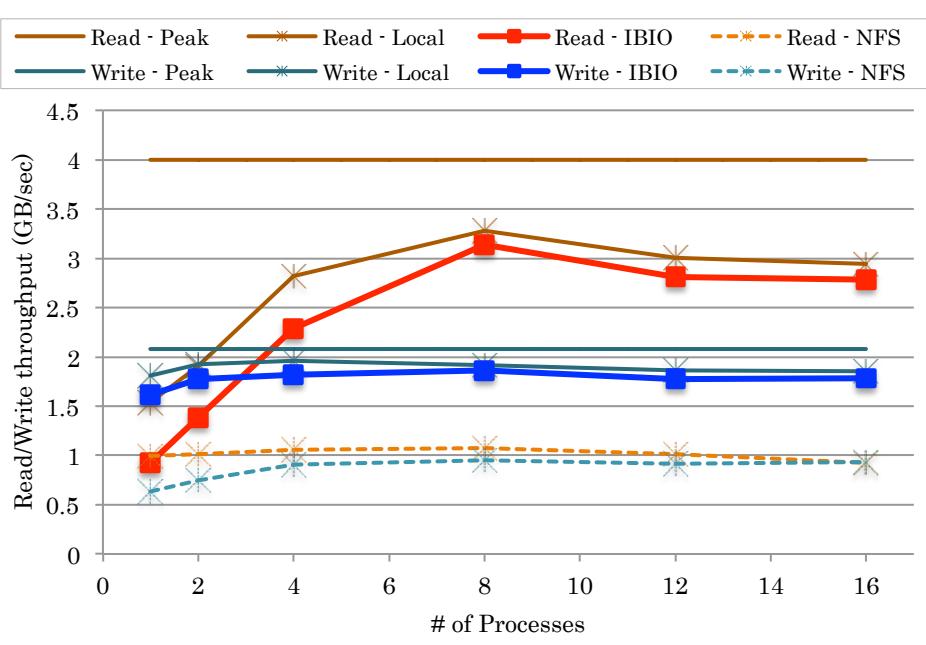
# Burst Buffers for Resilient Checkpoint/Restart

[CCGrid2014, Kento Sato et al.]

## TSUBAME3.0 EBD Prototype mSATA High I/O BW, low power & cost

- Provide POSIX-like I/O interfaces
  - open, read, write and close
  - Client can open any files on any servers
- IBIO use ibverbs for communication between clients and servers
  - Exploit network bandwidth of infiniBand

`open("hostname:/path/to/file", mode)`



### Node specification

CPU	Intel Core i7-3770K CPU (3.50GHz x 4 cores)
Memory	Cetus DDR3-1600 (16GB)
M/B	GIGABYTE GA-Z77X-UD5H
SSD	Crucial m4 msata 256GB CT256M4SSD3 (Peak read: 500MB/s, Peak write: 260MB/s)
SATA converter	KOUTECH IO-ASS110 mSATA to 2.5' SATA Device Converter with Metal Fram
RAID Card	Adaptec RAID 7805Q ASR-7805Q Single

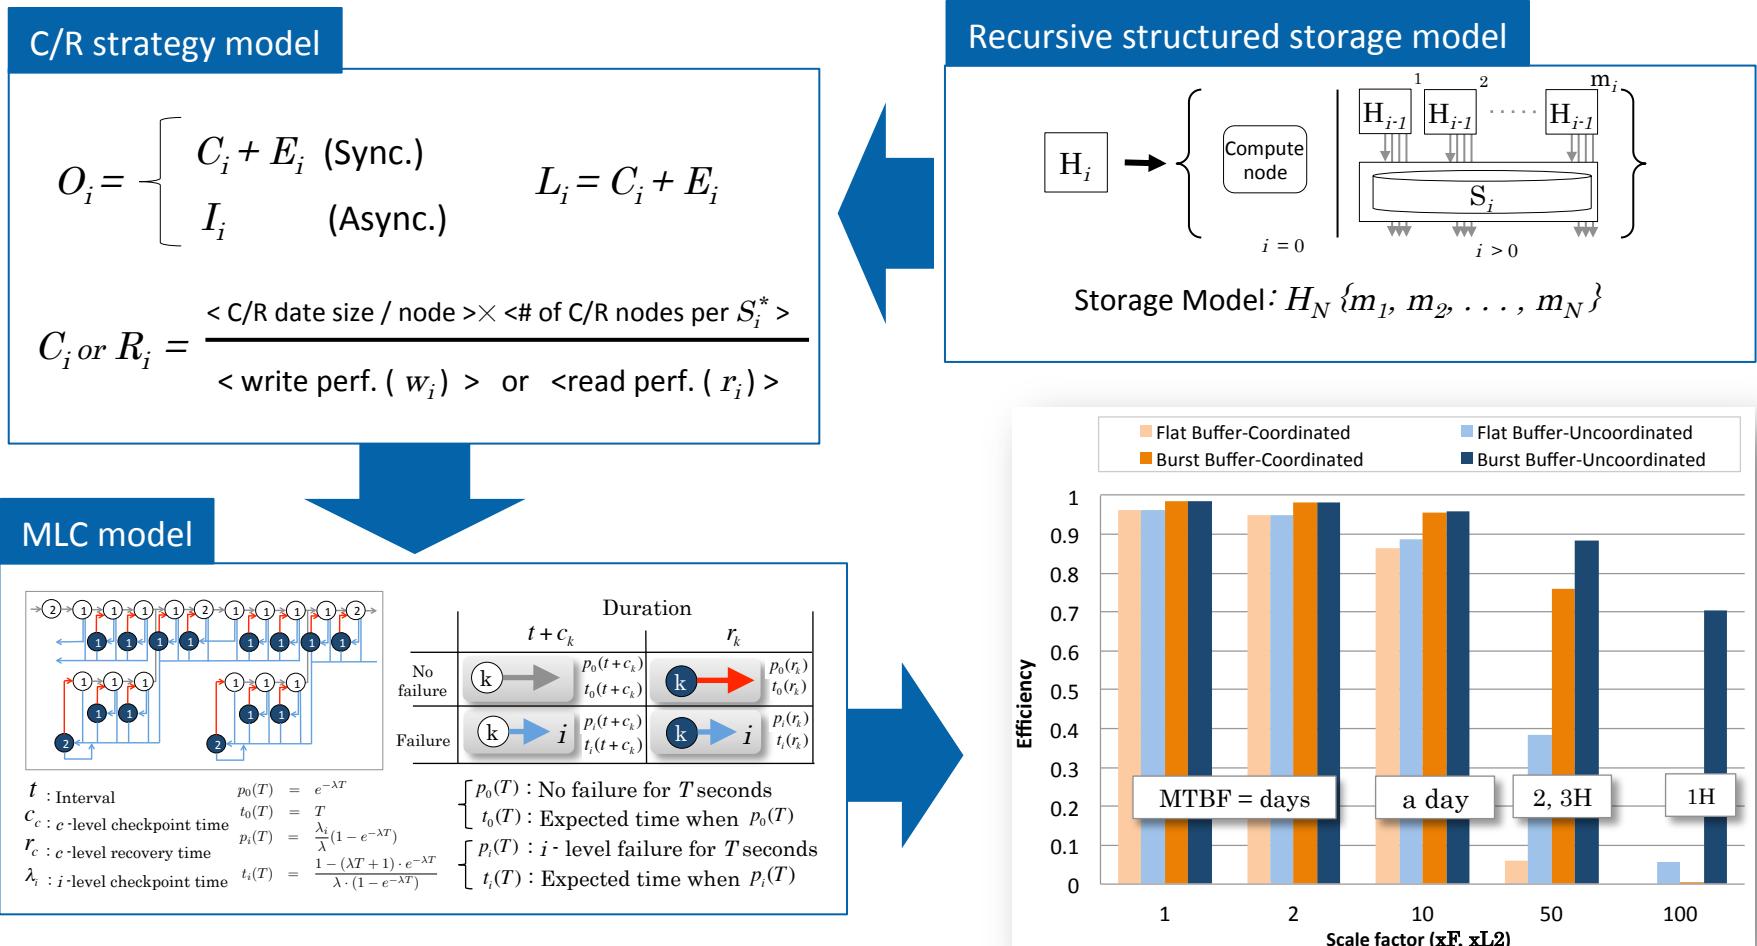
Interconnect :Mellanox FDR HCA (Model No.: MCX354A-FCBT)

# Burst Buffers for Resilient Checkpoint/Restart

[CCGrid2014, Kento Sato et al.]

## Resilience modeling overview

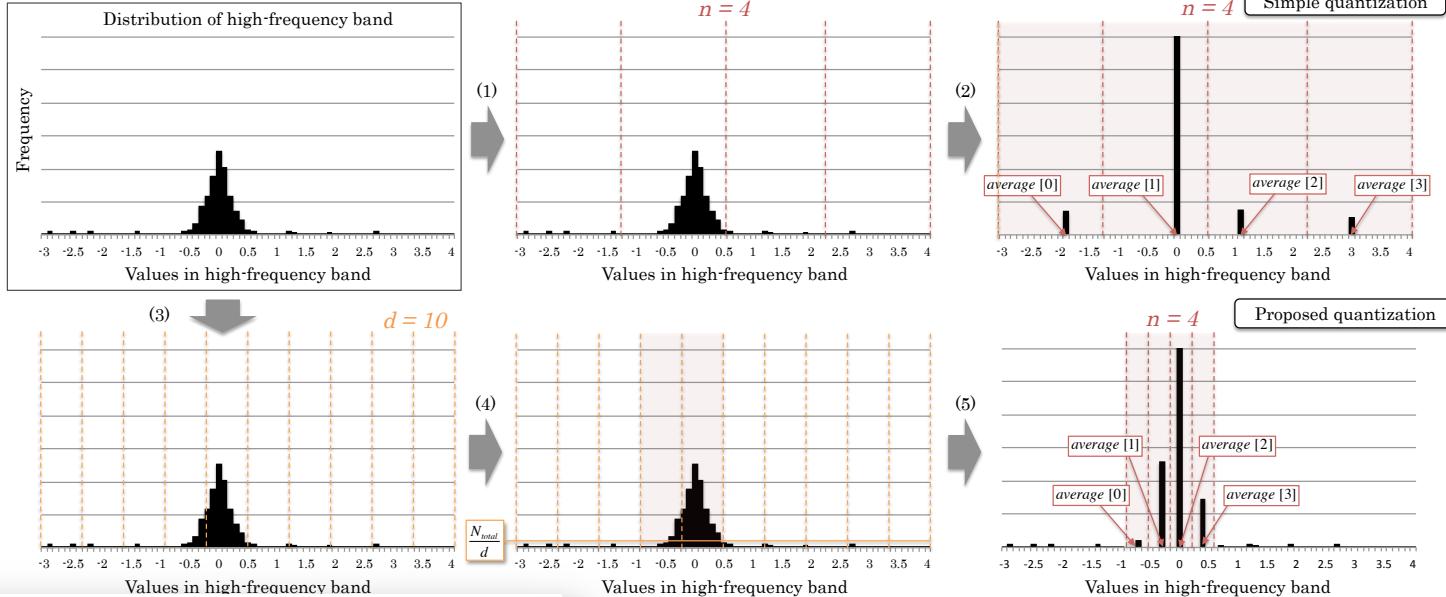
To find out the best checkpoint/restart strategy for systems with burst buffers, we model checkpointing strategies



# Lossy Floating-point compression

[Submitted to IPDPS2015, Naoto Sasaki et al.]

## Haar wavlet-based lossy compression

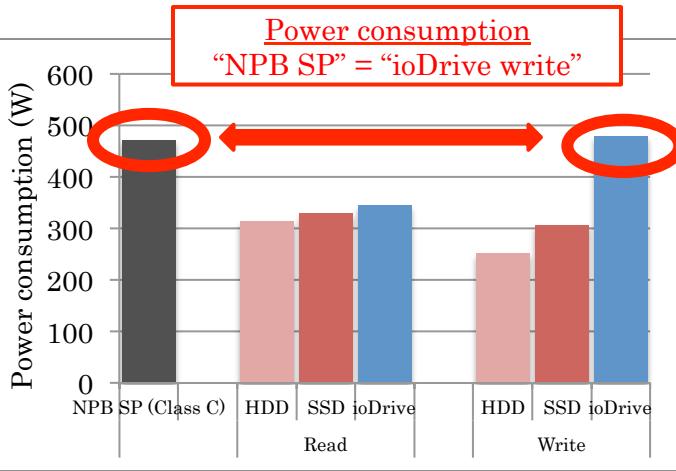


- Evaluation on a real application: NICAM
  - compression rate: 10%-15%
  - Errors: a few percent of errors
- More investigation is needed to deal with the errors

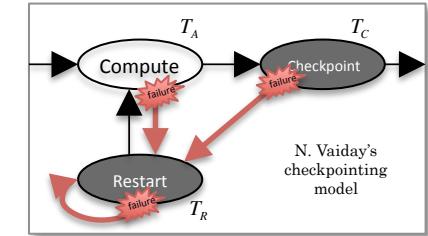
# NVM Energy-Away C/R Optimization

[FTXS2013, Takafumi Saito et al.]

## Energy optimization for C/R using DVFS



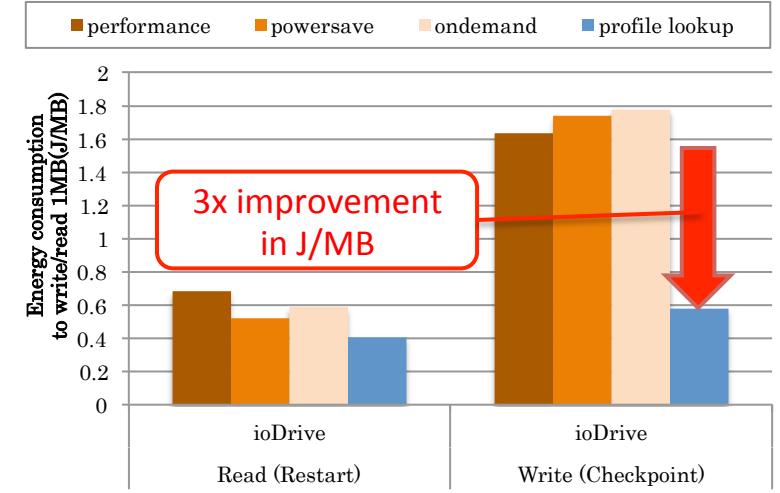
- ioDrive relies on CPU cores for
  - Grooming: a garbage collector that pre-erases unused blocks in background to accelerate future write operation
  - Wear leveling: a balanced write technique to extend the lifetime of a device
- When decreasing CPU frequency, I/O throughput of ioDrive is degraded like CPU



## Energy-Away C/R Model

	Expected {run   I/O} time	Power
Compute	$T_{\bar{A}} = \lambda^{-1} e^{\lambda(T_c+T_r)} (e^{\lambda T_a} - 1)$	$W_A$
Checkpoint	$T_{\bar{C}} = \lambda^{-1} (e^{\lambda T_c} - 1)$	$W_C$
Restart	$T_{\bar{R}} = \lambda^{-1} (e^{\lambda T_c} - 1)(e^{\lambda T_r} - 1)$	$W_R$

$$J = T_{\bar{A}} \cdot W_A + T_{\bar{C}} \cdot W_C + T_{\bar{R}} \cdot W_R$$



# Fail-in-Place Network Design

[SC14, Jens Domke et al.]

## Failure Analysis of HPC Systems and Fail-in-Place Strategy

### LANL Cluster 2 (97–05)

- Unknown configuration

### Deimos (07–12)

- 728 nodes
- 108 IB switches
- ≈1,600 links

### TSUBAME2.0/2.5 (10–?)

- 1,555 nodes (1,408 compute nodes)
- ≈500 IB switches
- ≈7,000 links

Software more reliable

High MTTR

≈1% annual failure rate

Repair/maintenance is expensive!

TABLE I. COMPARISON OF NETWORK-RELATED HARDWARE AND SOFTWARE FAILURES, MTBF/MTTR, AND ANNUAL FAILURE RATES

Fault Type	Deimos*	LANL Cluster 2	TSUBAME2.5
Percentages of network-related failures			
Software	13%	8%	1%
Hardware	87%	46%	99%
Unspecified		46%	
Percentages for hardware only			
NIC/HCA	59%	78%	1%
Link	27%	7%	93%
Switch	14%	15%	6%
Mean time between failure / mean time to repair			
NIC/HCA	X <sup>†</sup> / 10 min	10.2 d / 36 min	X / 5–72 h
Link	X / 24–48 h	97.2 d / 57.6 min	X / 5–72 h
Switch	X / 24–48 h	41.8 d / 77.2 min	X / 5–72 h
Annual failure rate			
NIC/HCA	1%	X	≥ 1%
Link	0.2%	X	0.9% <sup>‡</sup>
Switch	1.5%	X	1%

\* Deimos' failure data is not publicly available

† Not enough data for accurate calculation

‡ Excludes first month, i.e., failures sorted out during acceptance testing

### Fail-in-Place Strategy

- Replace only *critical* failures, and disable *non-critical* failed components
- Common in storage systems
- Applied when maintenance costs exceed maintenance benefits
- Example: IBM's Flipstone (uses RAID arrays; software disables failed HDD and migrates data)

Can we do fail-in-place in HPC networks?

# Fail-in-Place Network Design

[SC14, Jens Domke et al.]

Architecture

## Simulating Network Failures and Throughput Degradation

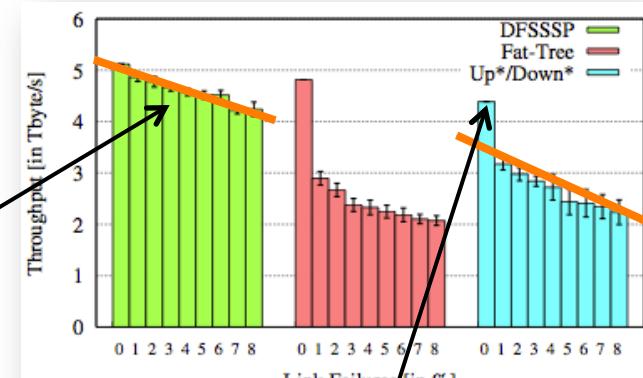
Routing is elementary component to enable fail-in-place networks

Tool-chain for checking fault tolerance of topology and routing algorithm

- Generate faulty topology based on artificial/real network topology
- Apply topology-[aware | agnostic] routing & check connectivity
- Flit-level simulation of InfiniBand hardware with uniform random injection or N-to-N exchange traffic

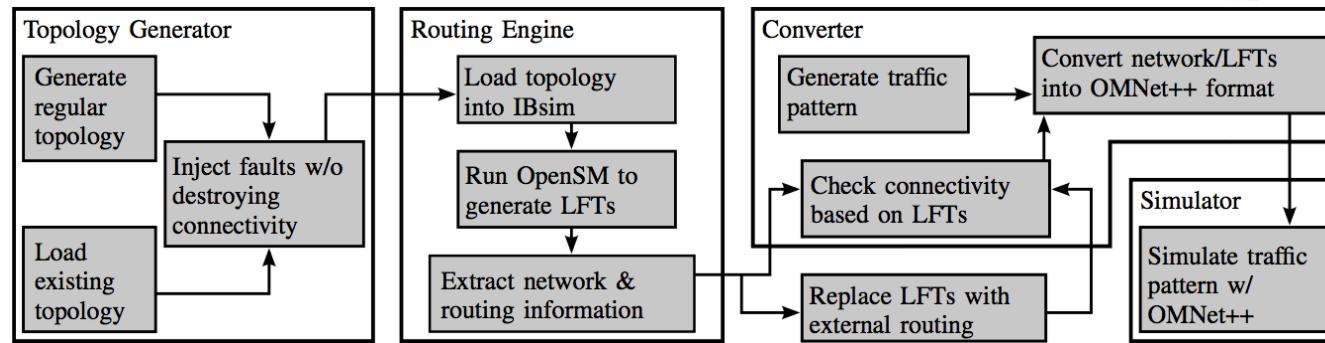
Simulated throughput degradation as a metric for network/routing reliability

- For each % of switch/link failures do multiple runs (diff. seeds)
- Calculate throughput
- Linear regression



(b) 14-ary 3-tree with 2,156 HCAs

Intercept



# Fail-in-Place Network Design

[SC14, Jens Domke et al.]

Architecture

## Implications for a real HPC System and Conclusions

All investigated routing algorithms show limitations

- Fat-tree, UpDown, DOR, Torus2QoS
- MinHop, SSSP, DFSSSP, LASH

Topology-aware routings

- High throughput decrease possible with small failure percentage
- Fail to route highly damaged netw.
- Routes not always DL-free (DOR)

Topology-agnostic routing algorithms

- Ignore deadlocks (MinHop, SSSP)
- Deadlock-avoidance via VLs can be impossible for large scale netw.

Changing from Up\*/Down\* (default) to DFSSSP routing on TSUBAME2.5 improves the throughput by **2.1x** for the fault-free network and increases TSUBAME's fail-in-place characteristics.

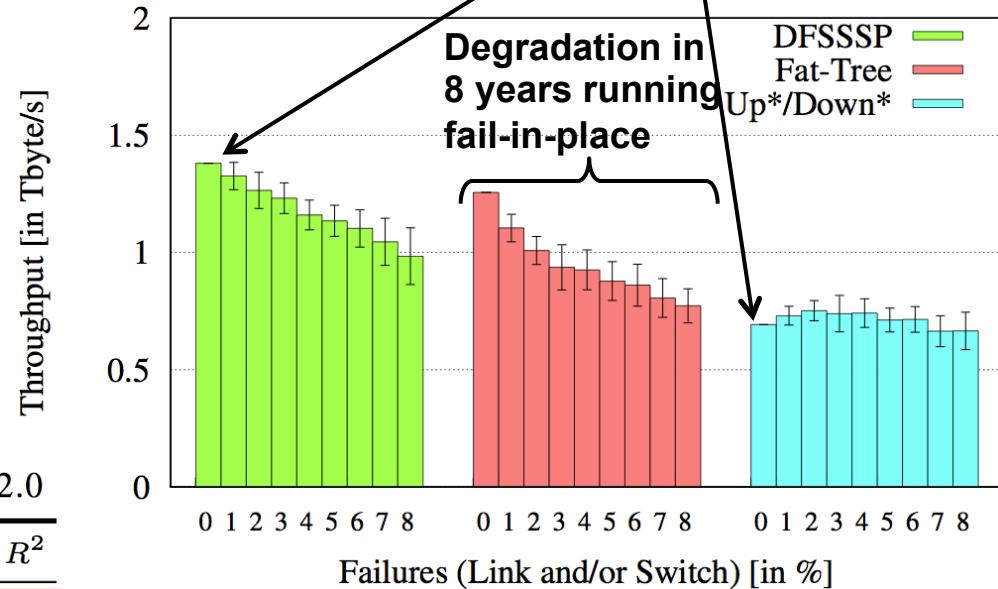
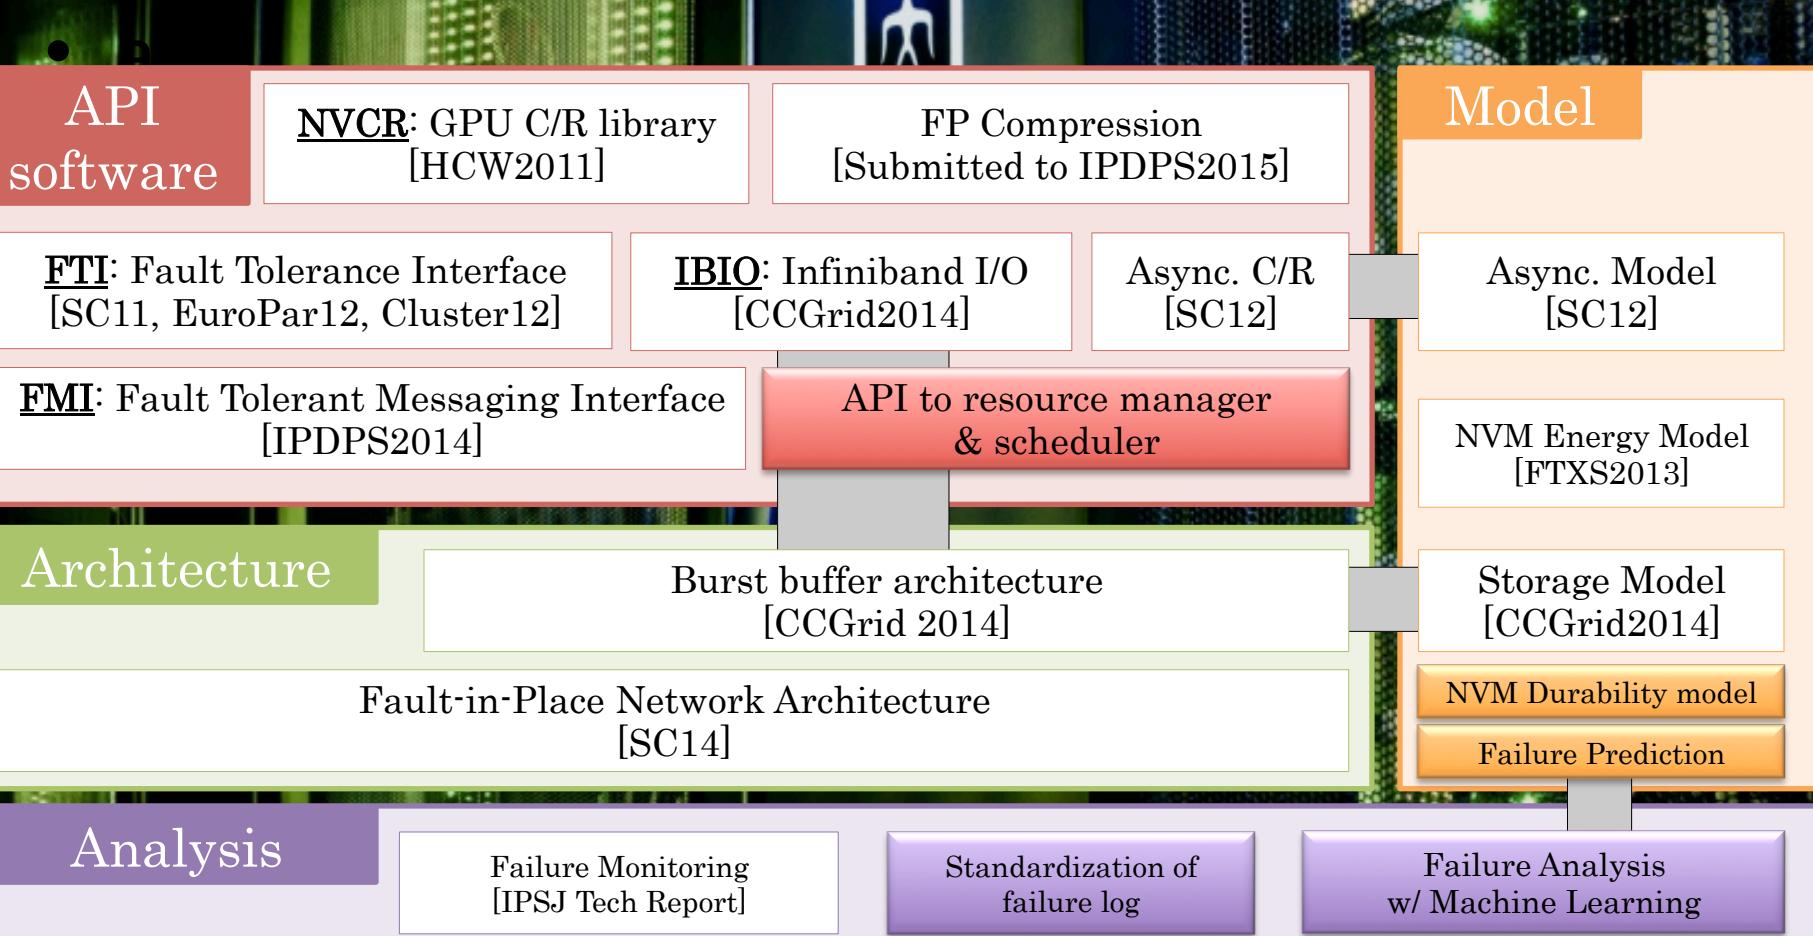


TABLE II. INTERCEPT, SLOPE, AND  $R^2$  FOR TSUBAME2.0

HPC system	Routing	Intercept [in Gbyte/s]	Slope	$R^2$
TSUBAME2.0	DFSSSP	1,393.40	-1.33	0.62
	Fat-Tree	1,187.19	-1.48	0.66
	Up*/Down*	717.76	-0.08	0.01

**Fail-in-Place Network Design is possible!  
(but we have to improve the routing)**

# Tokyo Tech. Billion-Way Resilience Project (2011-2015)



# Awards



## SC11 Technical Paper Perfect Score Award

(Leonardo Batista Gomez, Seiji Tsuboi, Dimitri Komatsch, Frank Cappello, Naoya Maruyama & Satoshi Matsuoka)



## CCGrid2014 Best Paper Award

(Kento Sato, Kathryn Mohror, Adam Moody, Todd Gamblin, Bronis R. de Supinski, Naoya Maruyama & Satoshi Matsuoka)

## API software

I: GPU C/R library [HPCW2011]

FP Compression  
[Submitted to IPDPS2014]

## Model

FTI: Fault Tolerance Interface [SC11, EuroPar12, Cluster12]

IBIO: Infiniband I/O [CCGrid2014]

FMI: Fault Tolerant Messaging Interface [IPDPS2014]

Async. Model [SC12]

Async. Model [SC12]

## Architecture

Burst buffer architecture [CCGrid 2014]

Fault-in-Place Network Architecture [SC14]

Storage Model [CCGrid2014]

NVM Durability model

Failure Prediction

## Analysis

Failure Monitoring [IPSJ Tech Report]

Standardization of failure log

Failure Analysis w/ Machine Learning

# SC14 Technical Paper



Wednesday 2:00PM-2:30PM  
Room 388-89-90

## Fail-in-Place Network Design: Interaction between Topology, Routing Algorithm and Failures

(Jens Domke, Torsten Hoefer, Satoshi Matsuoka)



F<sup>T</sup>I: Fault Tolerance Interface  
[SC11, EuroPar12, Cluster12]

IBIO: Infiniband I/O  
[14]

Async. C/R  
[SC12]

Async. Model  
[SC12]

FMI: Fault Tolerant Messaging Interface  
[IPDPS2014]

to resource manager  
& scheduler

NVM Energy Model  
[FTXS2013]

Architecture

Burst Buffer architecture  
[CCGrid 2014]

Storage Model  
[CCGrid2014]

Fault-in-Place Network Architecture  
[SC14]

NVM Durability model  
Failure Prediction

Analysis

Failure Monitoring  
[IPSJ Tech Report]

Standardization of  
failure log

Failure Analysis  
w/ Machine Learning

# Selected Publications

SC14	J. Domke and T. Hoefer and S. Matsuoka, "Fail-in-Place Network Design: Interaction between Topology, Routing Algorithm and Failures", IEEE/ACM International Conference on High Performance Computing, Networking, Storage and Analysis (SC14), New Orleans, LA, USA, 2014
CCGrid2014	Kento Sato, Kathryn Mohror, Adam Moody, Todd Gamblin, Bronis R. de Supinski, Naoya Maruyama and Satoshi Matsuoka, "A User-level InfiniBand-based File System and Checkpoint Strategy for Burst Buffers", In Proceedings of the 14 <sup>th</sup> IEEE/ACM International Symposium on Cluster, Cloud and Grid Computing (CCGrid2014), Chicago, USA, May, 2014.
IPDPS2014	Kento Sato, Adam Moody, Kathryn Mohror, Todd Gamblin, Bronis R. de Supinski, Naoya Maruyama and Satoshi Matsuoka, "FMI: Fault Tolerant Messaging Interface for Fast and Transparent Recovery", In Proceedings of the International Conference on Parallel and Distributed Processing Symposium 2014 (IPDPS2014), Phoenix, USA, May, 2014.
FTXS2013	Takafumi Saito, Kento Sato, Hitoshi Sato and Satoshi Matsuoka, "Energy-aware I/O Optimization for Checkpoint and Restart on a NAND Flash Memory System", In the Workshop on Fault-Tolerance for HPC at Extreme Scale 2013 (FTXS2013) in conjunction with the International Symposium on High Performance Parallel and Distributed Computing (HPDC13), New York, USA, June, 2013.
IPDPS2013	Improving the computing efficiency of HPC systems using a combination of proactive and preventive checkpointing - Mohamed Slim Bouguerra, Ana Gainaru, Leonardo Bautista-Gomez, Franck Cappello, Naoya Maruyama, Satoshi Matsuoka, IEEE International Parallel & Distributed Processing Symposium 2013 (IPDPS'13), Boston, MA, USA. (Acceptance rate 21.0%)
SNA-MC13	SAMPSON Parallel Computation for Sensitivity Analysis of TEPCO's Fukushima Daiichi Nuclear Power Plant Accident - Marco Pellegrini, Leonardo Bautista-Gomez, Naoya Maruyama, Masanori Naitoh, Satoshi Matsuoka, Franck Cappello, Joint International Conference on Supercomputing in Nuclear Applications and Monte Carlo 2013 (SNA-MC'13), Paris, France.
SC12	Kento Sato, Adam Moody, Kathryn Mohror, Todd Gamblin, Bronis R. de Supinski, Naoya Maruyama and Satoshi Matsuoka, "Design and Modeling of a Non-blocking Checkpointing System", In Proceedings of the International Conference on High Performance Computing, Networking, Storage and Analysis 2012 (SC12), Salt Lake, USA, Nov, 2012.
Cluster2012	Leonardo Bautista Gomez, Thomas Ropars, Naoya Maruyama, Franck Cappello, Satoshi Matsuoka. "Hierarchical Clustering Strategies for Fault Tolerance in Large Scale HPC Systems", In Proc. of IEEE Cluster 2012, IEEE Press, Sep. 2012.
EuroPar2012	L. Bautista Gomez, B. Nicolae, N. Maruyama, F. Cappello, S. Matsuoka. "Scalable Reed-Solomon-based Reliable Local Storage for HPC Applications on IaaS Clouds", In Proc. of International European Conference on Parallel and Distributed Computing (EuroPar 2012), Aug. 2012.
SC11	Leonardo Bautista, Naoya Maruyama, Dimitri Komatitsch, Tsuboi Seiji, Franck Cappello, Satoshi Matsuoka, Nakamura Takeshi. "FTI: High performance Fault Tolerance Interface for hybrid systems". In International Conference for High Performance Computing, Networking, Storage and Analysis (SC).Page 1-12.Nov. 2011.
IPDPSW2011	Nukada, A.; Takizawa, H.; Matsuoka, S., "NVCR: A Transparent Checkpoint-Restart Library for NVIDIA CUDA," Parallel and Distributed Processing Workshops and Phd Forum (IPDPSW), 2011 IEEE International Symposium on , vol., no., pp.104,113, 16-20 May 2011

# Question ?



**SC14**  
New Orleans, LA | **hpc matters.**



## Emerging Technologies Booth

Tuesday 5pm-6pm  
Wednesday 5pm-6pm



**Lawrence Livermore  
National Laboratory**