

Debugging/Testing Non-deterministic MPI Applications

ECP 2nd annual meeting
February 9th

Kento Sato




LLNL-PRES-741293

This work was performed under the auspices of the U.S.
Department of Energy by Lawrence Livermore National
Laboratory under contract DE-AC52-07NA27344. Lawrence
Livermore National Security, LLC


What is non-determinism in MPI applications ?

- Message receive orders change across executions
 - Unpredictable system noise (e.g. network, system daemon & OS jitter)
- Floating point arithmetic orders can also change across executions



Non-determinism increases debugging cost

- Control flows of an application can change across different runs




- Non-deterministic control flow
 - Successful run, seg-fault or hang
- Non-deterministic numerical results
 - Floating-point arithmetic is non-associative

$$(a+b)+c \neq a+(b+c)$$

In non-deterministic applications, it's hard to reproduce bugs and incorrect results.
It costs excessive amounts of time for "reproducing" target behaviors


Non-deterministic bugs cost substantial amounts of time and efforts in MPI applications

Diablo/HYPRE 2.10.1



- The bug manifested in particular clusters
- It hung only once every 30 runs after a few hours
- The scientists spent **2 months in the period of 18 months**, and then **gave up on debugging it**

ParaDis



- The bug intermittently manifested at 100 to 200 iteration
- The scientists **gave up debugging** by themselves

and more ...

How MPI introduces non-determinism ?

- It's typically due to communication with MPI_ANY_SOURCE
- In non-deterministic applications, each MPI rank doesn't know which other MPI ranks will send message and when


Non-deterministic code w/ MPI_ANY_SOURCE

```
MPI_Irecv(..., MPI_ANY_SOURCE, ...);  
while(1) {  
    MPI_Test(flag);  
    if (flag) {  
        <computation>  
        MPI_Irecv(..., MPI_ANY_SOURCE, ...);  
    }  
}
```

CORAL benchmark: MCB (Monte carlo benchmark)


- Use of MPI_ANY_SOURCE is not only source of non-determinism
 - MPI_Waitany/Waitsome/Testany/Testsome also introduce non-determinism

Example: Communications with neighbors



Non-deterministic code w/o MPI_ANY_SOURCE


```
MPI_Irecv(..., north_rank, ..., reqs[0]);  
MPI_Irecv(..., south_rank, ..., reqs[1]);  
MPI_Irecv(..., west_rank, ..., reqs[2]);  
MPI_Irecv(..., east_rank, ..., reqs[3]);  
while(1) {  
    MPI_Testsome(..., reqs, ..., flag, status);  
    if (flag) {  
        ...  
        <computation>  
        for (...) MPI_Irecv(..., status.MPI_SOURCE, ...);  
        ...  
    }  
}
```




MCB: Monte Carlo Benchmark

ReMPI deterministically reproduce order of message receives

- ReMPI is an MPI record-and-replay tool
 - Record an order of MPI message receives
 - Replay the exactly same order of MPI message receives
- Even if a bug manifests in a particular order of message receives, ReMPI can consistently reproduce the target bug
- ReMPI works with other existing debugging tools
 - STAT
 - Parallel debuggers (e.g., Totalview, DDT)





SC15: Kento Sato et al., " Clock delta compression for scalable order-replay of non-deterministic parallel applications"




MCB with/without ReMPI

- Performance metric: How many particles are tracked per second




Recording location: Local SSDe

- ReMPI becomes scalable by recording to local memory/storage
 - Each rank independently writes record → No communication across MPI ranks



ReMPI case study: ParaDiS


- ParaDis
 - non-deterministically crashed after 100 to 200 iterations
- ReMPI reproduced the bug at the exactly same iteration
- ReMPI is interoperable with parallel debuggers and makes debugging non-deterministic bug easier
 - We recorded a buggy behavior in record mode
 - We diagnosed with TotalView under replay mode




ReMPI is also useful for “Testing” ReMPI

- “Testing” is also important for maintaining software quality
- However, non-deterministic MPI applications present significant challenges to testing
 - The non-determinism can produce different results from run to run by nature
 - It’s difficult to reason the different numerical results are due to MPI non-determinism or software bug
- Using ReMPI, computational scientists can easily reproduce MPI behaviors, which facilitate testing

Testing deterministic apps



Testing non-deterministic apps



MPI is not only source of non-determinism

- Applications have been going towards hybrid programming model
 - E.g.) MPI + OpenMP
- OpenMP code adds another level of non-determinism
 - Reduction, critical section or data racy access
 - OpenMP non-determinism affects MPI function call behaviors
 - Need to record both MPI and OpenMP events

CORAL benchmarks

Scalable Science Benchmarks	Priority Level	Lines of Code	Parallelism	
			MPI	OpenMP/ Pthreads
LSMS	TR-1	200,000	X	X
QBOX	TR-1	47,000	X	X
HACC	TR-1	35,000	X	X
Nekbone	TR-1	48,000	X	

Throughput Benchmarks	Priority Level	Lines of Code	Parallelism	
			MPI	OpenMP/ Pthreads
CAM-SE	TR-1	150,000	X	X
UMT2013	TR-1	51,000	X	X
AMG2013	TR-1	75,000	X	X
MCB	TR-1	13,000	X	X
QMCPACK	TR-2	200,000	X	X
NAMD	TR-2	180,000	X	X
LULESH	TR-2	5,000	X	X
SNAP	TR-2	3,000	X	X
miniFE	TR-2	50,000	X	X




MPI+OpenMP


Providing record-and-replay
for MPI+OpenMP application is our future work

Unintended message races in MPI

- Many applications are written as a series of communication and computation routines executed by all processes (i.e., data parallel, SPMD)
- Developers must make sure all communication routines are “isolated”
- Example (Routine A and Routine B)
 - Different MPI_TAG or synchronization (e.g. MPI_Barrier) between the two routines
- If not isolated, message race bugs potentially occur
 - E.g.) A message sent in Routine B is received in Routine A
- Unintended message races are non-deterministic and infrequently occur



- NINJA (Noise Injection Agent Tool) exposes message race bugs by injecting noise




- NINJA detects suspicious communication routines
 - Communication routine using the same MPI_TAG without synchronization
- NINJA injects a delay to MPI messages in order to enforce message races
- NINJA can test if the application has unintended message races


NINJA cast study: Diablo/HYPRE-2.10.1(in ParaSail module)

- Unintended message races in Hypre
- Prior to NINJA, the bug does not manifest itself in Hypre test code
- NINJA consistently exposed message races in the test code

Diablo/HYPRE 2.10.1



The bug manifested in particular machines
It hung only once every 30 runs after a few hours
The scientists spent **2 months in the period of 18 months**, and then **gave up on debugging it**



Summary

- Non-determinism in MPI applications costs significant time for debugging and testing
- ReMPI and NINJA facilitate debugging/testing non-deterministic MPI applications
 - ReMPI is MPI record-and-replay for reproducing particular errors
 - NINJA is an noise injection tool for exposing message-race bugs
- We will extend our tools for supporting MPI+OpenMP applications in future



PRUNERS ReMPI



OR <https://github.com/PRUNERS/ReMPI>



NINJA

PRUNERS NINJA



OR <https://github.com/PRUNERS/NINJA>



**Lawrence Livermore
National Laboratory**