

PRUNERS

Providing Reproducibility for Uncovering Non-Deterministic Errors in Runs on Supercomputers

PMCD/JOWOG February 14th, 2018

Dong H. Ahn (PI) , Chris Chambreau, Ignacio Laguna, Gregory L. Lee, Kento Sato,
Martin Schulz, Simone Atzeni, Michael Bentley, Goef Sawaya, Ganesh
Gopalakrishnan, Zvonimir Rakamaric, Joachim Protze



LLNL-PRES-741293

This work was performed under the auspices of the U.S.
Department of Energy by Lawrence Livermore National
Laboratory under contract DE-AC52-07NA27344. Lawrence
Livermore National Security, LLC

Debugging/Testing large-scale applications is challenging

“On average, software developers spend
50% of their programming time finding and fixing bugs.” [1]



Debugging/Testing are inevitable software development processes.
Tools facilitating Debugging/Testing are indispensable

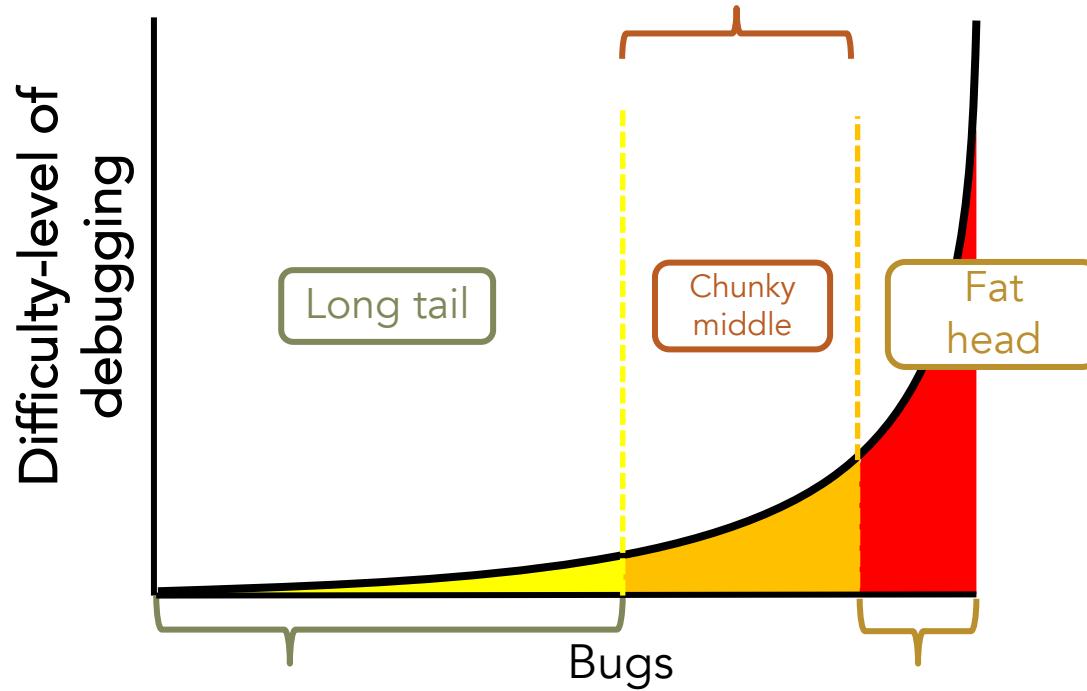
[1] Source: <http://www.prweb.com/releases/2013/1/prweb10298185.htm>, CAMBRIDGE, UK (PRWEB) JANUARY 08, 2013

Easy-to-fix bugs VS. Difficult-to-fix bugs



Totalview/DDT

Some bugs require to use Totalview or DDT



Many bugs can be quickly fixed by simple printf debugging

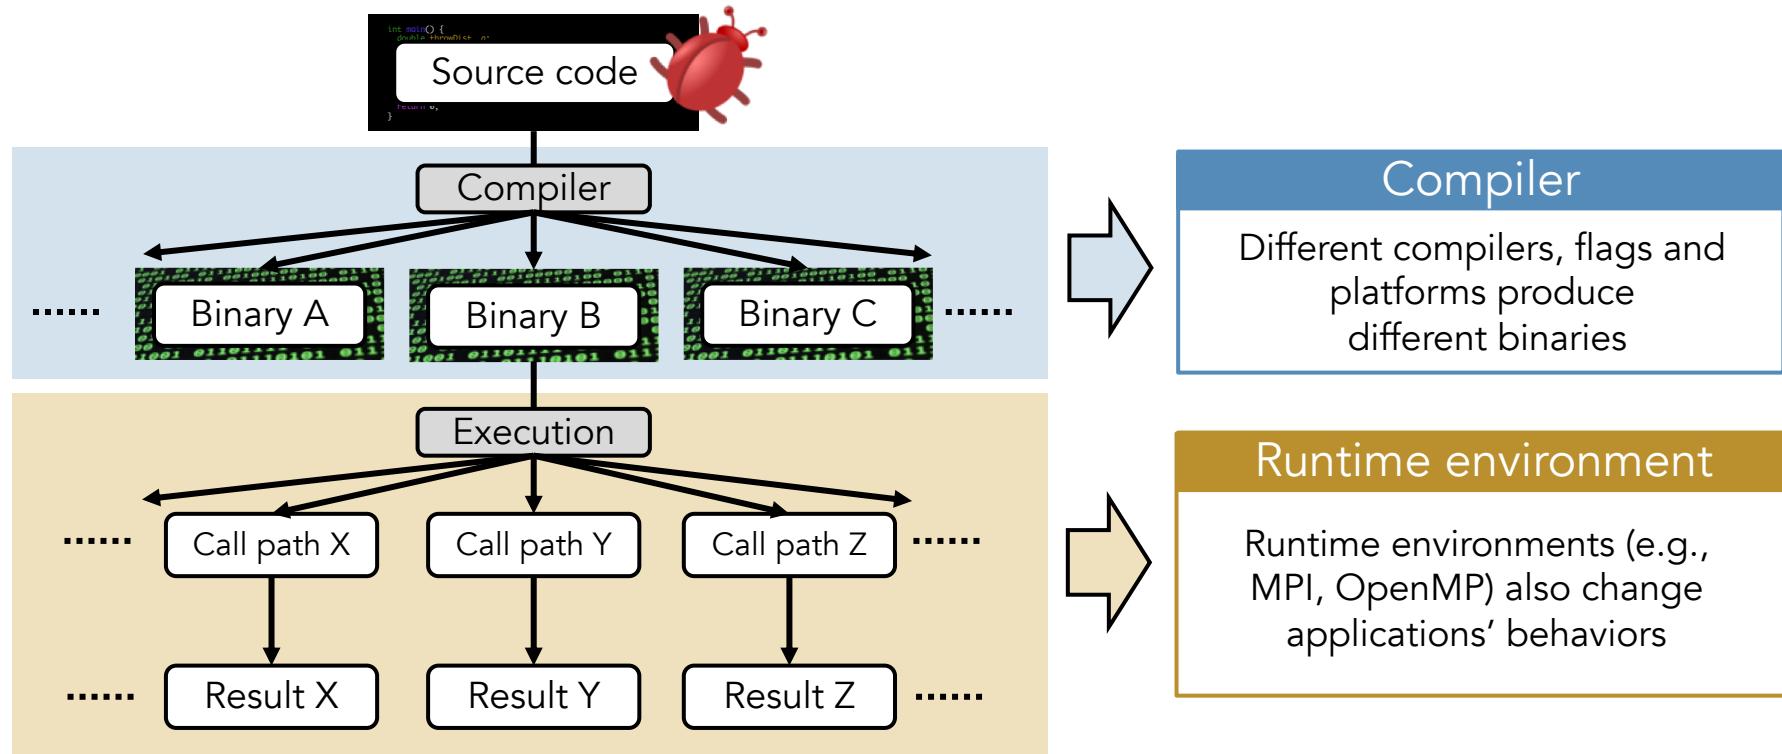


printf

~~printf~~ ~~Totalview/DDT~~

Bugs are not created equal !

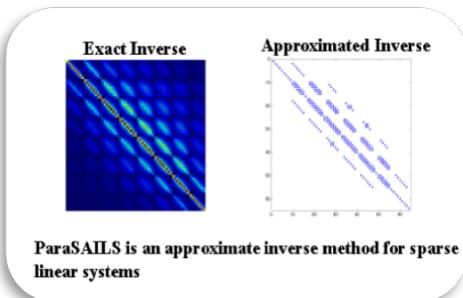
- When debugging/testing, reproducibility is very important



- Examples
 - Bugs that manifest themselves when using `-O3`, but do not with `-O0`
 - Bugs that manifest in a run and do not in the next run

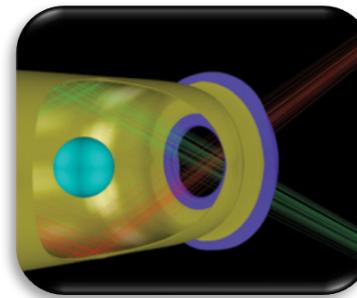
Non-deterministic bugs cost substantial amounts of time and efforts

Diablo/HYPRE 2.10.1



- The bug manifested in particular machines
- It hung only once every 30 runs after a few hours
- The scientists spent **2 months** in **the period of 18 months**, and then **gave up on debugging it**

HYDRA (porting on Sequoia)

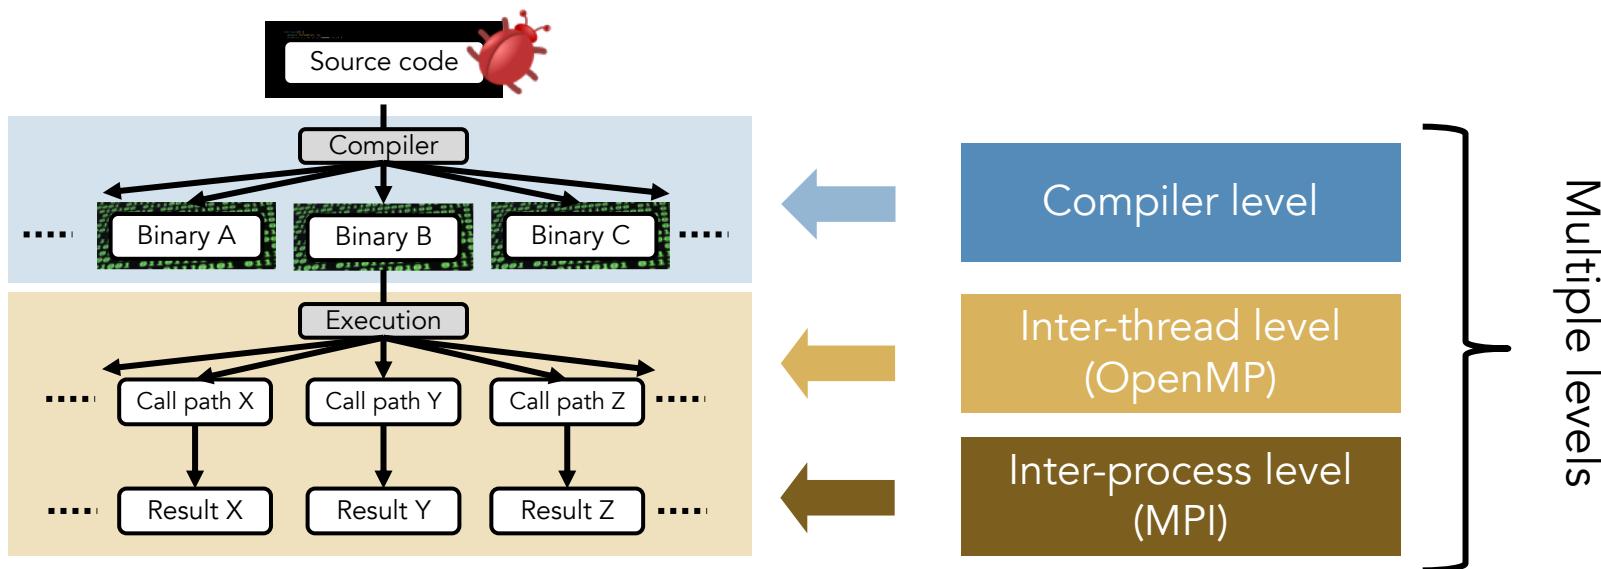


- MPI/OpenMP application non-deterministically crashed in an OpenMP region when compiling with optimization levels
- Manifested intermittently at or above 8K MPI processes
- The scientists spent **months**, and then **ended up disabling OpenMP**

and more ...

Non-deterministic bugs are introduced at multiple levels

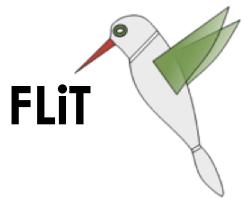
- Introduced at the compiler level or at different runtime levels
- A monolithic tool won't work for all cases
- Debugging/testing toolset
 - Individual tool works effectively
 - Interoperable and composable each other
 - Make debugging/testing easier even under other existing debuggers



PRUNERS

Multi-level debugging/testing capabilities

The PRUNERS Toolset comprises four individual tools



Compiler-induced floating-point computation variability tester

Data race detector for OpenMP programs

MPI record-and-replay tool
for reproducing non-deterministic MPI bugs

Noise injection tool
for exposing message race bugs

Compiler
level

OpenMP

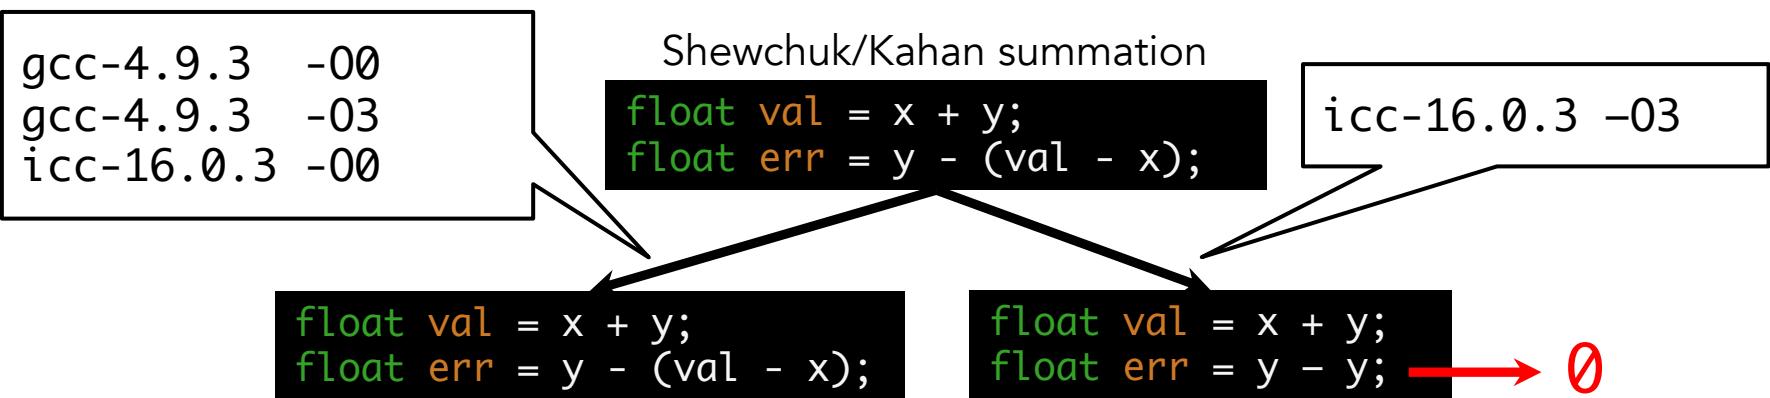
MPI

Multiple levels



Different compilers, compiler flags and platforms produce different numerical results

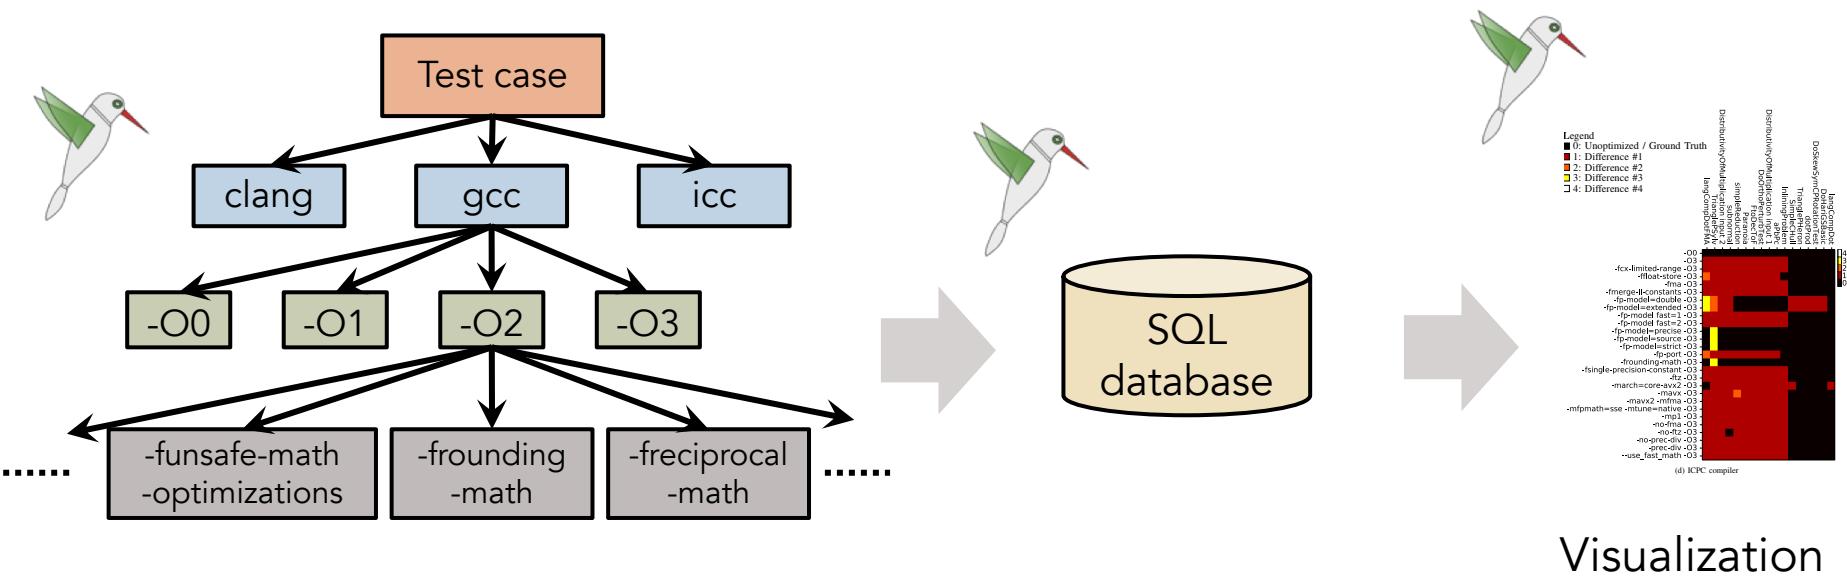
- No guarantee that floating-point computation on one platform is the same on another platform
 - E.g.) Apply associativity rules of real arithmetic



It's important to understand how sensitive your numerical algorithm is w.r.t. round-off errors introduced by different compilers for code verification and validation

FLiT (Floating-point Litmus Tester)

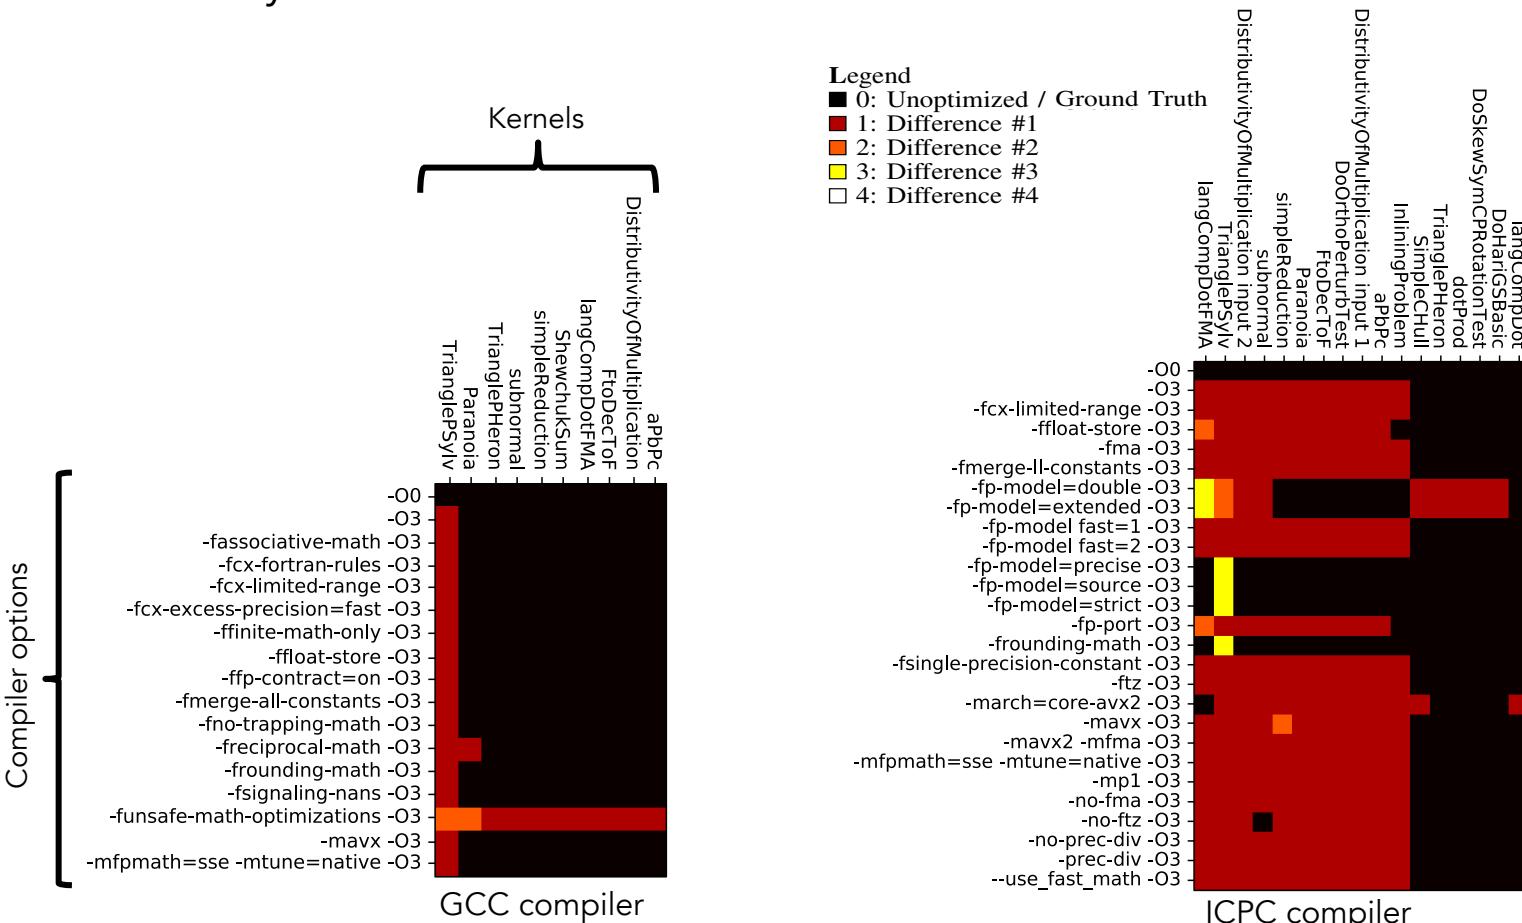
- FLiT is a reproducibility test framework
 - Test floating-point arithmetic variability in any user-given collection of programs
- FLiT tests the variability through hundreds of combinations
 - Different compilers, compiler flags, and also different hosts
- Results are stored in SQL database and used for visualization and for further analysis



Visualization

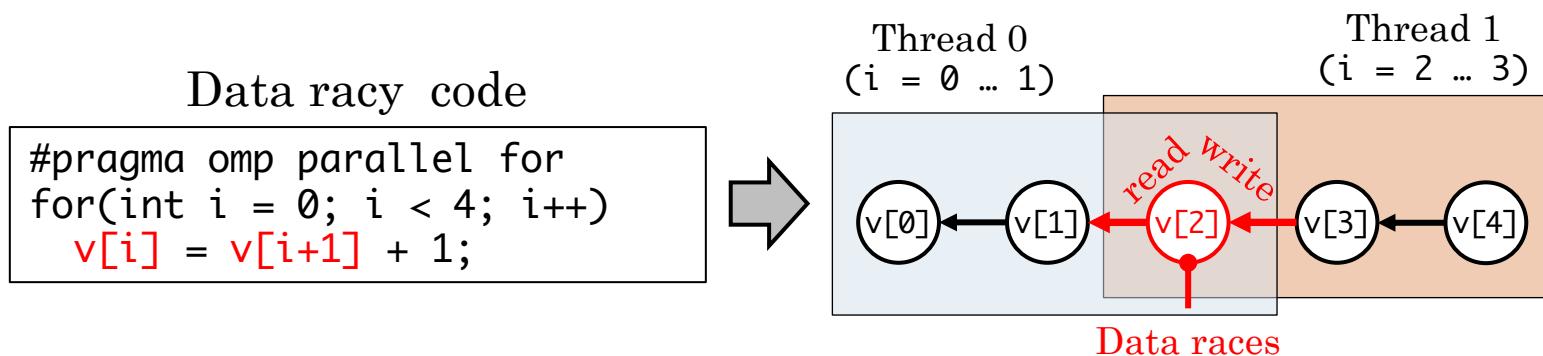
FLiT case study

- We tested several kernels which have compiler-induced FP variability
 - Difference in numerical results across different compilers, flags and kernels
- Example
 - When you want to find a compiler option that makes your applications faster while reproducing the exactly same results as non-optimized code, FLiT becomes very useful tool



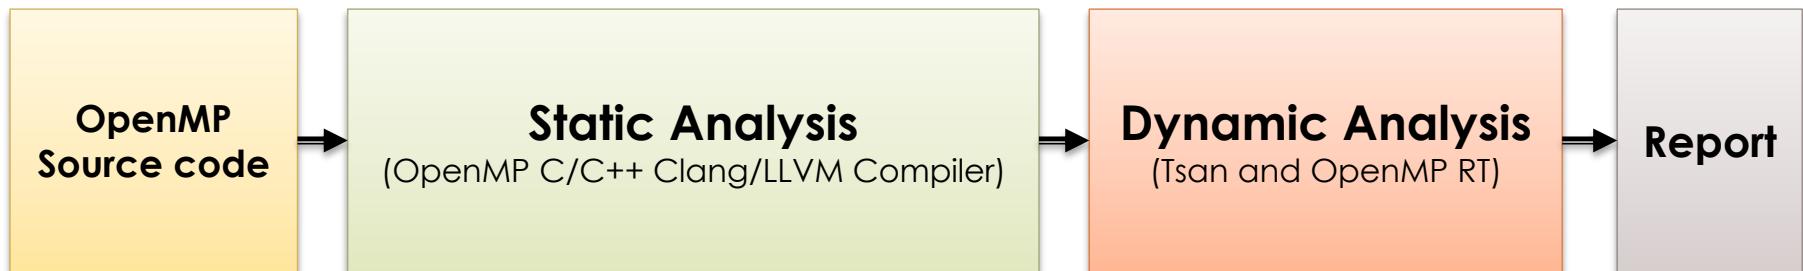
OpenMP easily creates non-deterministic bugs

- Data races in OpenMP are the most malignant non-deterministic bugs
- Depending on orders of read and write, numerical results change
- Orders of read and write are non-deterministic, it introduces non-deterministic bugs



In large applications,
it is difficult to manually identify data races

- Archer is a data race detector combining static and dynamic analysis



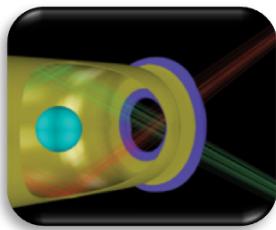
- Static analysis
 - Exclude regions that can be statically detected to be race-free for dynamic analysis (Blacklisting)
- Dynamic analysis
 - Detect data races based on LLVM/Clang ThreadSanitizer combined with OMPT-based annotation

Archer significantly reduce runtime overhead while providing high precision and accuracy

Archer case study: HYDRA

- Archer easily detected data races !

HYDRA (porting on Sequoia)

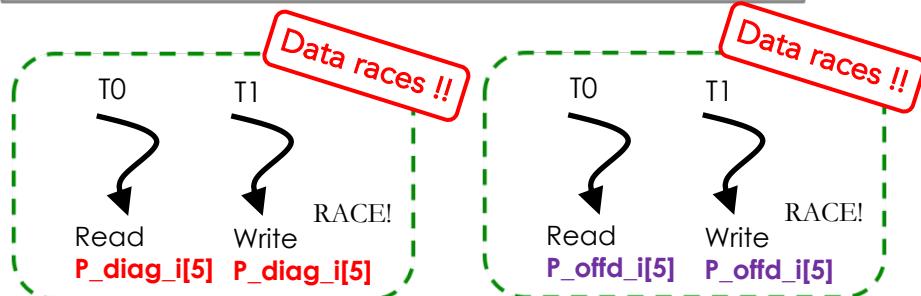


- MPI/OpenMP application non-deterministically crashed in an OpenMP region when compiling with optimization levels
- Manifested intermittently at or above 8K MPI processes
- The scientists spent **months**, and then **ended up disabling OpenMP**

Archer

```
1002: hypre_BoomerAMGInterpTruncation(...) {  
...  
1007:   int *P_diag_i = hypre_CSRMatrixI(P_diag);  
1014:   int *P_offd_i = hypre_CSRMatrixI(P_offd);  
  
1062: #pragma omp parallel private(...)  
1064: {  
...  
1172:   if(max_elmst>0) {  
...  
1179:     for(i=start; i<stop; i++) {  
...  
1183:       last_index = P_diag_i[i+1];  
1184:       last_index_offd = P_offd_i[i+1];  
1248:       P_diag_i[i] = cnt_diag;  
1249:       P_offd_i[i] = cnt_offd;  
...  
1484:   } /* end_parallel_region */  
1491:   return ieer;  
1492: }
```

423 lines
About 50 variables



How MPI introduces non-determinism ?

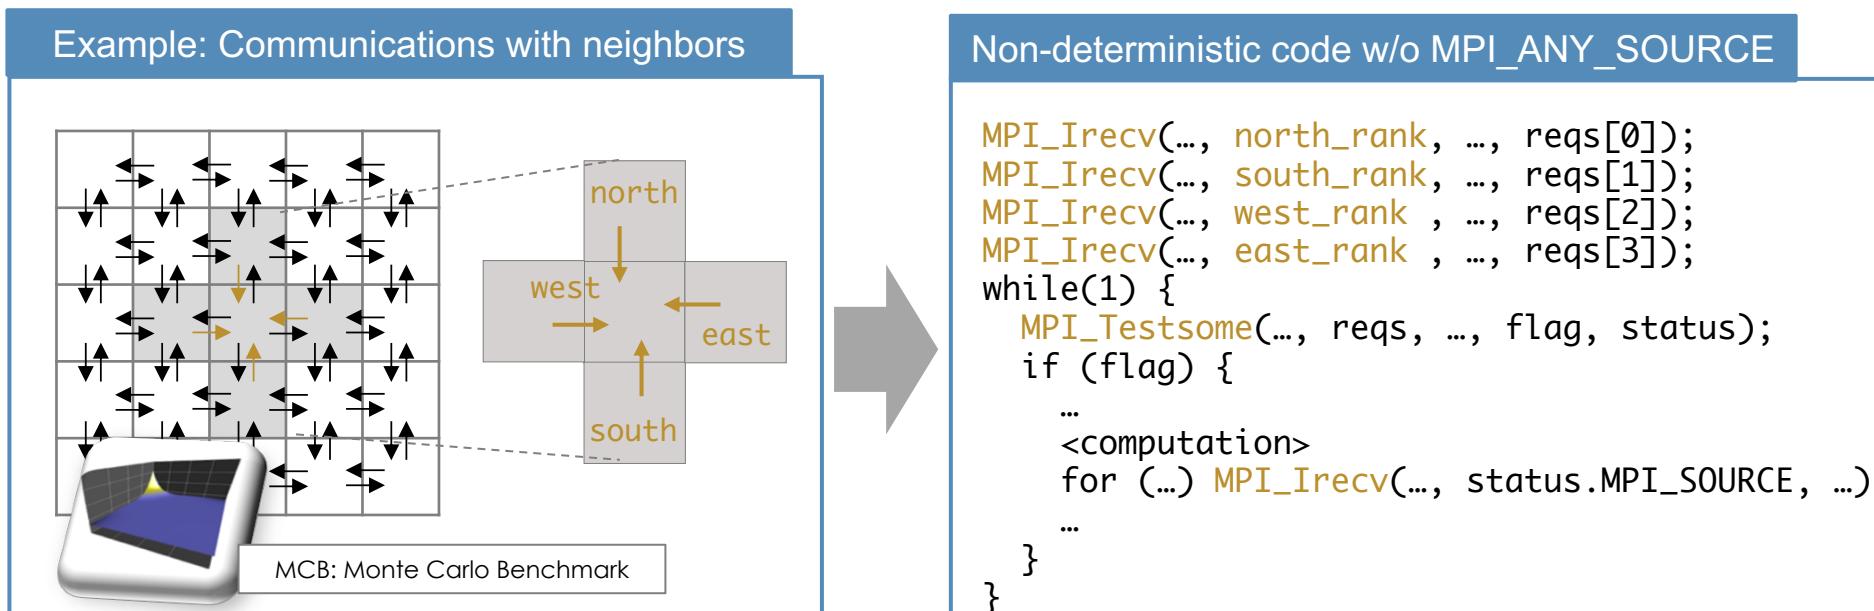
- It's typically due to communication with MPI_ANY_SOURCE
- In non-deterministic applications, each MPI rank doesn't know which other MPI ranks will send message and when

Non-deterministic code w/ MPI_ANY_SOURCE

```
MPI_Irecv(..., MPI_ANY_SOURCE, ...);  
while(1) {  
    MPI_Test(flag);  
    if (flag) {  
        <computation>  
        MPI_Irecv(..., MPI_ANY_SOURCE, ...);  
    }  
}
```

CORAL benchmark: MCB (Monte carlo benchmark)

- Use of MPI_ANY_SOURCE is not only source of non-determinism
 - MPI_Waitany/Waitsome/Testany/Testsome also introduce non-determinism
- MCB produces different numerical results from run to run



In non-deterministic applications,
if a bug manifests through a particular message receive order,
it's hard to reproduce the bug, thereby, hard to debug it

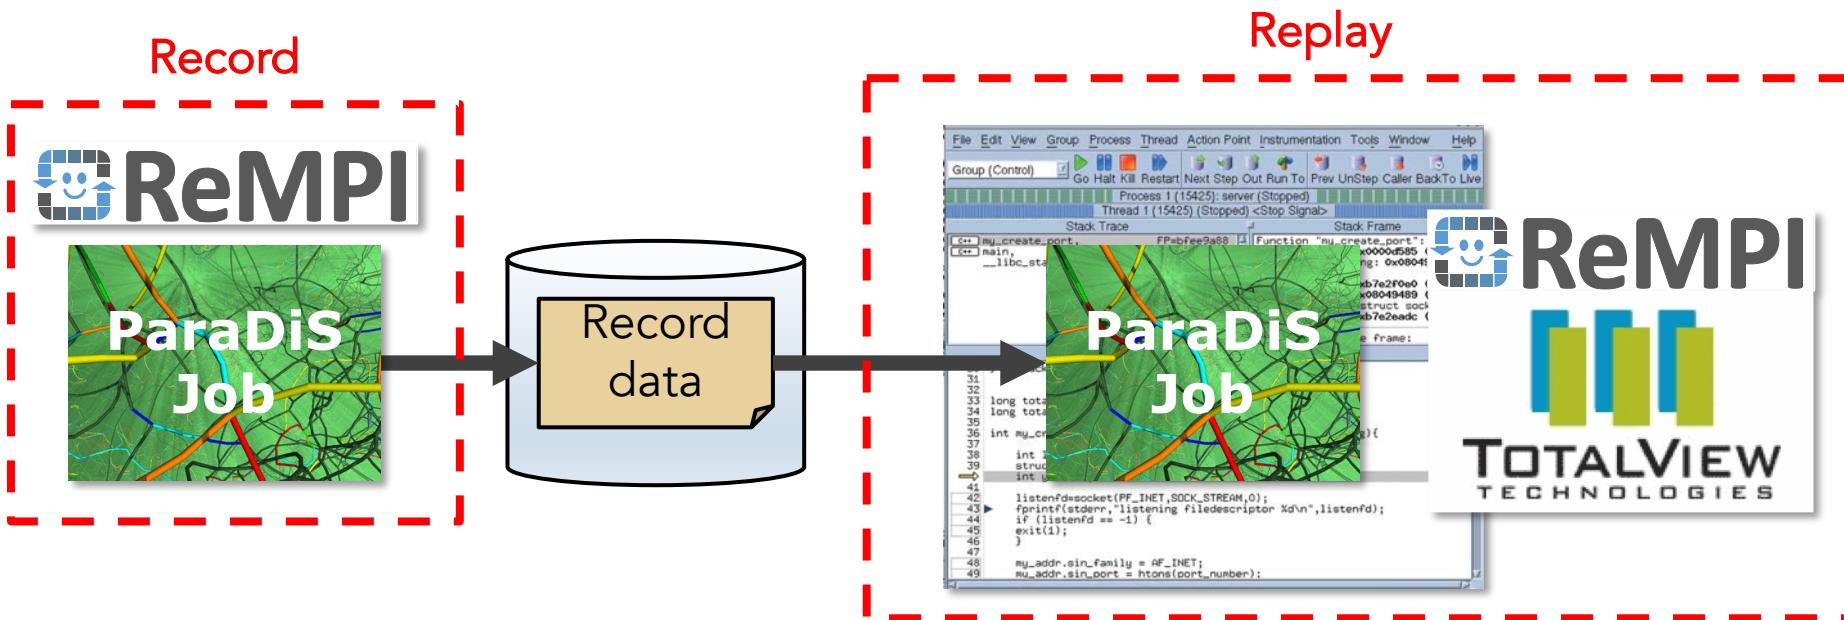
ReMPI deterministically reproduce order of message receives

- ReMPI is an MPI record-and-replay tool
 - Record an order of MPI message receives
 - Replay the exactly same order of MPI message receives
- Even if a bug manifests in a particular order of message receives, ReMPI can consistently reproduce the target bug



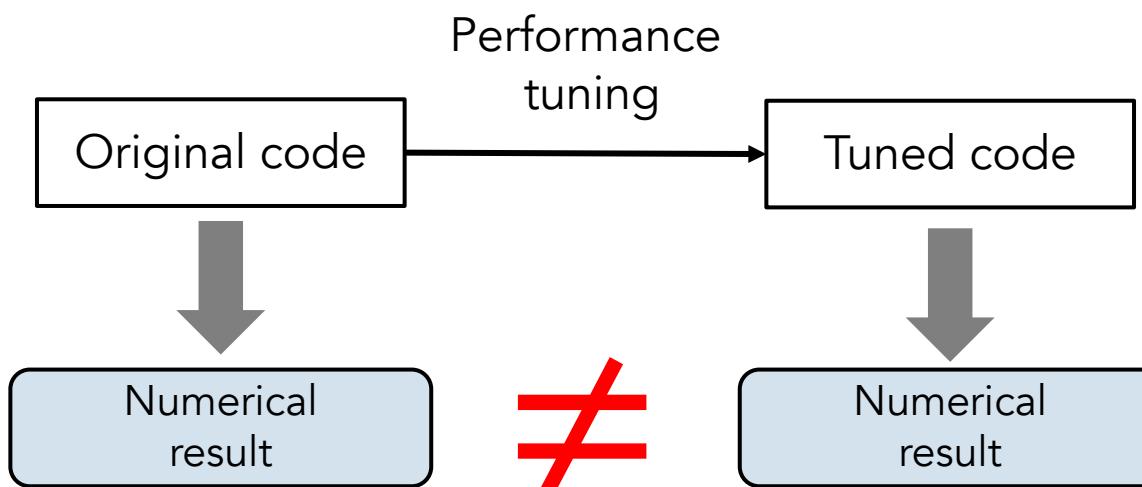
ReMPI case study: ParaDiS

- ParaDis
 - non-deterministically crashed after 100 to 200 iterations
- ReMPI reproduced the bug at the exactly same iteration
- ReMPI is interoperable with parallel debuggers and makes debugging non-deterministic bug easier
 - We recorded a buggy behavior in record mode
 - We diagnosed with TotalView under replay mode



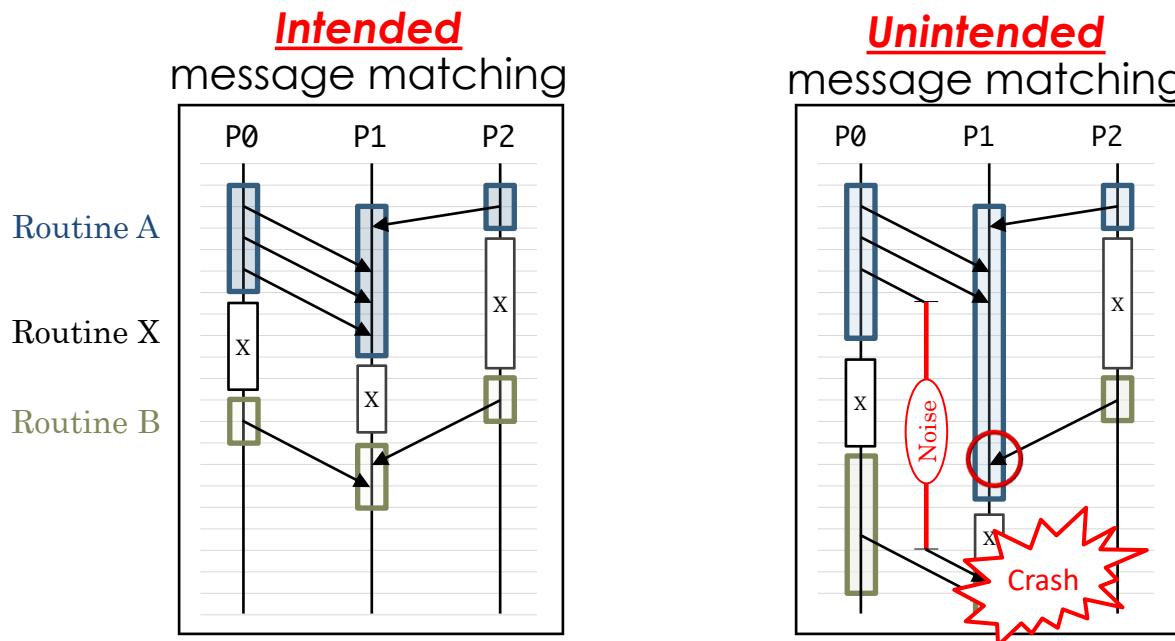
ReMPI is also useful for “Testing” ReMPI

- “Testing” is also important for maintaining software quality
- However, non-deterministic MPI applications present significant challenges to testing
 - The non-determinism can produce different results from run to run by nature
 - It’s difficult to reason the different numerical results are due to MPI non-determinism or software bug
- Using ReMPI, computational scientists can easily reproduce MPI behaviors, which facilitate testing

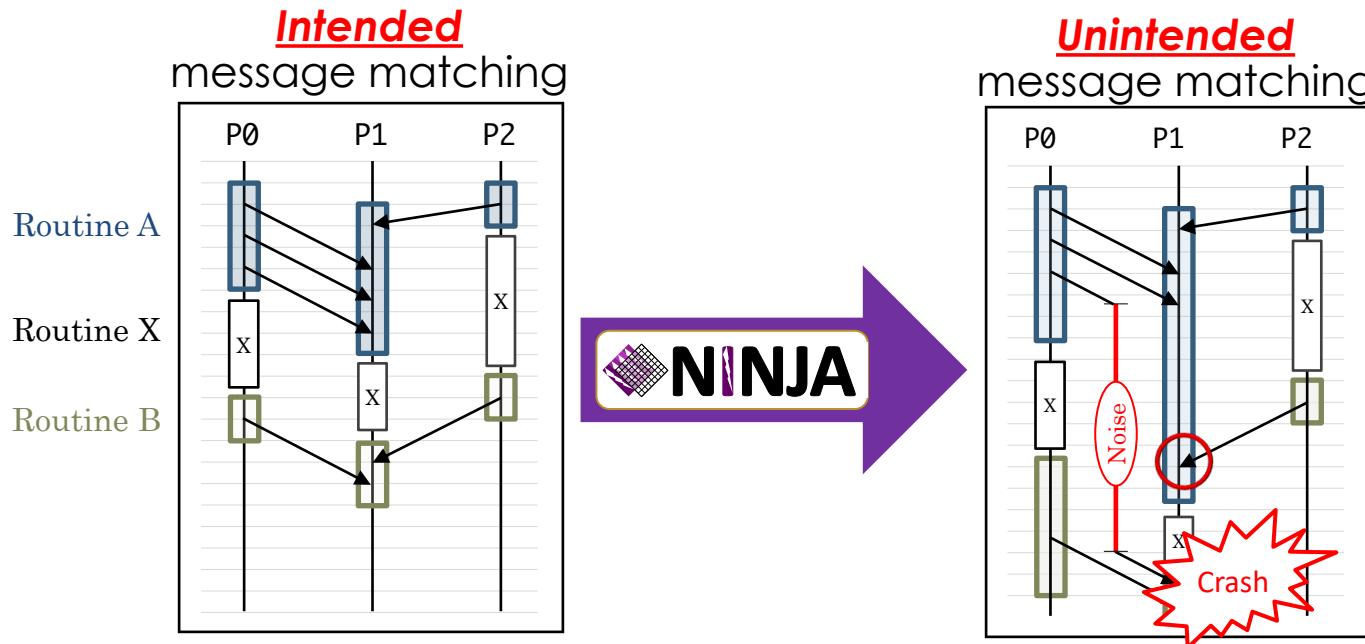


Unintended message races in MPI

- Many applications are written as a series of communication and computation routines executed by all processes (i.e., data parallel, SPMD)
- Developers must make sure all communication routines are “isolated”
- Example (Routine A and Routine B)
 - Different MPI_TAG or synchronization (e.g. MPI_Barrier) between the two routines
- If not isolated, message race bugs potentially occur
 - E.g.) A message sent in Routine B is received in Routine A
- Unintended message races are non-deterministic and infrequently occur



- NINJA (Noise Injection Agent Tool) exposes message race bugs by injecting noise

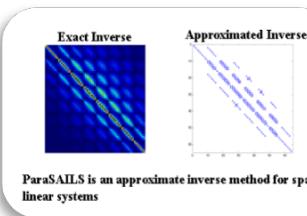


- NINJA detects suspicious communication routines
 - Communication routine using the same MPI_TAG without synchronization
- NINJA injects a delay to MPI messages in order to enforce message races
- NINJA can test if the application has unintended message races

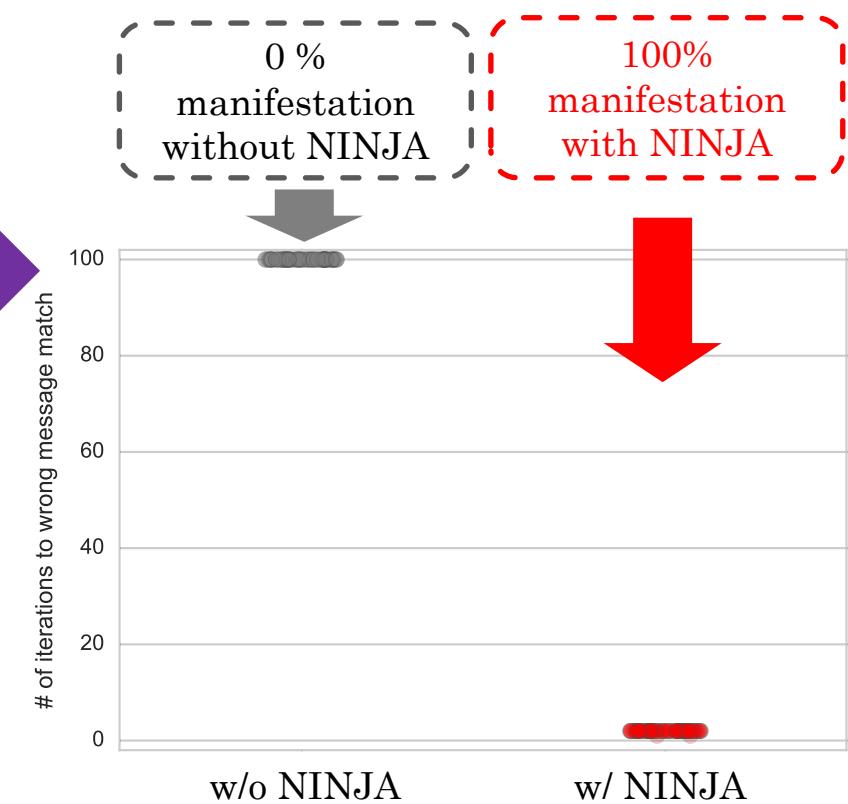
NINJA cast study: Diablo/HYPRE-2.10.1(in ParaSail module)

- Unintended message races in Hypre
- Prior to NINJA, the bug does not manifest itself in Hypre test code
- NINJA consistently exposed message races in the test code

Diablo/HYPRE 2.10.1

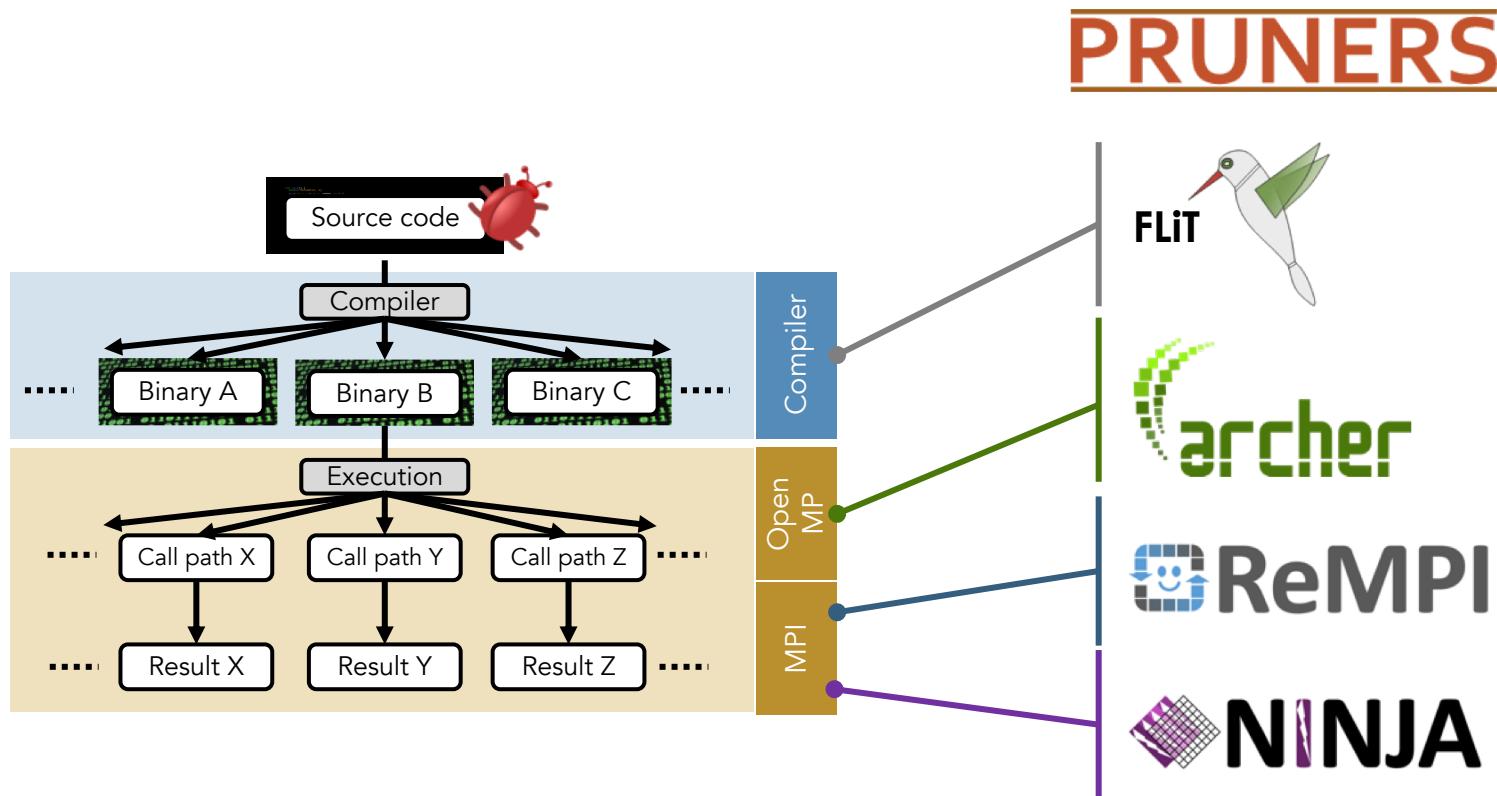


The bug manifested in particular machines
It hung only once every 30 runs after a few hours
The scientists spent **2 months in the period of 18 months**, and then **gave up on debugging it**



Summary

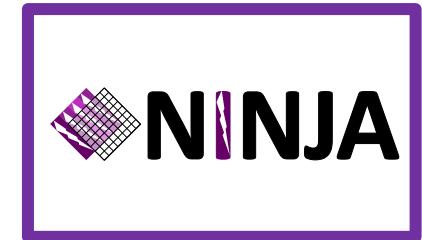
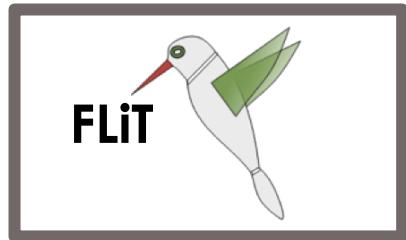
- Debugging and testing are becoming more challenging due to non-determinism in HPC applications
- The PRUNERS toolset facilitates debugging and testing for non-deterministic applications



PRUNERS

<https://pruners.github.io/>

PRUNERS toolset 



Dong H. Ahn (PI)
Chris Chambreau
Ignacio Laguna
Gregory L. Lee
Kento Sato
Martin Schulz



Simone Atzeni
Michael Bentley
Goef Sawaya
Prof. Ganesh Gopalakrishnan
Prof. Zvonimir Rakamaric



Joachim Protze



**Lawrence Livermore
National Laboratory**