

PRUNERS

Providing Reproducibility for Uncovering Non-Deterministic Errors in Runs on Supercomputers

Computational Reproducibility at Exascale (CRE2017)

November 12th, 2017

Kento Sato



LLNL-PRES-741293

This work was performed under the auspices of the U.S.
Department of Energy by Lawrence Livermore National
Laboratory under contract DE-AC52-07NA27344. Lawrence
Livermore National Security, LLC

Debugging/Testing large-scale applications is challenging

“On average, software developers spend
50% of their programming time finding and fixing bugs.” [1]




Debugging/Testing are inevitable software development processes.
Tools facilitating Debugging/Testing are indispensable

[1] Source: <http://www.prweb.com/releases/2013/1/prweb10298185.htm>, CAMBRIDGE, UK (PRWEB) JANUARY 08, 2013

Bugs are not created equal !


- When debugging/testing, reproducibility is very important



- Examples
 - Bugs that manifest themselves when using `-O3`, but do not with `-O0`
 - Bugs that do not manifest themselves


Non-deterministic bugs cost substantial amounts of time and efforts

Diablo/HYPRE 2.10.1



- The bug manifested in particular machines
- It hung only once every 30 runs after a few hours
- The scientists spent **2 months** in **the period of 18 months**, and then **gave up on debugging it**

HYDRA (porting on Sequoia)




- MPI/OpenMP application non-deterministically crashed in an OpenMP region when compiling with optimization levels
- Manifested intermittently at or above 8K MPI processes
- The scientists spent **months**, and then **ended up disabling OpenMP**

and more ...

Non-deterministic bugs are introduced at multiple levels


- Introduced at the compiler level or at different runtime levels
- A monolithic tool won't work for all cases
- Debugging/testing toolset
 - Individual tool works effectively
 - Interoperable and composable each other
 - Make debugging/testing easier even under other existing debuggers



PRUNERS

Multi-level debugging/testing capabilities

The PRUNERS Toolset comprises four individual tools that can co-operate



Compiler-induced floating-point computation variability tester

Data race detector for OpenMP programs

MPI record-and-replay tool
for reproducing non-deterministic MPI bugs

Noise injection tool
for exposing message race bugs

Compiler
level


OpenMP

MPI

Multiple levels

Different compilers, compiler flags and platforms produce different numerical results


- No guarantee that floating-point computation on one platform is the same on another platform
 - E.g.) Apply associativity rules of real arithmetic



Understanding how sensitive your algorithm is w.r.t. to different round-off errors introduced by different compilers and flags are important for code verification and validation


FLiT (Floating-point Litmus Tester)

- FLiT is a reproducibility test framework
 - Test floating-point arithmetic variability in any user-given collection of programs
- FLiT tests the variability through hundreds of combinations
 - Different compilers, compiler flags, and also different hosts
- Results are stored in SQL database and used for visualization and for further analysis




FLiT case study

- We tested several kernels which have compiler-induced FP variability
 - Difference in numerical results across different compilers, flags and kernels
 - Example
 - When you want to find a compiler option that makes your applications faster while reproducing the exactly same results as non-optimized code, FLiT becomes very useful tool




OpenMP easily creates non-deterministic bugs

- Data races in OpenMP are the most malignant non-deterministic bugs
- Depending on orders of read and write, numerical results change
- Orders of read and write are non-deterministic, it introduces non-deterministic bugs



In large-scale applications,
it is difficult to identify data races manually

- Archer is a data race detector combining static and dynamic analysis




- Static analysis
 - Exclude regions that can be statically detected to be race-free for dynamic analysis (Blacklisting)
- Dynamic analysis
 - Detect data races based on LLVM/Clang ThreadSanitizer combined with OMPT-based annotation

Archer significantly reduce runtime overhead while providing high precision and accuracy

Archer case study: HYDRA

- Archer easily detected data races !

HYDRA (porting on Sequoia)




- MPI/OpenMP application non-deterministically crashed in an OpenMP region when compiling with optimization levels
- Manifested intermittently at or above 8K MPI processes
- The scientists spent **months**, and then **ended up disabling OpenMP**

Archer

```
1002: hypre_BoomerAMGInterpTruncation(...) {  
...  
1007:   int *P_diag_i = hypre_CSRMatrixI(P_diag);  
1014:   int *P_offd_i = hypre_CSRMatrixI(P_offd);  
  
1062: #pragma omp parallel private(...)  
1064: {  
...  
1172:   if(max_elmst>0) {  
...  
1179:     for(i=start; i<stop; i++) {  
...  
1183:       last_index = P_diag_i[i+1];  
1184:       last_index_offd = P_offd_i[i+1];  
1248:       P_diag_i[i] = cnt_diag;  
1249:       P_offd_i[i] = cnt_offd;  
...  
1484:   } /* end parallel region */  
1491:   return ieer;  
1492: }
```

423 lines

About 50 variables




MPI can also introduce non-determinism

- It's typically due to communication with MPI_ANY_SOURCE
- In non-deterministic applications, each MPI rank doesn't know which other MPI rank will send message and when
- Example
 - If processes communicate with neighbors, messages can arrive in any order from neighbors

MPI_ANY_SOURCE communication


```
MPI_Irecv(..., MPI_ANY_SOURCE, ...);  
while(1) {  
    MPI_Test(flag);  
    if (flag) {  
        <computation>  
        MPI_Irecv(..., MPI_ANY_SOURCE, ...);  
    }  
}
```



If a bug manifests through a particular message receive order,
It's hard to reproduce the bug, thereby, hard to debug it

ReMPI deterministically reproduces order of message receives


- ReMPI is an MPI record-and-replay tool
 - Record an order of MPI message receives
 - Replay the exactly same order of MPI message receives
- Even if a bug manifests in a particular order of message receives, ReMPI can consistently reproduce the target bug



 **ReMPI**

ReMPI case study: ParaDiS

- ParaDis
 - non-deterministically crashed after 100 to 200 iterations
- ReMPI reproduced the bug at the exactly same iteration
- ReMPI is interoperable with parallel debuggers and makes debugging non-deterministic bug easier
 - We recorded a buggy behavior in record mode
 - We diagnosed with TotalView under replay mode




ReMPI is also useful for “Testing”




- “Testing” is also important for maintaining software quality
- MPI non-deterministic applications present significant challenges to testing
 - The non-determinism can produce different results from run to run by nature
- Using ReMPI, computational scientists can easily reproduce MPI behaviors, which facilitate testing

Testing deterministic apps




Testing non-deterministic apps




Unintended message races in MPI

- Many applications are written as a series of communication and computation routines (i.e., data parallel, SPMD)
- Developers must make sure all communication routines are “isolated”
- Example (Routine A and Routine B)
 - Different MPI_TAG or synchronization (e.g. MPI_Barrier) between the two routines
- If not isolated, message race bugs potentially occur
 - E.g.) A message sent in Routine B is received in Routine A
- Unintended message races are non-deterministic and infrequently occur



NINJA: Noise Injection Agent Tool

- NINJA exposes message race bugs by injecting noise




- NINJA detects suspicious communication routines
 - Communication routine using the same MPI_TAG without synchronization
- NINJA injects a delay to MPI messages in order to enforce message races
- NINJA can test if the application has unintended message races


NINJA cast study: Diablo/HYPRE-2.10.1(in ParaSail module)

- Unintended message races in Hypre
- Prior to NINJA, the bug does not manifest itself in Hypre test code
- NINJA consistently exposed message races in the test code

Diablo/Hypre 2.10.1




The bug manifested in particular machines
It hung only once every 30 runs after a few hours
The scientists spent **2 months in the period of 18 months**, and then **gave up on debugging it**



Summary

- Debugging and testing large-scale HPC applications are becoming more challenging
- The PRUNERS toolset facilitates debugging and testing
- Exscale computing will be the culmination of non-deterministic execution for unprecedentedly high performance, and PRUNERS leads the way to solve its debugging and testing challenges.



PRUNERS



PRUNERS

<https://pruners.github.io/>

PRUNERS toolset 



Dong H. Ahn
Chris Chambreau
Ignacio Laguna
Gregory L. Lee
Kento Sato
Martin Schulz



Simone Atzeni
Michael Bentley
Goef Sawaya
Prof. Ganesh Gopalakrishnan
Prof. Zvonimir Rakamaric



Joachim Protze



**Lawrence Livermore
National Laboratory**